

Artificial Neurons

Deep Learning
CS 435/635

Course Instructor: Chandresh
AI Lab Coordinator @IIT Indore

Course Content

Module I: History of Deep Learning, **Sigmoid Neurons, Perceptrons**, and learning algorithms. Multilayer Perceptrons (MLPs), Representation Power of MLPs,.

Module II: Feedforward Neural Networks. Backpropagation. first and second-order training methods. NN Training tricks.

Module III: Introduction to Autoencoders and their characteristics, relation to PCA, Regularization in autoencoders, and Types of autoencoders.

Module IV: Architecture of Convolutional Neural Networks (CNN), types of CNNs.

Module V: Architecture of Recurrent Neural Networks (RNN), Backpropagation through time. Encoder-Decoder Models, Attention Mechanism. Advanced Topics: Transformers and BERT.

Acknowledgement

- Neural Networks and Deep Learning course by Danna Gurari University of Colorado Boulder

Today's Topics

- Binary classification applications
- Evaluating classification models
- Biological neurons: inspiration
- Artificial neuron: Perceptron

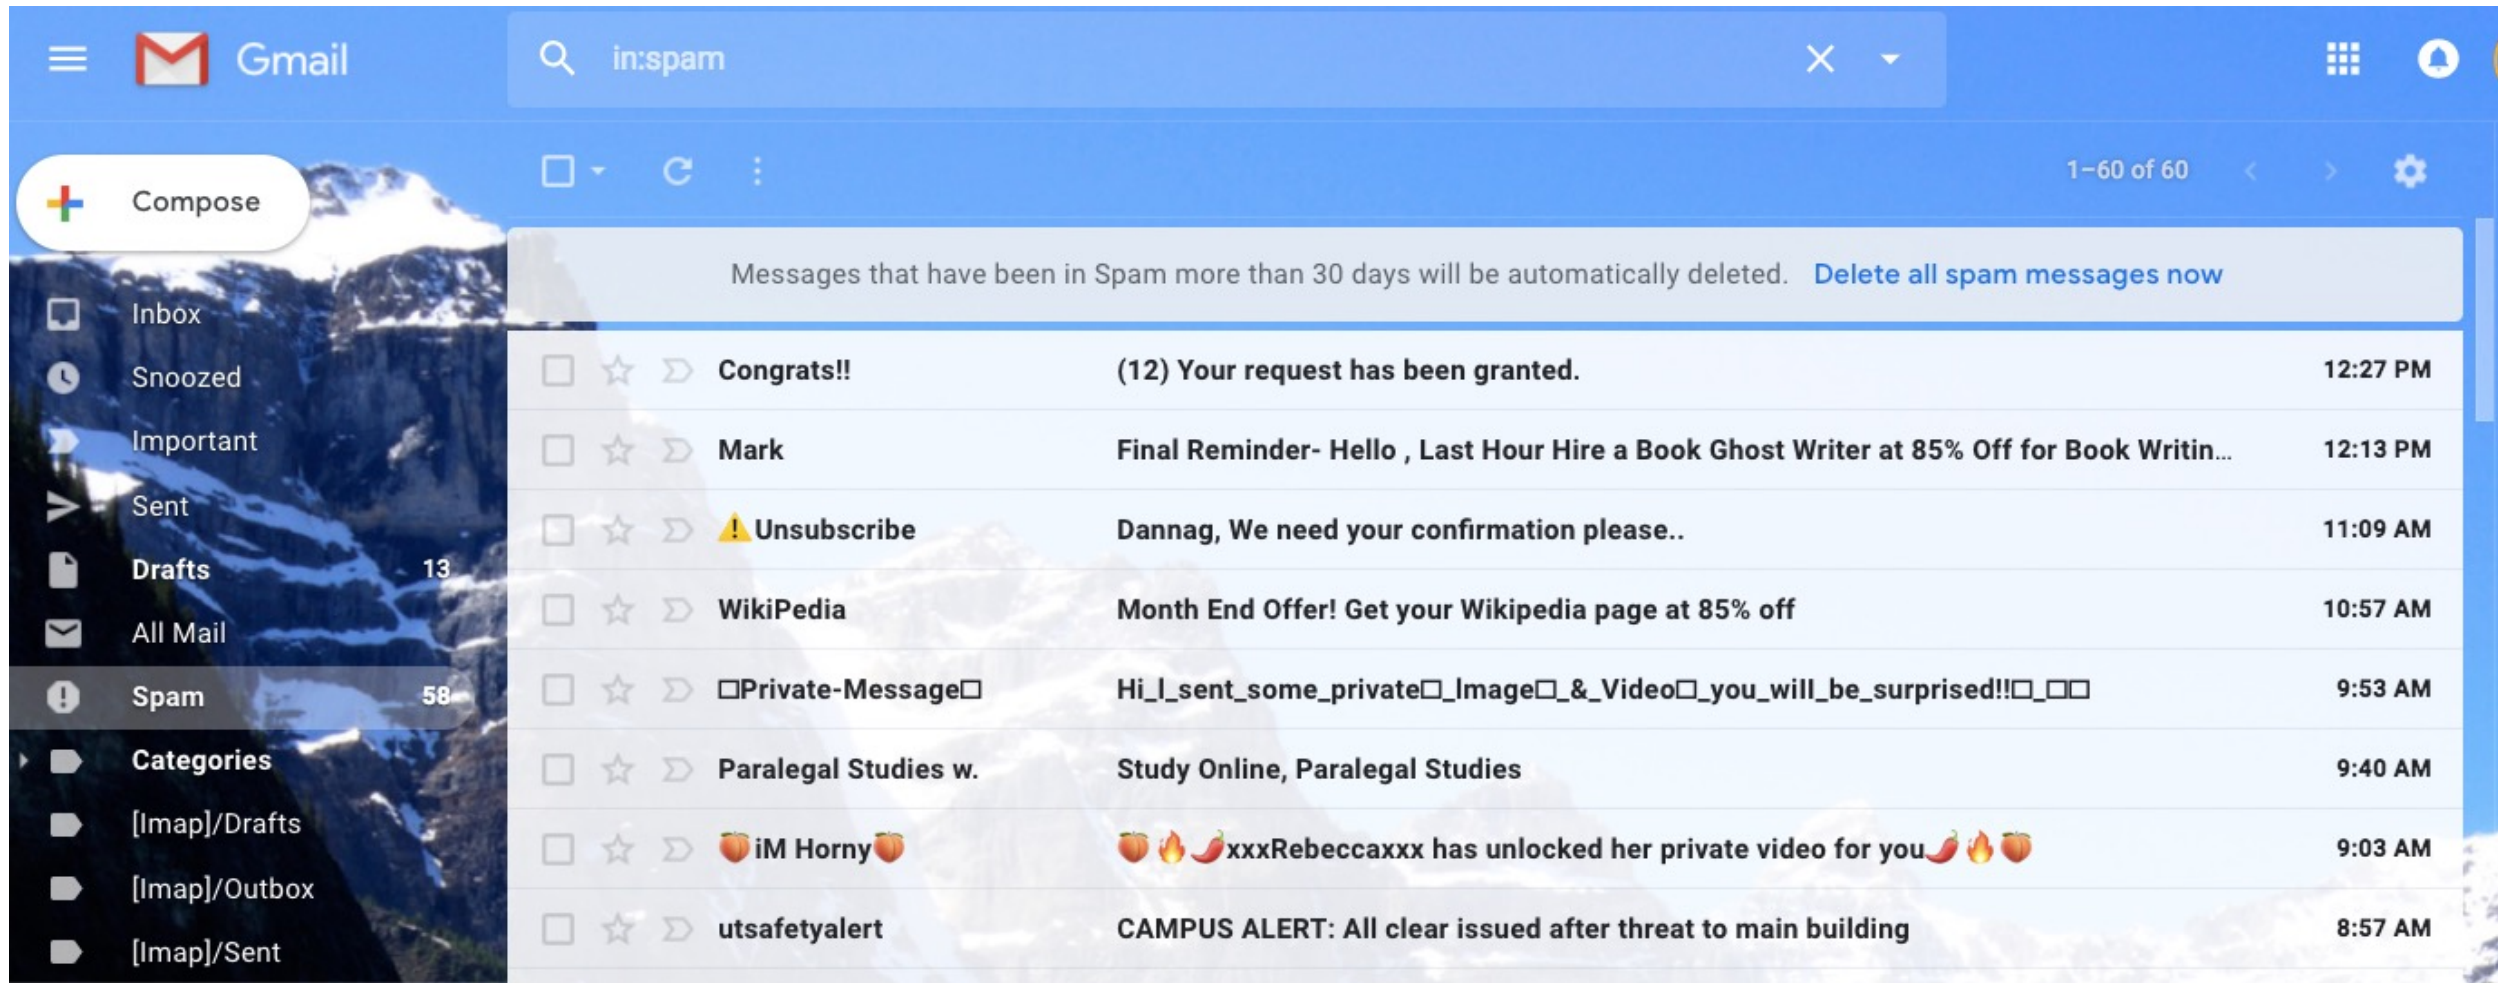
Today's Topics

- Binary classification applications
- Evaluating classification models
- Biological neurons: inspiration
- Artificial neuron: Perceptron

Today's Scope: Binary Classification

Distinguish 2 classes


Binary Classification: Spam Detection




The screenshot displays the Gmail interface with the search bar set to "in:spam". The left sidebar shows the "Spam" folder selected, containing 58 messages. The main area lists several spam messages with their subjects, senders, and times. A notification at the top of the message list states: "Messages that have been in Spam more than 30 days will be automatically deleted. [Delete all spam messages now](#)".

Message Icon	Star	Sender	Subject	Time	
<input type="checkbox"/>	☆	>	Congrats!!	(12) Your request has been granted.	12:27 PM
<input type="checkbox"/>	☆	>	Mark	Final Reminder- Hello , Last Hour Hire a Book Ghost Writer at 85% Off for Book Writin...	12:13 PM
<input type="checkbox"/>	☆	>	! Unsubscribe	Dannag, We need your confirmation please..	11:09 AM
<input type="checkbox"/>	☆	>	WikiPedia	Month End Offer! Get your Wikipedia page at 85% off	10:57 AM
<input type="checkbox"/>	☆	>	Private-Messag	Hi_I_sent_some_private_Image_&_Video_you_will_be_surprised!!_	9:53 AM
<input type="checkbox"/>	☆	>	Paralegal Studies w.	Study Online, Paralegal Studies	9:40 AM
<input type="checkbox"/>	☆	>	iM Horny	xxxRebeccaxxx has unlocked her private video for you	9:03 AM
<input type="checkbox"/>	☆	>	utsafetyalert	CAMPUS ALERT: All clear issued after threat to main building	8:57 AM

Binary Classification: Resume Pre-Screening



Book demo Free trial




ABOUT PLATFORM ROADMAP COIN TEAM WHITEPAPER SIGN IN BUY COIN


AI AND BLOCKCHAIN STAFFING & RECRUITMENT

- ◆ Decentralizing and simplifying the staffing industry
- ◆ Eliminate the intermediaries between job seeker, through an open ecosystem of hiring managers by using Blockchain, AI, and other technologies, that ultimately makes hiring more cost-effective

BUY COIN WHITE PAPER

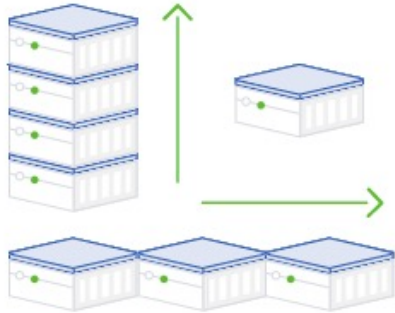


HOME PRODUCT ABOUT US BLOG REQUEST A DEMO



Automatically screen resumes

Trained with over 20 million diverse profiles, Skillate's AI algorithm helps to screen and shortlist resume with the click of a button. Seamlessly integrate with all external channels and ATS to source resume directly



Matching algorithm and candidate recommendation

Skillate's matching engine maps all the relevant profiles with the job requirements - be it skills, education or experience and recommends the best candidate

Binary Classification: Cancer Diagnosis

 [What we do](#) [About us](#) [News](#) [Careers](#) [Pathologists](#) [Partner with Us](#) [Partner Login](#)



Pathology Evolved.


Advanced learning toward faster,
more accurate diagnosis of disease.

Binary Classification: Cognitive Impairment Recognition by Apple App Usage




Image Credit: <https://www.techradar.com/news/the-10-best-phones-for-seniors>
https://www.technologyreview.com/f/615032/the-apps-you-use-on-your-phone-could-help-diagnose-your-cognitive-health/?utm_medium=tr_social&utm_campaign=site_visitor.unpaid.engagement&utm_source=Twitter#Echobox=1579899156

Binary Classification: Sentiment Analysis


What's the Tomatometer®?CriticsSIGN UPLOG INMOVIESTV SHOWSNew RT PODCASTNEWSSHOWTIMES

CRITIC REVIEWS FOR *MULAN*




Its cast, its attitude, its overall eagerness to please -- all benefits, one would think -- don't add up to a good movie. They add up to a blueprint of the movie this ought to be.

October 23, 2020 | Rating: 2.5/5 | [Full Review...](#)




K. Austin Collins
Rolling Stone
★ TOP CRITIC



While glorious to look at, the movie still feels slightly hollow. All the right pieces are there, but an emotional connection to the characters is lacking.

September 10, 2020 | Rating: 6.8/10 | [Full Review...](#)



Amy Amatangelo
Paste Magazine
★ TOP CRITIC

Binary Classification: Food Quality Control



Machine Learning: Using Algorithms to Sort Fruit

Demo: <https://www.youtube.com/watch?v=Bl3XzBWpZbY>

Can you think of other binary
classification applications?

Today's Topics

- Binary classification applications
- Evaluating classification models
- Biological neurons: inspiration
- Artificial neuron: Perceptron

Goal: Design Models that **Generalize Well** to New, Previously Unseen Examples

Example:



Label:

Hairy

Hairy

Not Hairy



Hairy




Goal: Design Models that **Generalize Well** to New, Previously Unseen Examples

1. Split data into a “**training set**” and “**test set**”

Training Data

Test Data

Example:



...



...

Label:

Hairy

Hairy

Not Hairy

...


Hairy

...

Goal: Design Models that **Generalize Well** to New, Previously Unseen Examples


2. Train model on “**training set**” to try to minimize prediction error on it

Training Data

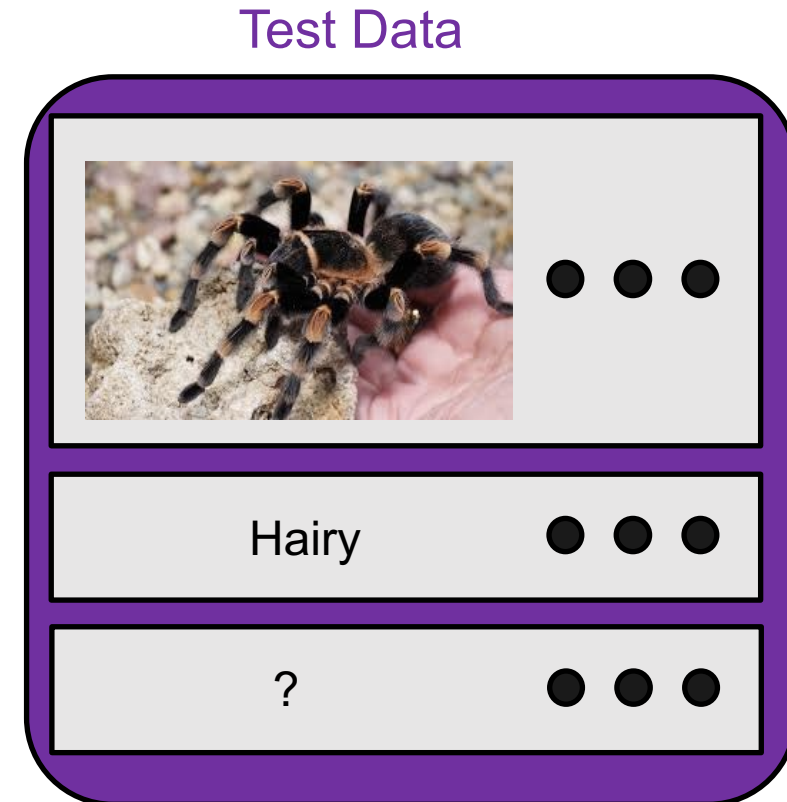


Goal: Design Models that **Generalize Well** to New, Previously Unseen Examples

3. Apply trained model on “**test set**” to measure generalization error




Example:



Goal: Design Models that **Generalize Well** to New, Previously Unseen Examples


3. Apply trained model on “**test set**” to measure generalization error


Prediction Model



Example:


Test Data



Goal: Design Models that **Generalize Well** to New, Previously Unseen Examples

3. Apply trained model on “**test set**” to measure generalization error


Prediction Model



Example:

Test Data

	...
Label:	Hairy ...
Predicted Label:	Not X hairy ...

Evaluation Methods: Confusion Matrix

		Actual	
		Spam	Not spam
Predicted	Spam	TP	FP
	Not spam	FN	TN

TP = true positive

TN = true negative

FP = false positive

FN = false negative

Evaluation Methods : Descriptive Statistics

Commonly-used statistical descriptions:

e.g.,

		Actual	
		Spam	Not spam
Predicted	Spam	50	10
	Not spam	15	100

- How many **actual spam** results are there? - 65
- How many **actual trusted** results are there? - 110
- How many **correctly classified instances**? - 150/175 ~ 86%
- How many **incorrectly classified instances**? - 25/175 ~ 14%

- What is the **precision**?
 - $50/(50+10) \sim 83\%$

$$\frac{TP}{TP + FP}$$

- What is the **recall**?
 - $50/(50+15) \sim 77\%$

$$\frac{TP}{TP + FN}$$

Group Discussion

- Which of these evaluation metrics would you use versus not use and why?
 - Accuracy (percentage of correctly classified examples)
 - Precision
 - Recall
- Scenario 1: Medical test for a rare disease affecting one in every million people.
- Scenario 2: Deciding which emails to flag as spam.

Today's Topics

- Binary classification applications
- Evaluating classification models
- **Biological neurons: inspiration**
- Artificial neuron: Perceptron

Inspiration: Animal's Computing Machinery

Neuron

- basic unit in the nervous system for receiving, processing, and transmitting information; e.g., messages such as...

“hot”



<https://www.clipart.email/clipart/dont-touch-hot-stove-clipart-73647.html>

“loud”



<https://kisselpaso.com/if-the-sun-city-music-fest-gets-too-loud-there-is-a-phone-number-you-can-call-to-complain/>

“spicy”




https://www.babycenter.com/404_when-can-my-baby-eat-spicy-foods_1368539.bc

Inspiration: Animal's Computing Machinery



<https://en.wikipedia.org/wiki/Nematode#/media/File:CelestansGoldsteinLabUNC.jpg>

Nematode worm: 302 neurons




© 2006 Encyclopædia Britannica, Inc.

<https://www.britannica.com/science/human-nervous-system>


Human: ~100,000,000,000 neurons

Inspiration: Animal's Computing Machinery



Demo (0-1:14): <https://www.youtube.com/watch?v=oa6rvUJlg7o>

Inspiration: Basic Understanding of Neurons



- When the input signals exceed a certain threshold within a short period of time, a neuron “fires”
- Neuron “firing” is an “all-or-none” process, where either a signal is sent or nothing happens

Image Source: <https://becominghuman.ai/introduction-to-neural-networks-bd042ebf2653>

Sidenote: It Remains An Open Research Problem
to Understand How Individual Neurons Work

Today's Topics

- Binary classification applications
- Evaluating classification models
- Biological neurons: inspiration
- Artificial neuron: Perceptron

Historical Context: Artificial Neurons


First mathematical
model of neuron



Warren McCulloch
(Neurophysiologist)

http://web.csulb.edu/~cwallis/artificialn/warren_mcculloch.html



Emerges from
interdisciplinary
collaboration



Walter Pitts
(Mathematician)

https://en.wikipedia.org/wiki/Walter_Pitts

Artificial Neuron: McCulloch-Pitts Neuron

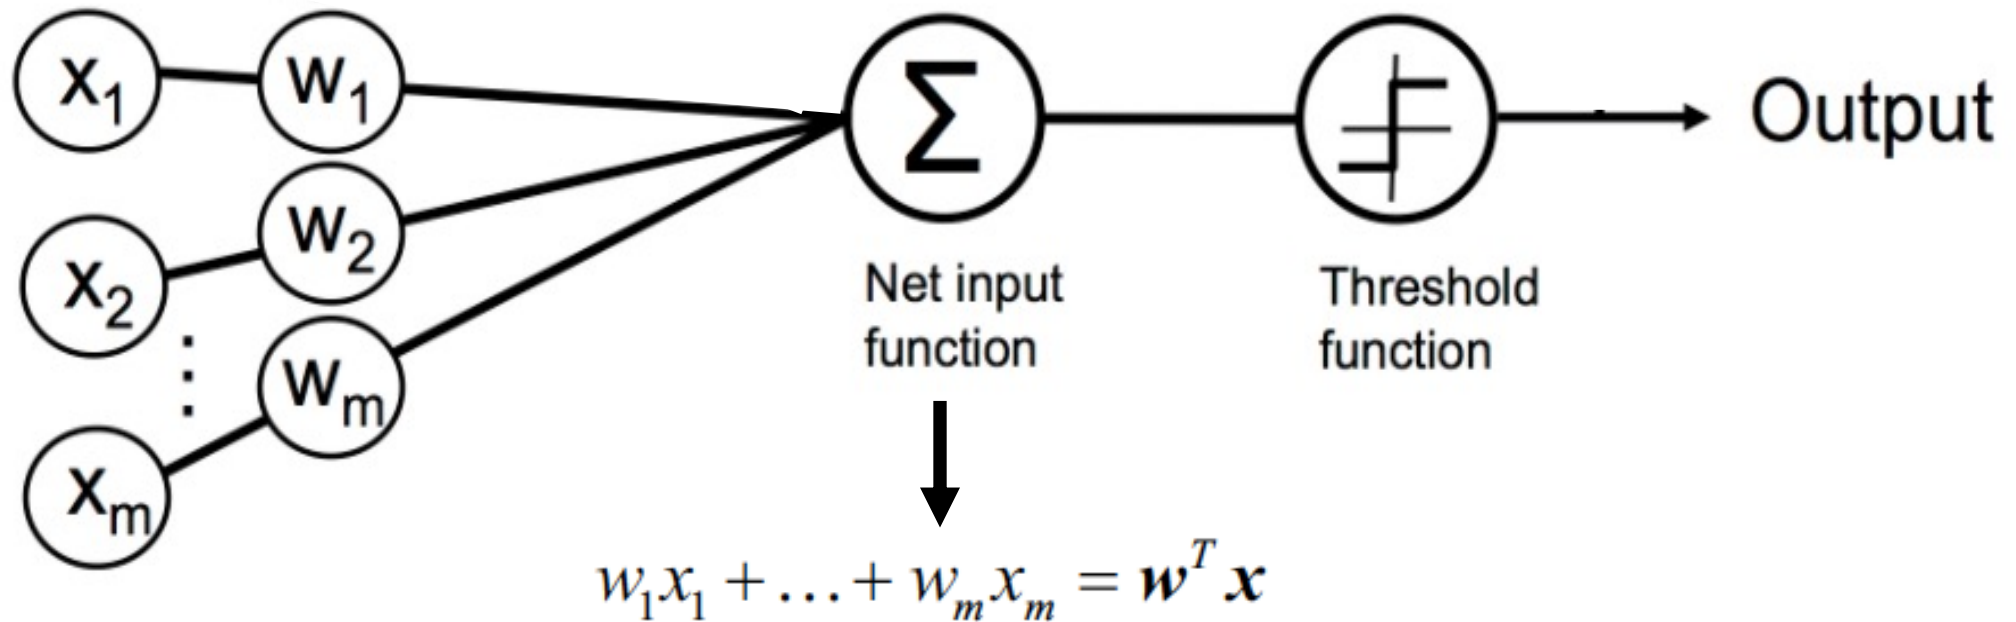


Python Machine Learning; Raschka & Mirjalili

Image Source: <https://becominghuman.ai/introduction-to-neural-networks-bd042ebf2653>

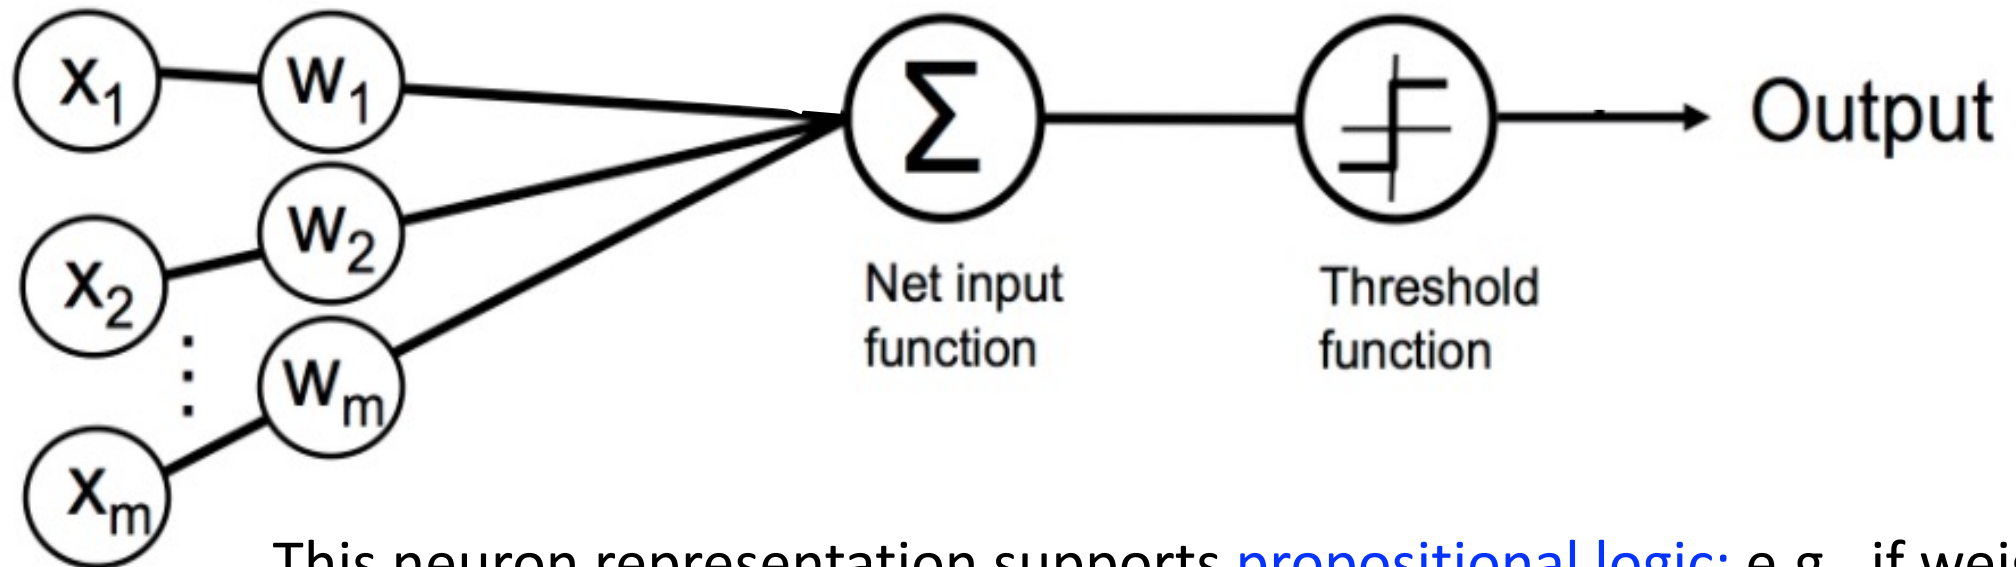
Artificial Neuron: McCulloch-Pitts Neuron

- inputs (x) and weights (w) can be 0 or 1
- weights (w) and threshold values are fixed
- outputs 1 or 0 (mimics neurons by “firing” only when aggregate value exceeds threshold)



Artificial Neuron: McCulloch-Pitts Neuron


- inputs (x) and weights (w) can be 0 or 1
- weights (w) and threshold values are fixed
- outputs 1 or 0 (mimics neurons by “firing” only when aggregate value exceeds threshold)



This neuron representation supports [propositional logic](#); e.g., if weights equal 1 and there are 3 inputs, how is the [AND function](#) achieved?

Artificial Neuron: McCulloch-Pitts Neuron

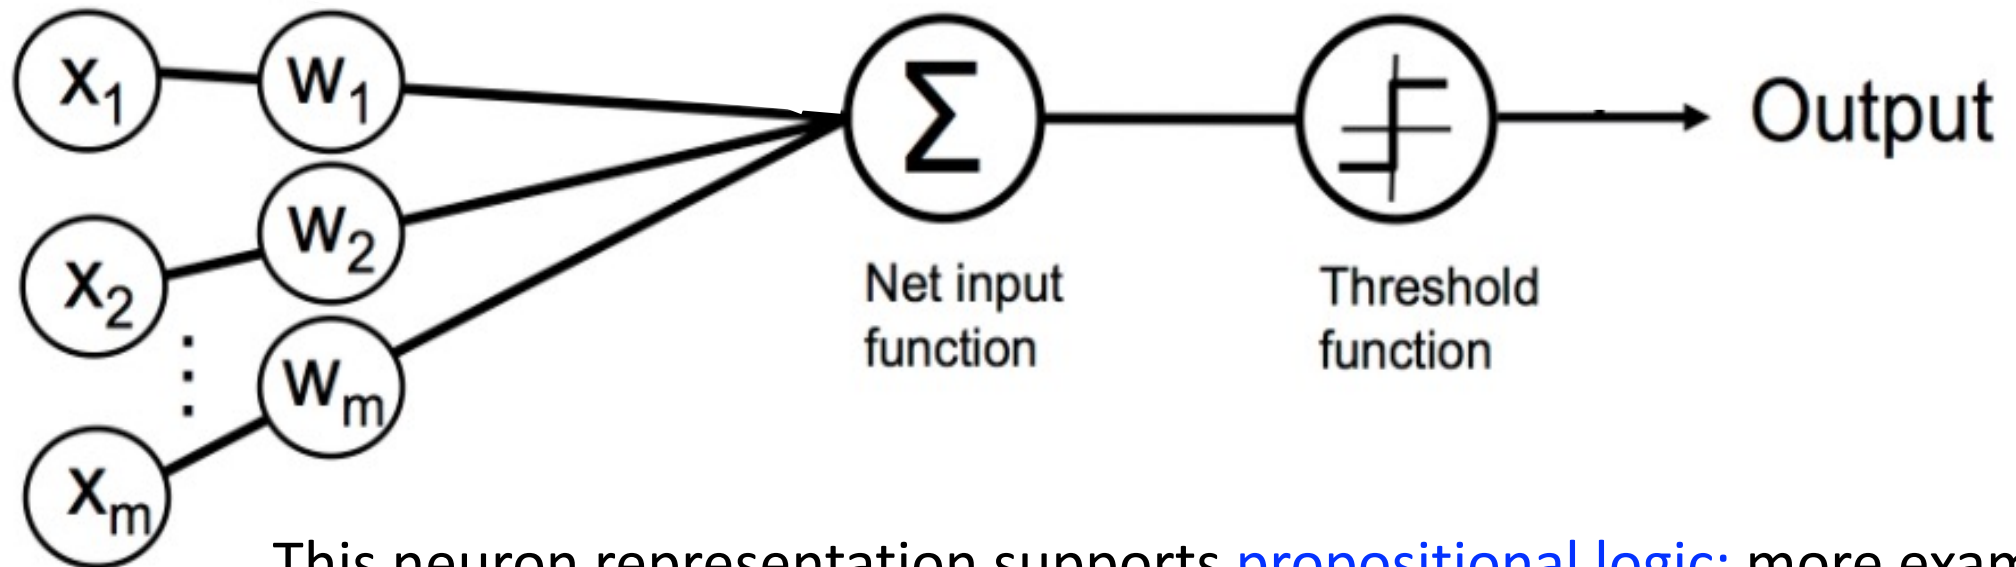
- inputs (x) and weights (w) can be 0 or 1
- weights (w) and threshold values are fixed
- outputs 1 or 0 (mimics neurons by “firing” only when aggregate value exceeds threshold)



This neuron representation supports [propositional logic](#); e.g., if weights equal 1 and there are 3 inputs, how is the [OR function](#) achieved?

Artificial Neuron: McCulloch-Pitts Neuron

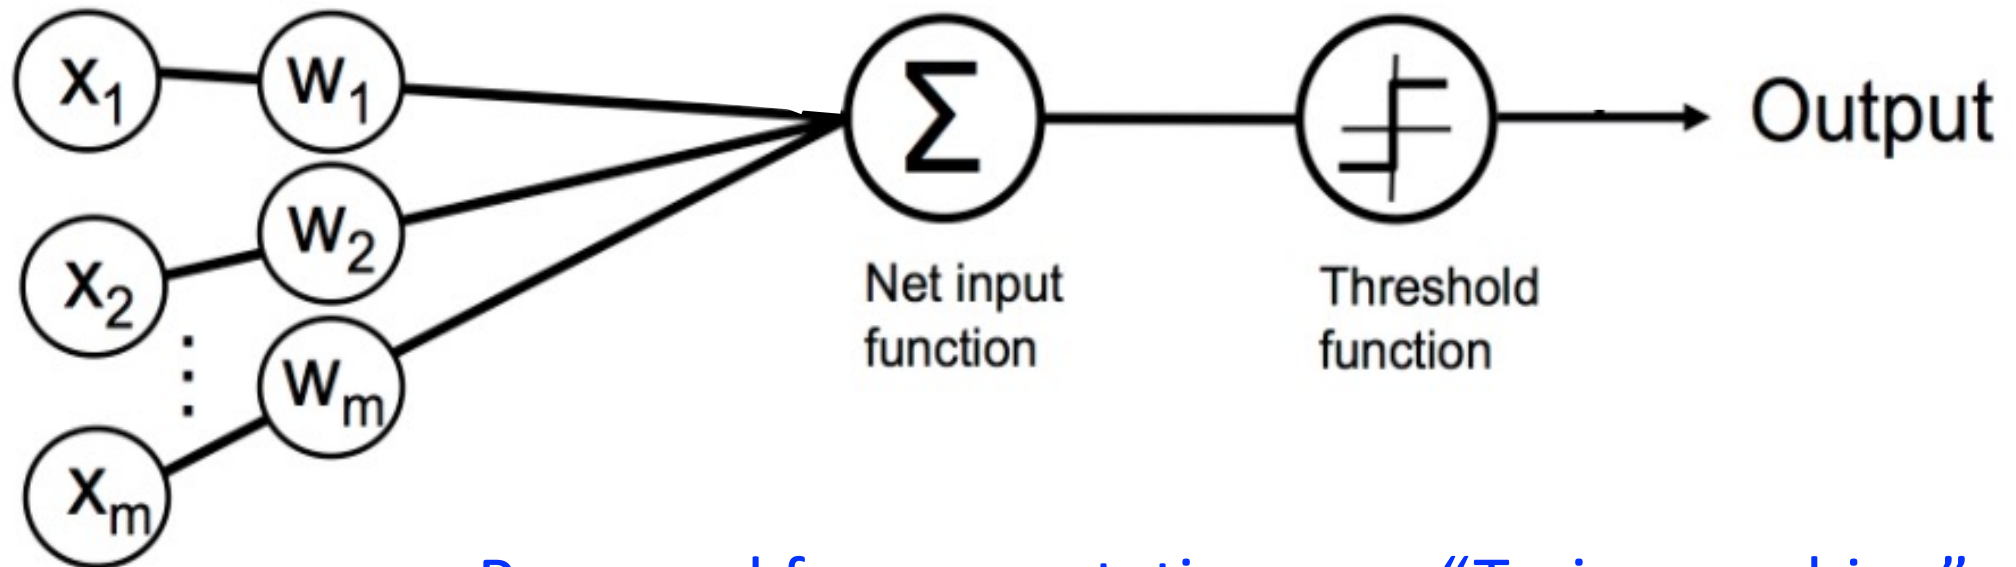
- inputs (x) and weights (w) can be 0 or 1
- weights (w) and threshold values are fixed
- outputs 1 or 0 (mimics neurons by “firing” only when aggregate value exceeds threshold)



This neuron representation supports [propositional logic](https://home.csulb.edu/~cwallis/artificialn/History.htm); more examples found at <https://home.csulb.edu/~cwallis/artificialn/History.htm>


Artificial Neuron: McCulloch-Pitts Neuron

- inputs (x) and weights (w) can be 0 or 1
- weights (w) and threshold values are fixed
- outputs 1 or 0 (mimics neurons by “firing” only when aggregate value exceeds threshold)



Proposed for computation on a “Turing machine”

Historical Context: Artificial Neurons



Perceptron: Innovator and Vision



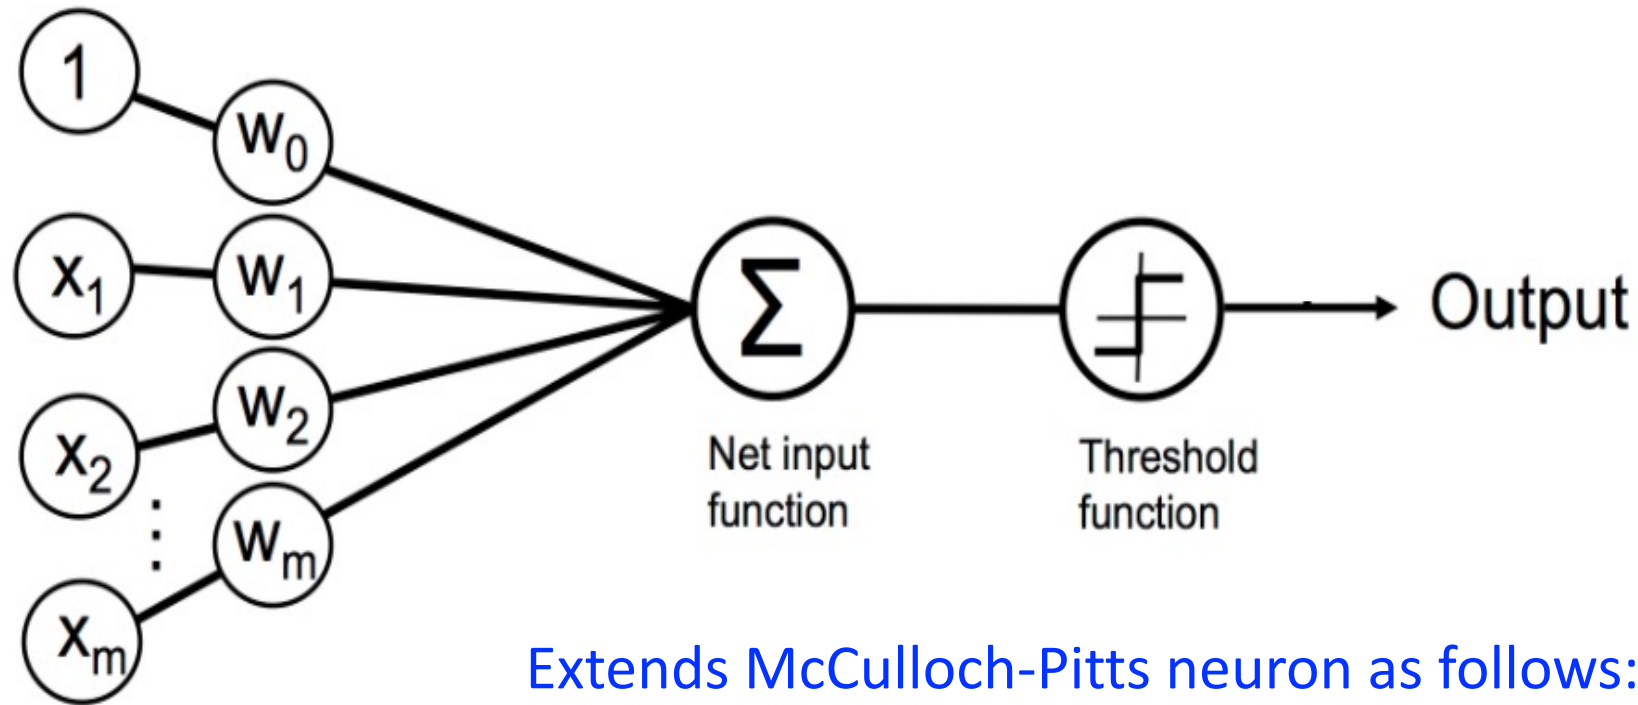
Frank Rosenblatt
(Psychologist)

“[The perceptron is] the embryo of an electronic computer that [the Navy] expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence.... [It] is expected to be finished in about a year at a cost of \$100,000.”

1958 New York Times article: <https://www.nytimes.com/1958/07/08/archives/new-navy-device-learns-by-doing-psychologist-shows-embryo-of.html>

https://en.wikipedia.org/wiki/Frank_Rosenblatt

Perceptron: Architecture (Linear Threshold Unit)



Extends McCulloch-Pitts neuron as follows:

- inputs and weights can be any value
- weights (W) are learned

Perceptron: Architecture (Linear Threshold Unit)

- Function deciding output value (“fire” or not):

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ -1 & \text{otherwise} \end{cases}$$

- Rewriting function:

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

*** Note:** Kamath textbook offers two common conventions for Perceptrons of using two possible output values of $\{-1, 1\}$ and $\{0, 1\}$, in Chapters 2.5 and 4. The output choice dictates whether the threshold should be set to 0.5 or 0.


- Where:

$$z = w_0 x_0 + w_1 x_1 + \dots + w_m x_m = \mathbf{w}^T \mathbf{x}$$


Bias $-\theta$ 1

Perceptron: Architecture (Linear Threshold Unit)

Graphical representation:



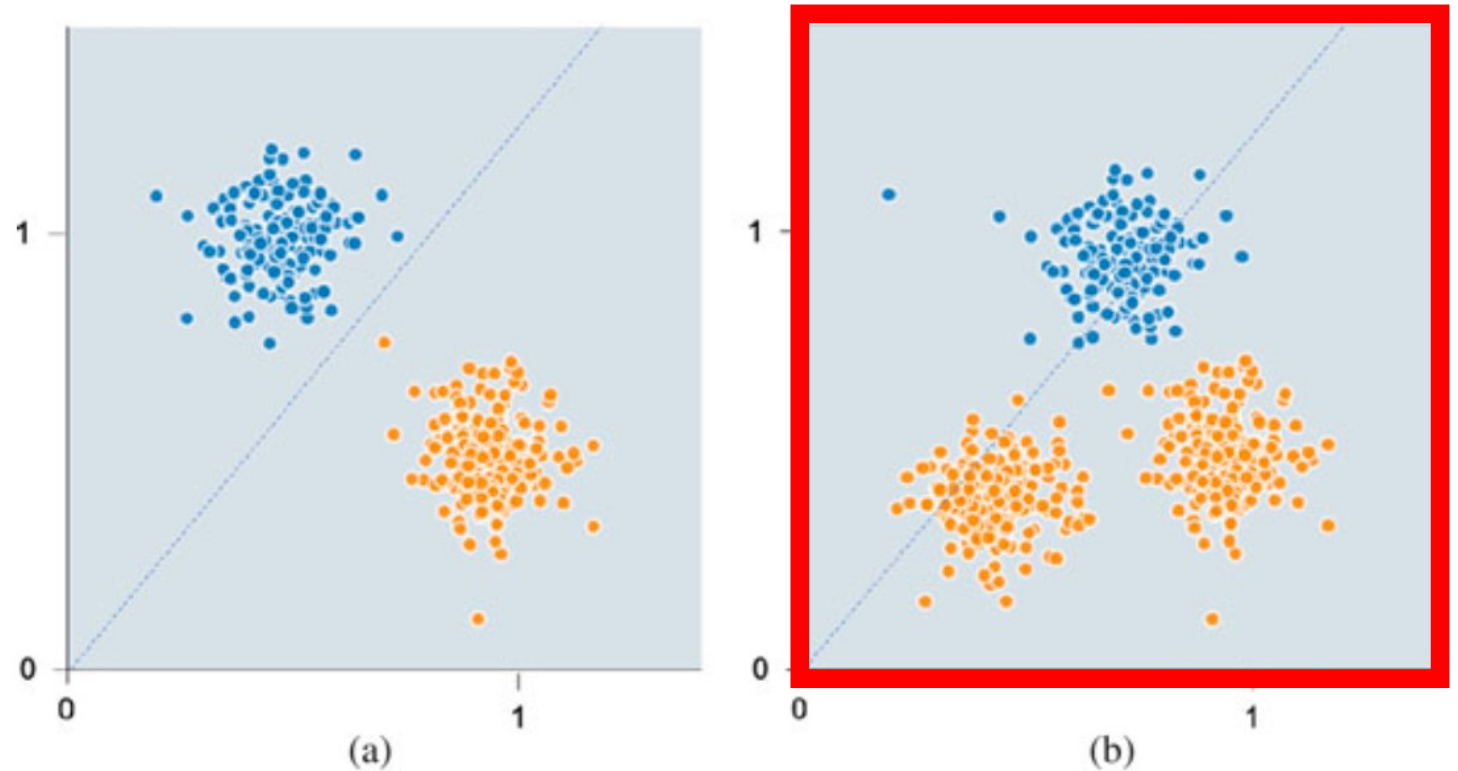
Perceptron: Architecture (Linear Threshold Unit)



What is the motivation for weights? e.g.,
for predicting if you will like a movie?

Perceptron: Architecture (Linear Threshold Unit)

Motivation for **bias**: without it, model must go through origin



$$z = w_0 x_0 + w_1 x_1 + \dots + w_m x_m = \mathbf{w}^T \mathbf{x}$$


Bias

$-\theta$

1

Perceptron: Architecture (Linear Threshold Unit)

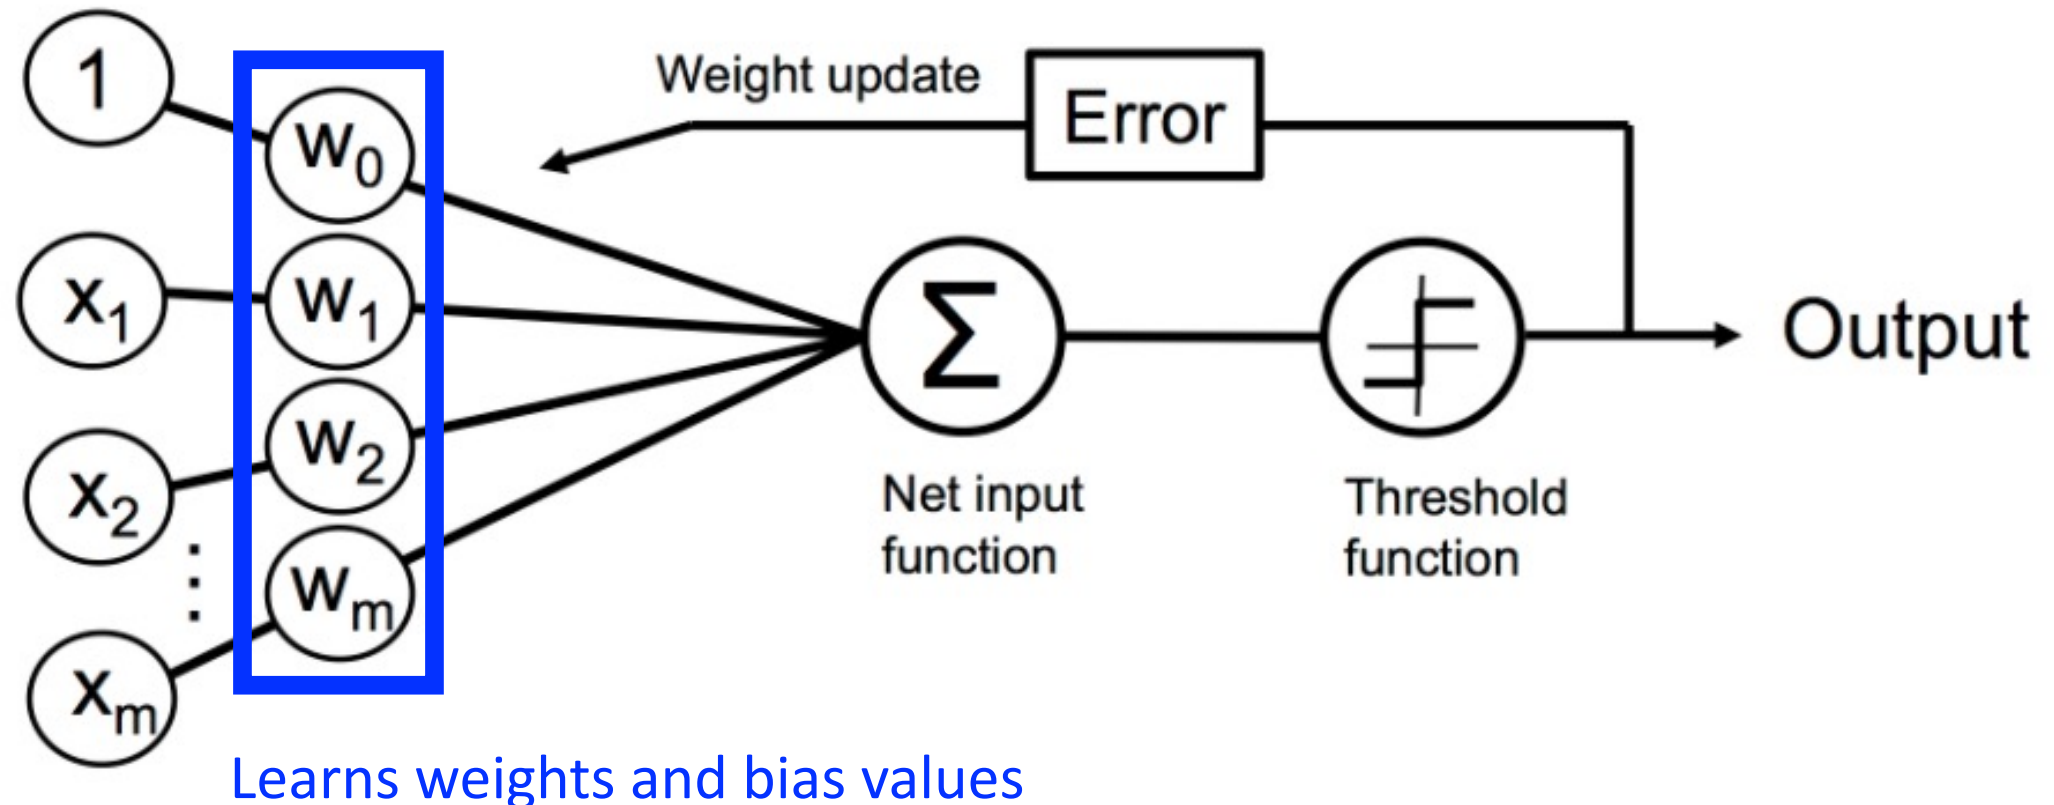
Motivation for **bias**: with it, model does not have to go through origin



$$z = w_0 x_0 + w_1 x_1 + \dots + w_m x_m = \mathbf{w}^T \mathbf{x}$$


Bias $-\theta$ 1

Perceptron: Learning Algorithm




Perceptron: Learning Algorithm

Process: iteratively update boundary with observation of each additional example:



Perceptron: Learning Algorithm

Process: iteratively update boundary with observation of each additional example:



Perceptron: Learning Algorithm

1. Initialize weights/bias to 0 or small random numbers
2. For each training sample (i.e., i) :

1. Compute predicted value (i.e., $\{-1, 1\}$): $\sum_{j=0}^m \mathbf{x}_j \mathbf{w}_j = \mathbf{w}^T \mathbf{x}$

2. Update parameters based on prediction success: $\mathbf{w}_j := \mathbf{w}_j + \Delta \mathbf{w}_j$

$$\Delta \mathbf{w}_j = \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$$

Learning Rate
(set a priori and
held constant)

True Class Label

Predicted Class Label

Perceptron: Learning Algorithm

1. Initialize weights/bias to 0 or small random numbers
2. For each training sample (i.e., i) :

1. Compute predicted value (i.e., $\{-1, 1\}$): $\sum_{j=0}^m \mathbf{x}_j \mathbf{w}_j = \mathbf{w}^T \mathbf{x}$

2. Update parameters based on prediction success: $\mathbf{w}_j := \mathbf{w}_j + \Delta \mathbf{w}_j$.

$$\Delta \mathbf{w}_j = \eta (\text{target}^{(i)} - \text{output}^{(i)}) \mathbf{x}_j^{(i)}$$

What happens to the weights when the model predicts the **correct** class label?

- no weight update since result is 0

Perceptron: Learning Algorithm

1. Initialize weights/bias to 0 or small random numbers
2. For each training sample (i.e., i) :

1. Compute predicted value (i.e., $\{-1, 1\}$): $\sum_{j=0}^m \mathbf{x}_j \mathbf{w}_j = \mathbf{w}^T \mathbf{x}$

2. Update parameters based on prediction success: $\mathbf{w}_j := \mathbf{w}_j + \Delta \mathbf{w}_j$

$$\Delta \mathbf{w}_j = \eta (\text{target}^{(i)} - \text{output}^{(i)}) \mathbf{x}_j^{(i)}$$


What happens to the weights when the model predicts the **wrong** class label?

- weights change since result is “2” or “-2”

Perceptron: Example

- True Model: Y is 1 if at least two of the three inputs are equal to 1.


X_1	X_2	X_3	Y
1	0	0	?
1	0	1	
1	1	0	
1	1	1	
0	0	1	
0	1	0	
0	1	1	
0	0	0	



Perceptron: Example

- True Model: Y is 1 if at least two of the three inputs are equal to 1.


X_1	X_2	X_3	Y
1	0	0	-1
1	0	1	?
1	1	0	
1	1	1	
0	0	1	
0	1	0	
0	1	1	
0	0	0	



Perceptron: Example

- True Model: Y is 1 if at least two of the three inputs are equal to 1.


X_1	X_2	X_3	Y
1	0	0	-1
1	0	1	1
1	1	0	?
1	1	1	
0	0	1	
0	1	0	
0	1	1	
0	0	0	



Perceptron: Example

- True Model: Y is 1 if at least two of the three inputs are equal to 1.

X_1	X_2	X_3	Y
1	0	0	-1
1	0	1	1
1	1	0	1
1	1	1	1
0	0	1	-1
0	1	0	-1
0	1	1	1
0	0	0	-1



Perceptron: Example (Training with 1st Sample)

- Compute predicted value: $\sum_{j=0}^m x_j w_j = \mathbf{w}^T \mathbf{x}$; $\phi(\mathbf{w}^T \mathbf{x}) = \begin{cases} 1 & \text{if } \phi(\mathbf{w}^T \mathbf{x}) \geq 0 \\ -1 & \text{otherwise} \end{cases}$

X_1	X_2	X_3	Y	Predicted	w_0	w_1	w_2	w_3
1	0	0	-1	?	0	0	0	0

Perceptron: Example (Training with 1st Sample)

- Update params: $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y	Predicted
1	0	0	-1	1

w_0	w_1	w_2	w_3
0	0	0	0
?	?	?	?

$$\Delta w_0 = \eta (\text{target}^{(i)} - \text{output}^{(i)})$$

$$\Delta w_1 = \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_1^{(i)}$$

$$\Delta w_2 = \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_2^{(i)}$$

$$\Delta w_3 = \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_3^{(i)}$$

Perceptron: Example (Training with 1st Sample)

- Update params: $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y
1	0	0	-1

Predicted
1

w_0	w_1	w_2	w_3
0	0	0	0
?	?	?	?

$$\Delta w_0 = 0.1(-1-1)*1 = -0.2$$

$$\Delta w_1 = 0.1(-1-1)*1 = -0.2$$

$$\Delta w_2 = 0.1(-1-1)*0 = 0$$

$$\Delta w_3 = 0.1(-1-1)*0 = 0$$

Updates make weights more negative so that the model is more likely to classify the sample as -1 next time

Perceptron: Example (Training with 2nd Sample)

- Compute output value: $\sum_{j=0}^m x_j w_j = \mathbf{w}^T \mathbf{x}$; $\phi(\mathbf{w}^T \mathbf{x}) = \begin{cases} 1 & \text{if } \phi(\mathbf{w}^T \mathbf{x}) \geq 0 \\ -1 & \text{otherwise} \end{cases}$

X_1	X_2	X_3	Y	Predicted	w_0	w_1	w_2	w_3
1	0	0	-1	1	0	0	0	0
1	0	1	1	?	-0.2	-0.2	0	0

Perceptron: Example (Training with 2nd Sample)

- Update params: $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y	Predicted
1	0	0	-1	1
1	0	1	1	-1

w_0	w_1	w_2	w_3
0	0	0	0
-0.2	-0.2	0	0
?	?	?	?

$$\Delta w_0 = \eta (\text{target}^{(i)} - \text{output}^{(i)})$$

$$\Delta w_1 = \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_1^{(i)}$$

$$\Delta w_2 = \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_2^{(i)}$$

$$\Delta w_3 = \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_3^{(i)}$$

Perceptron: Example (Training with 2nd Sample)

- Update params: $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y
1	0	0	-1
1	0	1	1

Predicted
1
-1

w_0	w_1	w_2	w_3
0	0	0	0
-0.2	-0.2	0	0
?	?	?	?

$$\Delta w_0 = 0.1(1 - -1) * 1 = 0.2$$

$$\Delta w_1 = 0.1(1 - -1) * 1 = 0.2$$

$$\Delta w_2 = 0.1(1 - -1) * 0 = 0$$

$$\Delta w_3 = 0.1(1 - -1) * 1 = 0.2$$

Updates make weights more positive so that the model is more likely to classify the sample as 1 next time

Perceptron: Example (Training with 2nd Sample)

- Update params: $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y	Predicted
1	0	0	-1	1
1	0	1	1	-1

w_0	w_1	w_2	w_3
0	0	0	0
-0.2	-0.2	0	0
0	0	0	0.2

$$\Delta w_0 = 0.1(1 - -1) * 1 = 0.2$$

$$\Delta w_1 = 0.1(1 - -1) * 1 = 0.2$$

$$\Delta w_2 = 0.1(1 - -1) * 0 = 0$$

$$\Delta w_3 = 0.1(1 - -1) * 1 = 0.2$$

What is the influence of the learning rate? i.e., what would happen if the value was larger/smaller?

Perceptron: Example – One Epoch (Training with All Samples)

- $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y
1	0	0	-1
1	0	1	1
1	1	0	1
1	1	1	1
0	0	1	-1
0	1	0	-1
0	1	1	1
0	0	0	-1

	w_0	w_1	w_2	w_3
0	0	0	0	0
1	-0.2	-0.2	0	0
2	0	0	0	0.2
3	0	0	0	0.2
4	0	0	0	0.2
5	-0.2	0	0	0
6	-0.2	0	0	0
7	0	0	0.2	0.2
8	-0.2	0	0.2	0.2

Perceptron: Example – Six Epochs

- $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y
1	0	0	-1
1	0	1	1
1	1	0	1
1	1	1	1
0	0	1	-1
0	1	0	-1
0	1	1	1
0	0	0	-1

Epoch	w_0	w_1	w_2	w_3
0	0	0	0	0

Perceptron: Example – Six Epochs

- $w_j = w_j + \eta (\text{target}^{(i)} - \text{output}^{(i)}) x_j^{(i)}$; learning rate = 0.1

X_1	X_2	X_3	Y
1	0	0	-1
1	0	1	1
1	1	0	1
1	1	1	1
0	0	1	-1
0	1	0	-1
0	1	1	1
0	0	0	-1

	w_0	w_1	w_2	w_3
0	0	0	0	0
1	-0.2	-0.2	0	0
2	0	0	0	0.2
3	0	0	0	0.2
4	0	0	0	0.2
5	-0.2	0	0	0
6	-0.2	0	0	0
7	0	0	0.2	0.2
8	-0.2	0	0.2	0.2

Epoch	w_0	w_1	w_2	w_3
0	0	0	0	0
1	-0.2	0	0.2	0.2
2	-0.2	0	0.4	0.2
3	-0.4	0	0.4	0.2
4	-0.4	0.2	0.4	0.4
5	-0.6	0.2	0.4	0.2
6	-0.6	0.4	0.4	0.2

Perceptron: Learning Algorithm Choices

- Learning rate
- Number of epochs (passes over the dataset)

Today's Topics

- Binary classification applications
- Evaluating classification models
- Biological neurons: inspiration
- Artificial neuron: Perceptron



The End

Credits

- Image of Boulder: <http://boulderrunning.com/where2run/five-trails-for-hill-running-and-mountain-training/>
- Stick person figure: <https://drawception.com/game/AsPNcppPND/draw-yourself-blindfolded-pio/>
- Figure: <https://www.quora.com/What-is-meant-by-gradient-descent-in-laymen-terms>
- Figure and great reference: https://beamandrew.github.io/deeplearning/2017/02/23/deep_learning_101_part1.html