

# Today's Topics

- Introduction to natural language processing
- Text representation
- Neural word embeddings
- Programming tutorial

# Today's Topics

- Introduction to natural language processing
- Text representation
- Neural word embeddings
- Programming tutorial

# NLP: Computers that Can Understand (and So Also Communicate in) Human Language

**Most recent customer reviews**

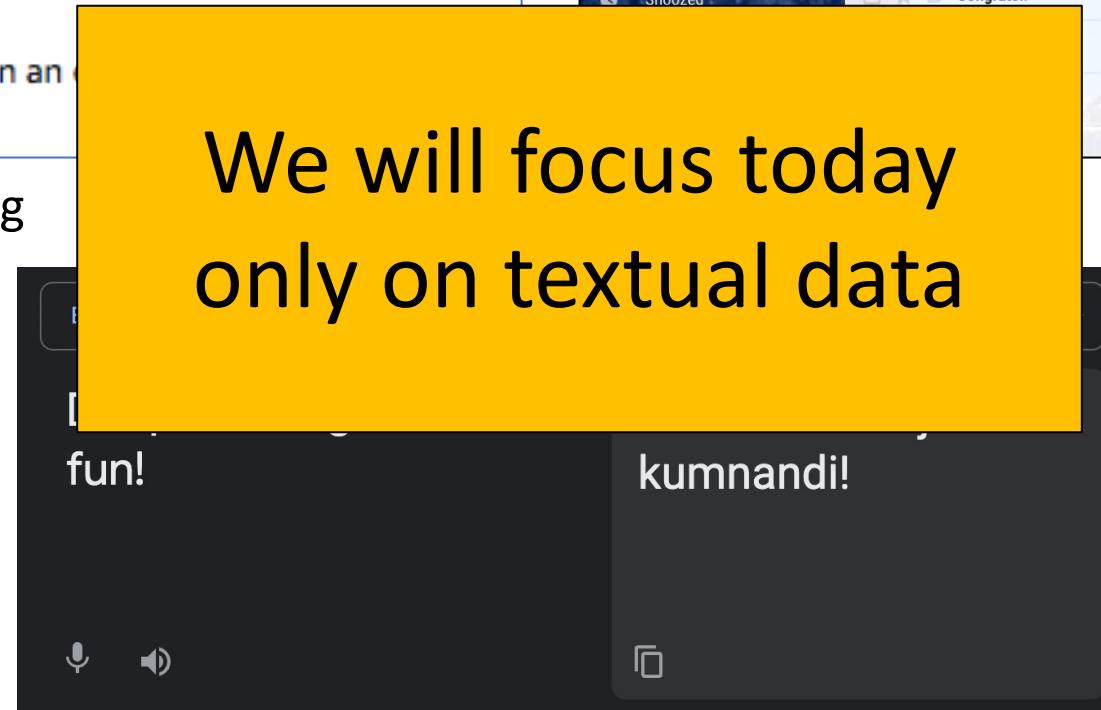
Vivek Chopra

★★★★★ Quality product, easy to setup

Fantastic product.  
Wanted to enable voice command on an...

Published 25 minutes ago

Opinion Mining



Language Translation

Compose

Inbox

Snoozed

Congrats!!

Messages that have been in Spam more than 30 days will be automatically deleted. Delete all spam messages now

(12) Your request has been granted. 12:27 PM

Final Reminder- Hello , Last Hour Hire a Book Ghost Writer at 85% Off for Book Writin... 12:13 PM

Dannag, We need your confirmation please.. 11:09 AM

Month End Offer! Get your Wikipedia page at 85% off 10:57 AM

Spam Detection

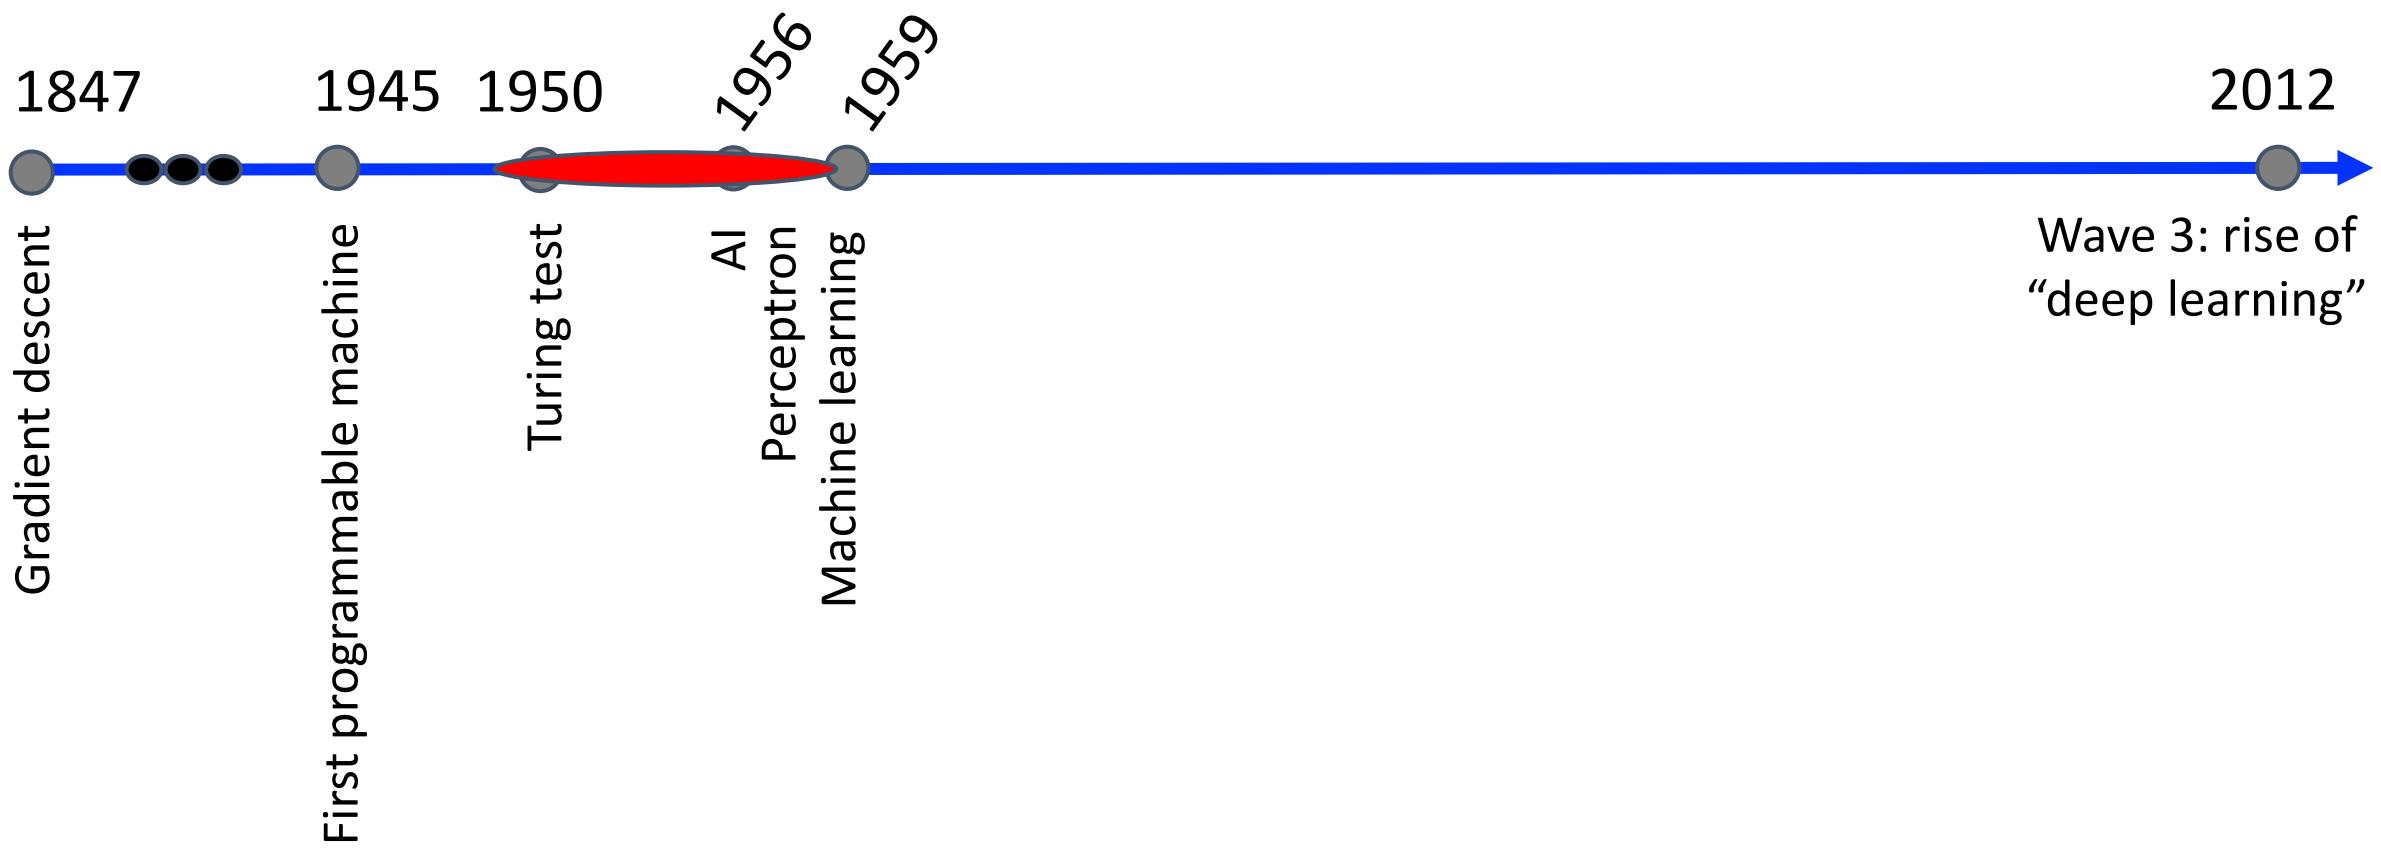
# Why Discuss NLP With RNNs?

- RNNs have a strong track record for NLP problems
- Text data's representation (i.e., sequential data) is a natural match for RNNs

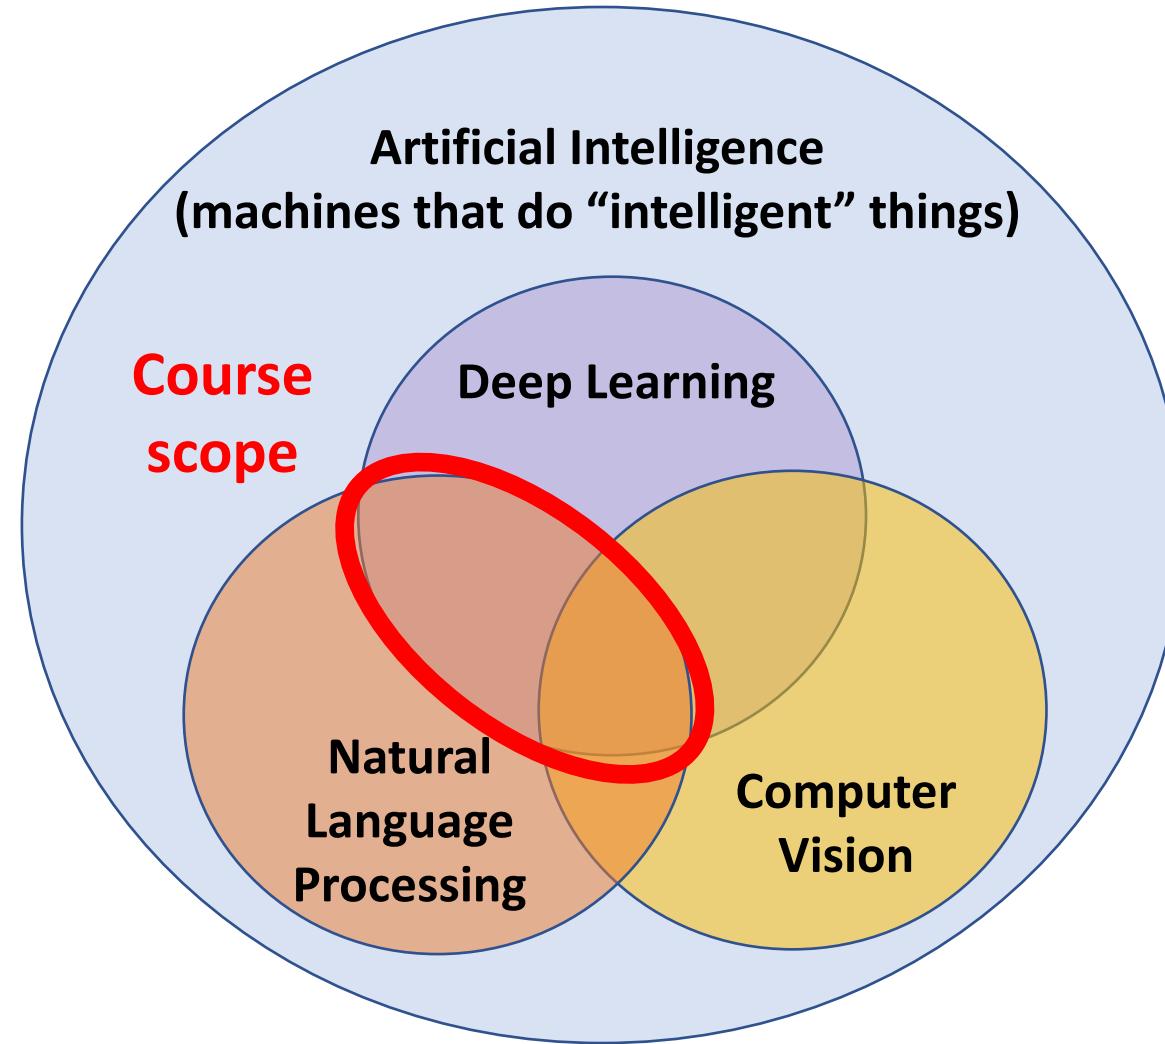


# Historical Context: Origins of NLP

Research community emerged mostly on  
the problem of translating languages



# NLP in Context

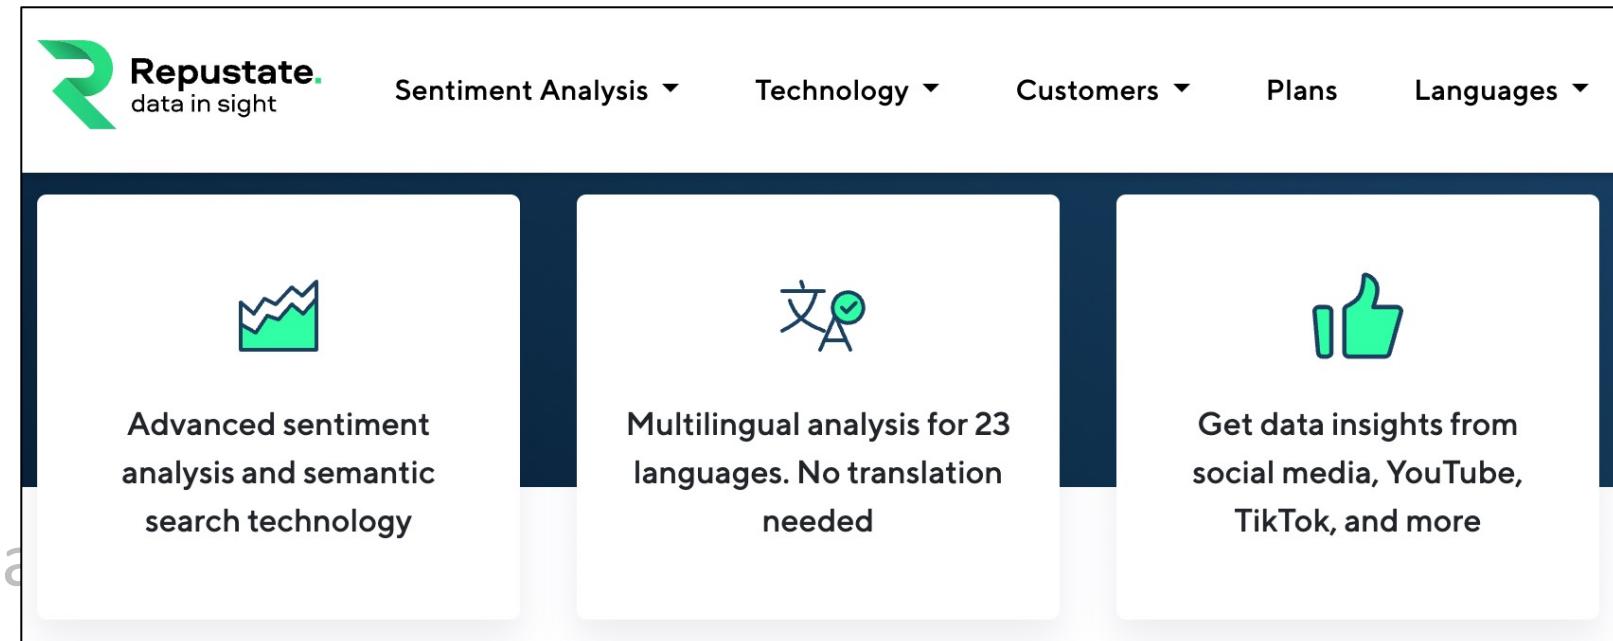


# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...

# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...



# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...

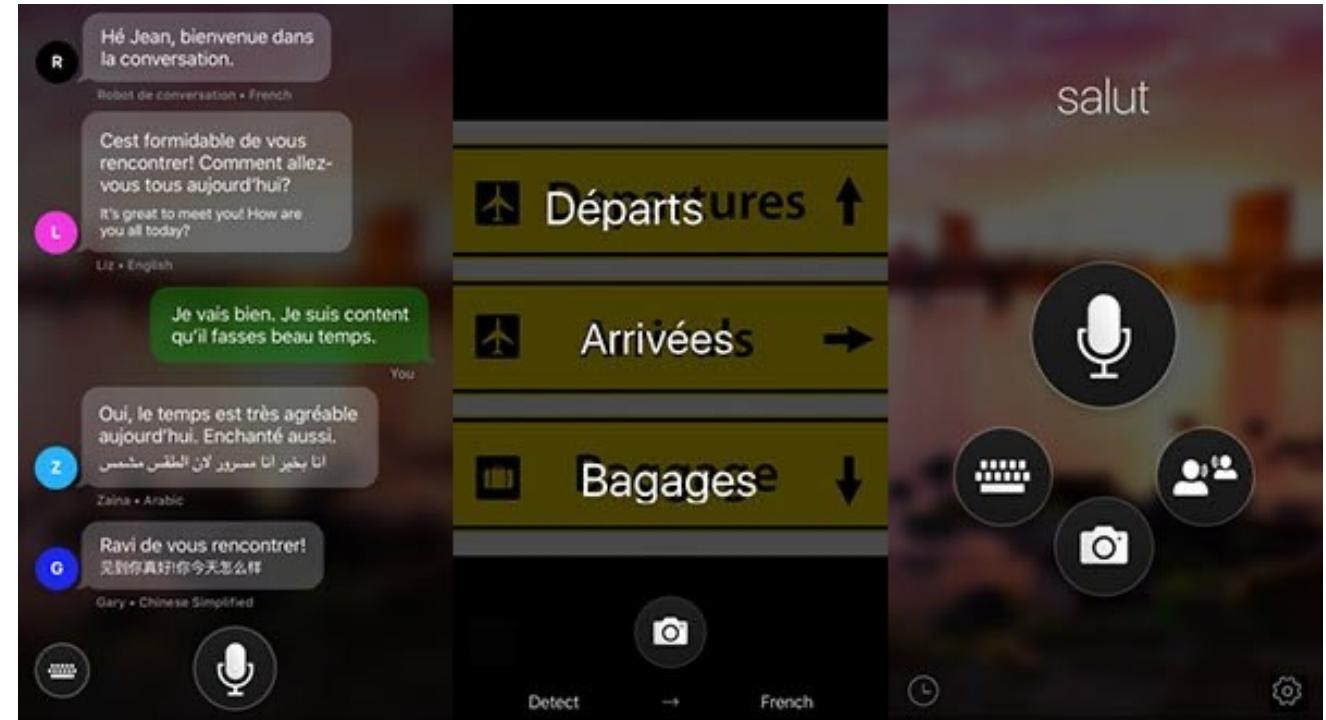
The screenshot shows the Fakespot homepage with a search bar at the top. Below it, a main headline reads "Hate returning stuff to Amazon? Get Fakespot". A subtext explains that with Fakespot, users can get the best products from the best sellers at the best price. A prominent blue button says "Add Fakespot — It's free". Below this, there are links to partner websites: eBay, Amazon, Best Buy, Sephora, and Walmart. To the right, there are three search results for "Apple Airpods Pro":

- amazon**: Apple Airpods Pro, Sold by SalesKingBest9393, with a "Seller Warning" badge.
- eBay**: Apple Airpods Pro, Sold by TechSeller33, with a "Seller Approved" badge.
- Walmart**: Apple Airpods Pro, Sold by WalmartSeller95, with a "Seller Approved" badge.

# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...

e.g., Microsoft translator

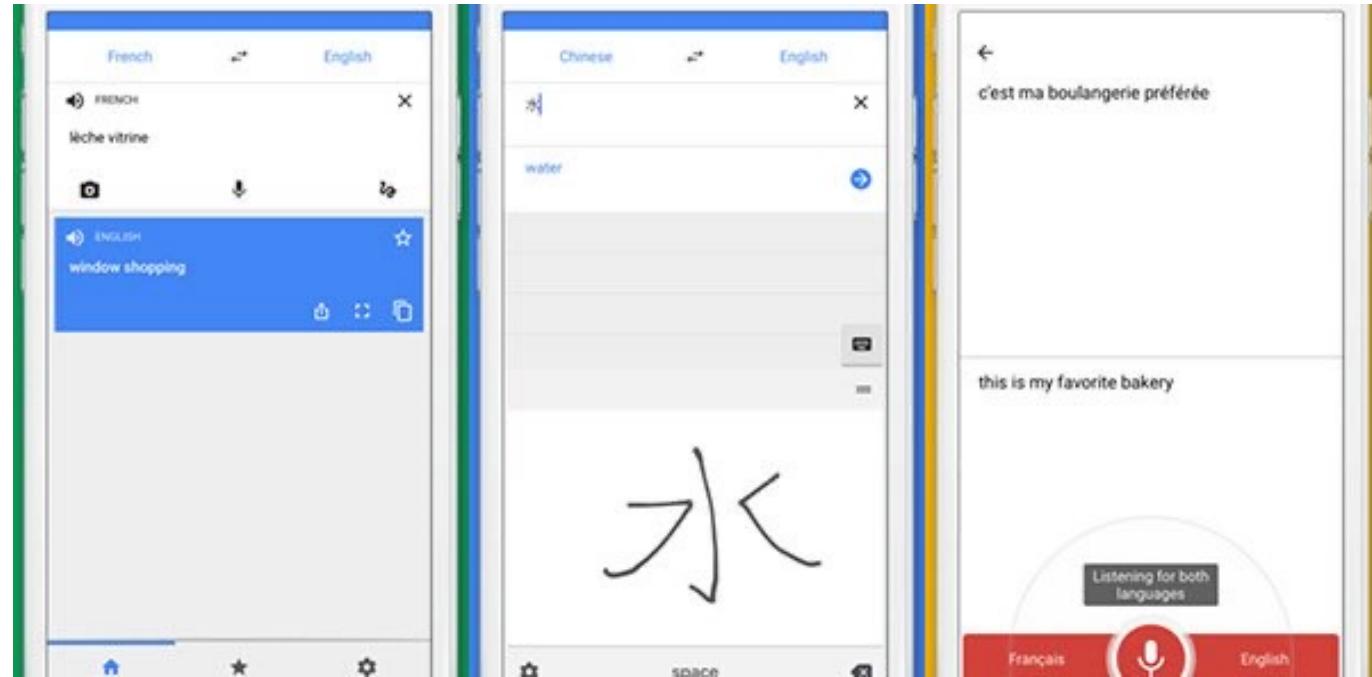


<https://uncubed.com/daily/best-translation-apps-for-travel-in-2019/>

# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...

e.g., Google translate



<https://uncubed.com/daily/best-translation-apps-for-travel-in-2019/>

# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...

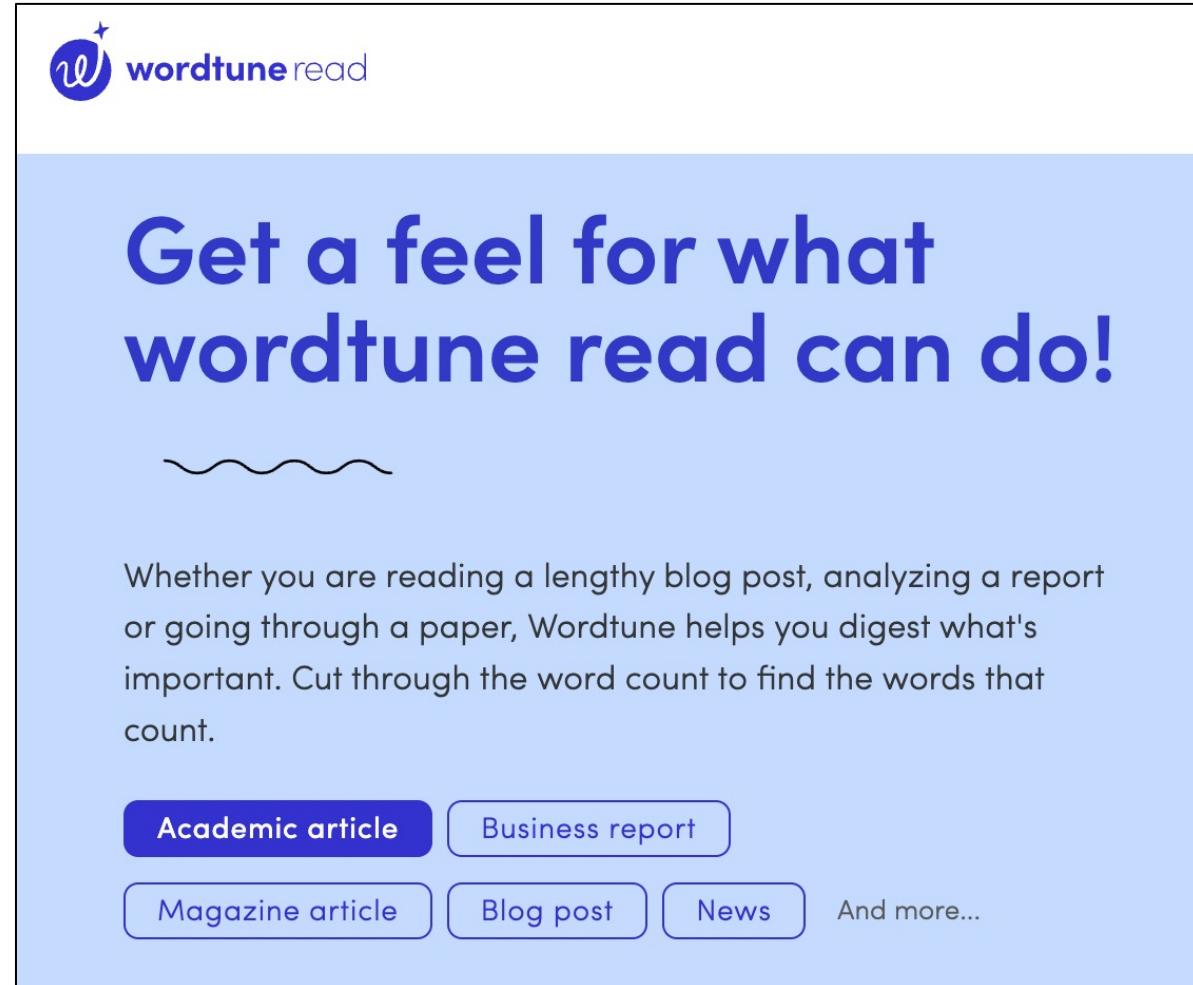
e.g., IBM Watson question answering system (and Jeopardy winner)



<https://www.nytimes.com/2011/02/17/science/17jeopardy-watson.html>

# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...



# Key Challenge: Replicate Language Understanding for So Many Tasks!

- Text classification
- Machine translation
- Question answering
- Automatic summarization
- And more...

# Other Key Challenges: Replicate Language Understanding for So Many Languages/Individuals!

- Need a computable characterization of all human languages that simultaneously captures nuances from individuals; e.g., 7000+ languages spoken around the world



# Today's Topics

- Introduction to natural language processing
- **Text representation**
- Neural word embeddings
- Programming tutorial

# Input: String (Collection of Characters)

Most Relevant ▾



Lives in Austin, Texas

**Keith C. McCormic** Let the food pantries have it instead of monetizing it.

Like · Reply · 1d

↪ 5 Replies



**Caty O'Neil Webb** The promo code isn't working but I found another one on line GETFIFTY% .

Like · Reply · 1d · Edited

↪ 2 Replies

## Machine learning

From Wikipedia, the free encyclopedia

*For the journal, see [Machine Learning \(journal\)](#).*

*"Statistical learning" redirects here. For statistical learning in linguistics, see [statistical learning in language](#)*

**Machine learning** is a field of [computer science](#) that uses statistical techniques to give [computer systems](#) the ability to "learn" (e.g., progressively improve performance on a specific task) with [data](#), without being explicitly programmed.<sup>[2]</sup>

The name *machine learning* was coined in 1959 by [Arthur Samuel](#).<sup>[1]</sup> Machine learning explores the study and construction of [algorithms](#) that can learn from and make predictions on [data](#)<sup>[3]</sup> – such algorithms overcome following strictly static [program instructions](#) by making data-driven predictions or decisions,<sup>[4]:2</sup> through building a [model](#) from sample inputs. Machine learning is employed in a range of computing tasks where designing and programming explicit algorithms with good performance is difficult or infeasible; example applications include [email filtering](#), detection of network intruders, and [computer vision](#).

- Common terms
  - **Corpus:** dataset
  - **Document:** example

# Input: Which “String” Feature Types Apply?

- Categorical data
  - Comes from a fixed list (e.g., education level)
- Structured string data
  - e.g., addresses, dates, telephone numbers,
- Text data

# How to Describe Text to a Computer?

- Challenge: input often varies in length

## Machine learning

From Wikipedia, the free encyclopedia

*For the journal, see [Machine Learning \(journal\)](#).*

*"Statistical learning" redirects here. For statistical learning in linguistics, see [statistical learning in language](#)*

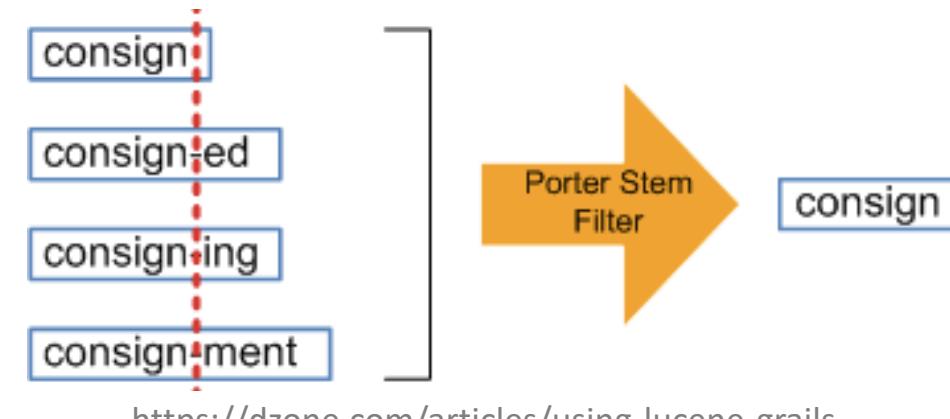
**Machine learning** is a field of [computer science](#) that uses statistical techniques to give [computer systems](#) the ability to "learn" (e.g., progressively improve performance on a specific task) with [data](#), without being explicitly programmed.<sup>[2]</sup>

The name *machine learning* was coined in 1959 by [Arthur Samuel](#).<sup>[1]</sup> Machine learning explores the study and construction of [algorithms](#) that can learn from and make predictions on [data](#)<sup>[3]</sup> – such algorithms overcome following strictly static [program instructions](#) by making data-driven predictions or decisions,<sup>[4]:2</sup> through building a [model](#) from sample inputs. Machine learning is employed in a range of computing tasks where designing and programming explicit algorithms with good performance is difficult or infeasible; example applications include [email filtering](#), detection of network intruders, and [computer vision](#).

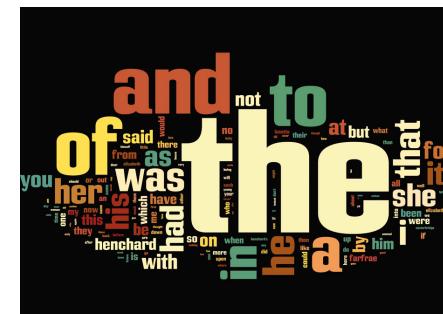
- Solution: convert text to numeric format that DL algorithms can handle

# Implementation Details – Possible Pre-processing

- Lower case all letters
  - Stemming: use each word's stem; e.g., singular to plural, resolve different verb forms
    - e.g.,



- Stop word removal: discard frequent words



<https://github.com/topics/stopwords-removal>

# Converting Text to Vectors

1. Tokenize training data
2. Learn vocabulary
3. Encode data as vectors

# Converting Text to Vectors

1. Tokenize training data; convert data into sequence of tokens (e.g., data ->“This is tokening”)
2. Learn vocabulary
3. Encode data as vectors

Two common approaches:

Character Level

[T] [h] [i] [s] [i] [s] [t] [o] [k] [e] [n] [i] [z] [i] [n] [g] [.]

Word Level

[This] [is] [tokenizing] [.]

# Converting Text to Vectors

1. Tokenize training data
2. Learn vocabulary by identifying all unique tokens in the training data
3. Encode data as vectors

Two common approaches:

Character Level

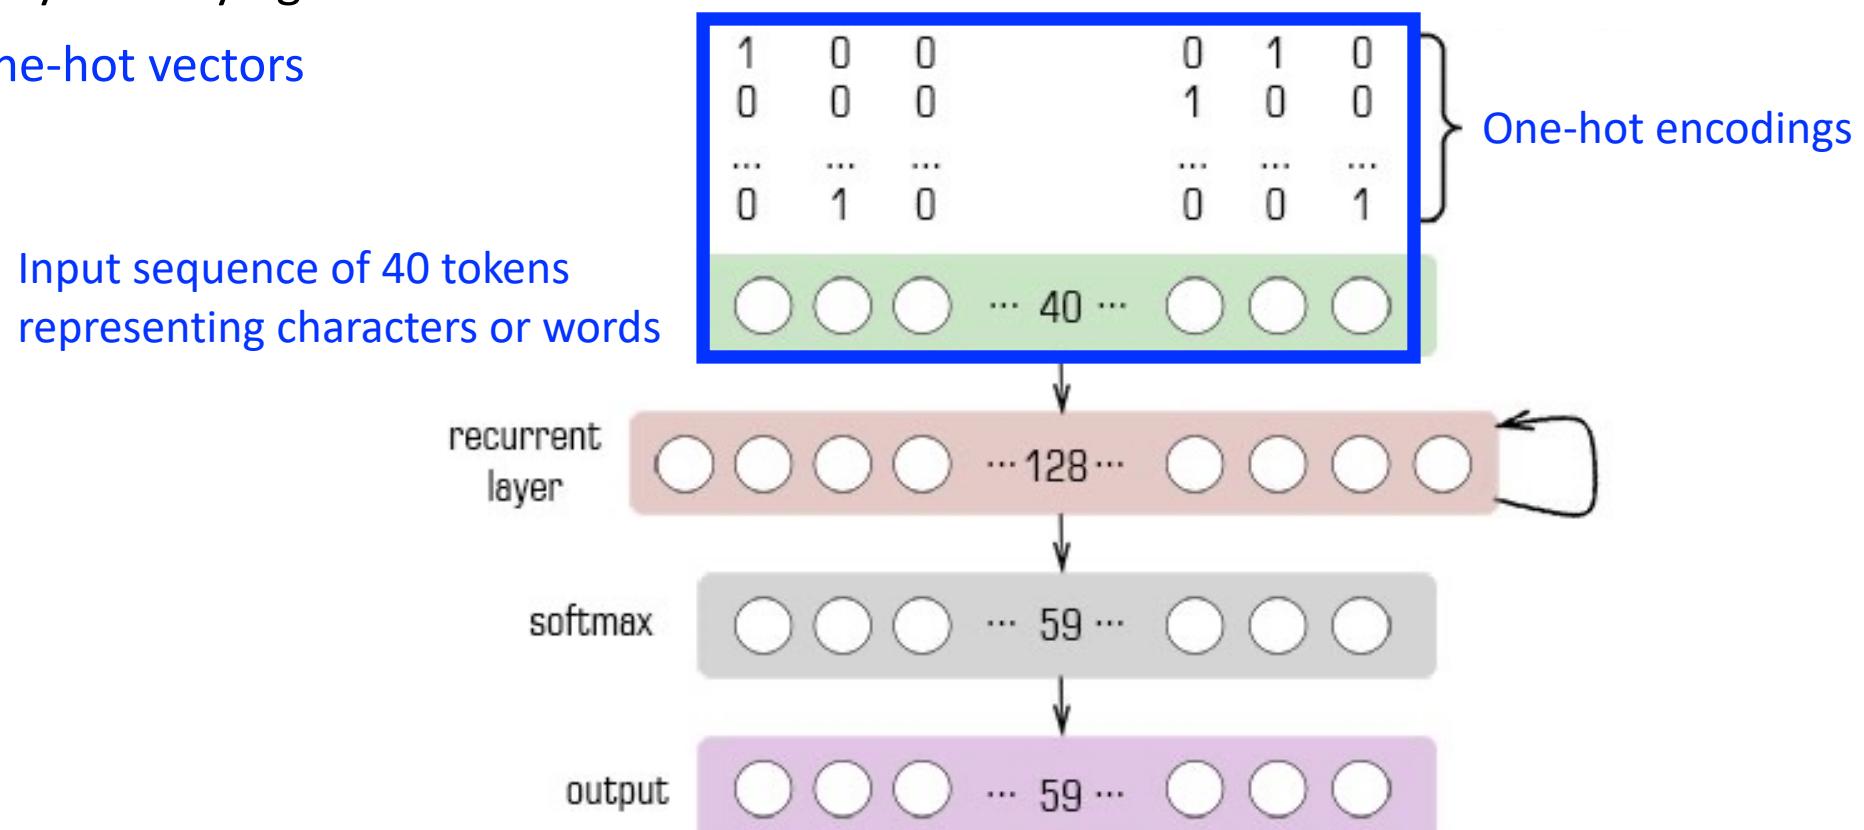
Token	a	b	c	***	0	1	***	!	@	***
Index	1	2	3	***	27	28	***	119	120	***

Word Level

Token	a	an	at	***	bat	ball	***	zipper	zoo	***
Index	1	2	3	***	527	528	***	9,842	9,843	***

# Converting Text to Vectors

1. Tokenize training data
2. Learn vocabulary by identifying all unique tokens in the training data
3. Encode data as one-hot vectors



# Converting Text to Vectors

What are the pros and cons for using word tokens instead of character tokens?

Character Level

Token	a	b	c	***	0	1	***	!	@	***
Index	1	2	3	***	27	28	***	119	120	***

Word Level

Token	a	an	at	***	bat	ball	***	zipper	zoo	***
Index	1	2	3	***	527	528	***	9,842	9,843	***

- Pros: length of input/output sequences is shorter, simplifies learning semantics
- Cons: “UNK” word token needed for out of vocabulary words; vocabulary can be large

# Converting Text to Vectors

Character Level

Token	a	b	c	***	0	1	***	!	@	***
Index	1	2	3	***	27	28	***	119	120	***

Word Level

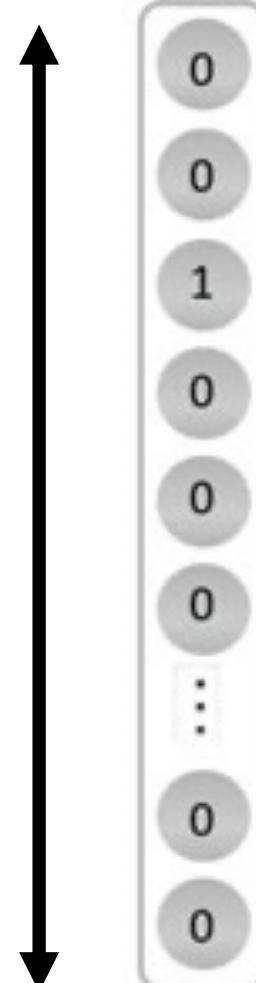
Token	a	an	at	***	bat	ball	***	zipper	zoo	***
Index	1	2	3	***	527	528	***	9,842	9,843	***

Word level representations are more commonly used

# Problems with One-Hot Encoding Words?

Dimensionality = vocabulary size

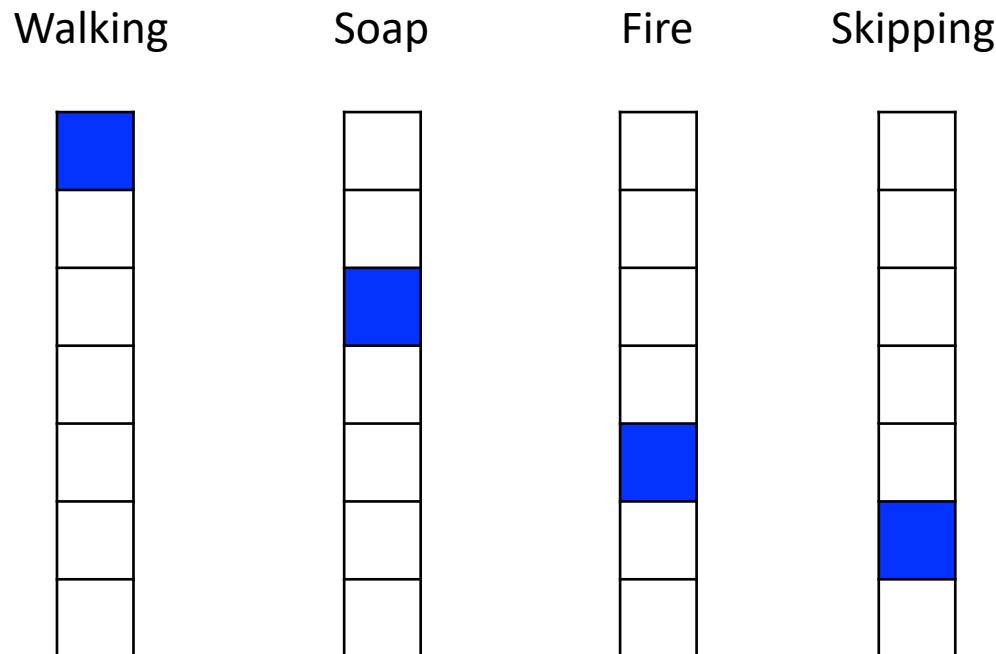
e.g., English has ~170,000 words  
with ~10,000 commonly used words



- Huge memory burden
- Computationally expensive

# Limitation of One-Hot Encoding Words

- No notion of which words are similar, yet such understanding can improve generalization
  - e.g., “walking”, “running”, and “skipping” are all suitable for “He was \_\_\_\_\_ to school.”

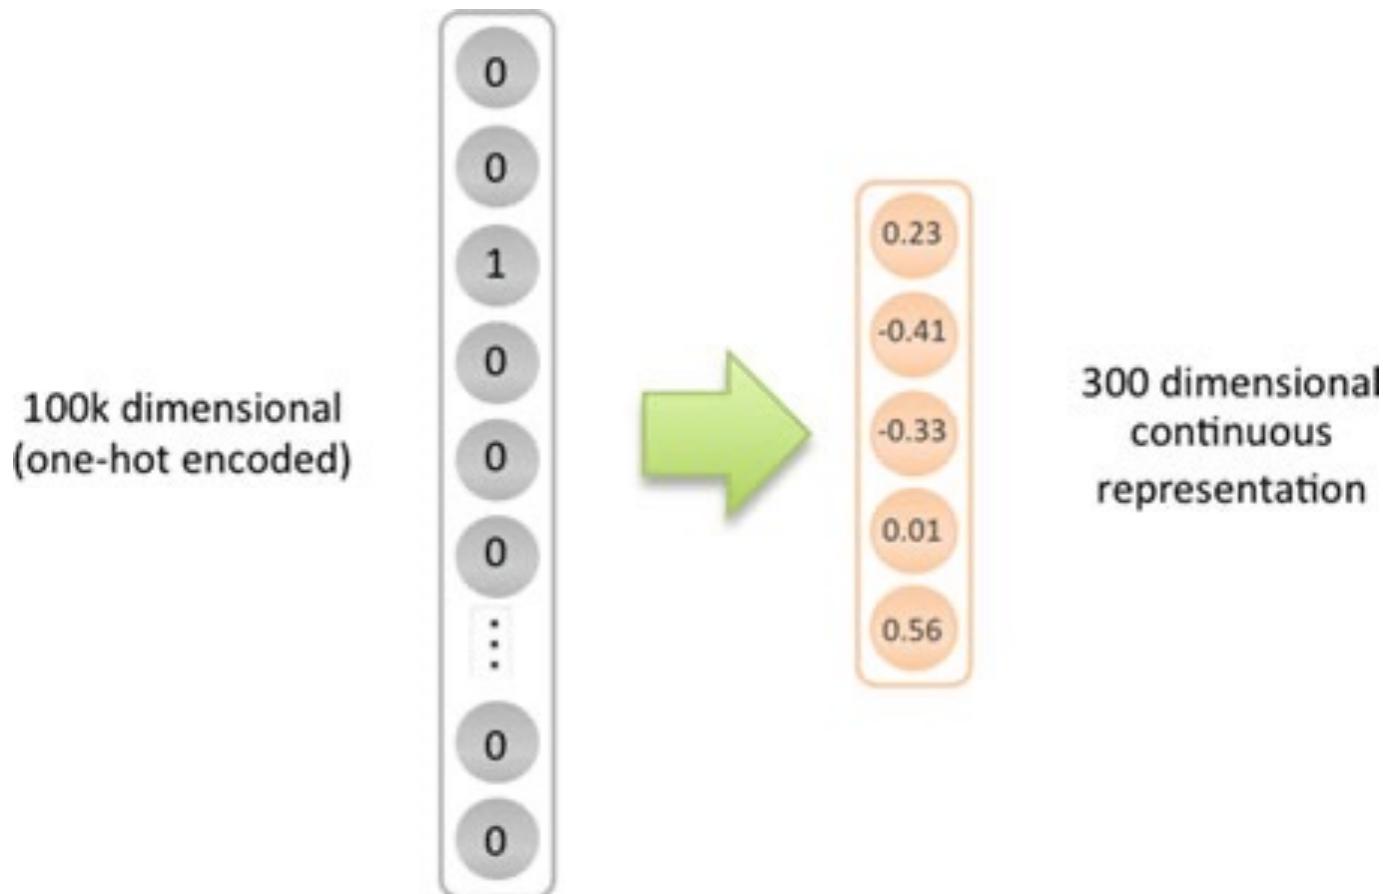


The distance between  
all words is equal!

# Today's Topics

- Introduction to natural language processing
- Text representation
- Neural word embeddings
- Programming tutorial

Idea: Represent Each Word Compactly in a Space Where Vector Distance Indicates Word Similarity



# Inspiration: Distributional Semantics

“The distributional hypothesis says that the meaning of a word is derived from the context in which it is used, and words with similar meaning are used in similar contexts.”

- Origins: Harris in 1954 and Firth in 1957

# Inspiration: Distributional Semantics

“The distributional hypothesis says that **the meaning of a word is derived from the context in which it is used**, and words with similar meaning are used in similar contexts.”

# Inspiration: Distributional Semantics

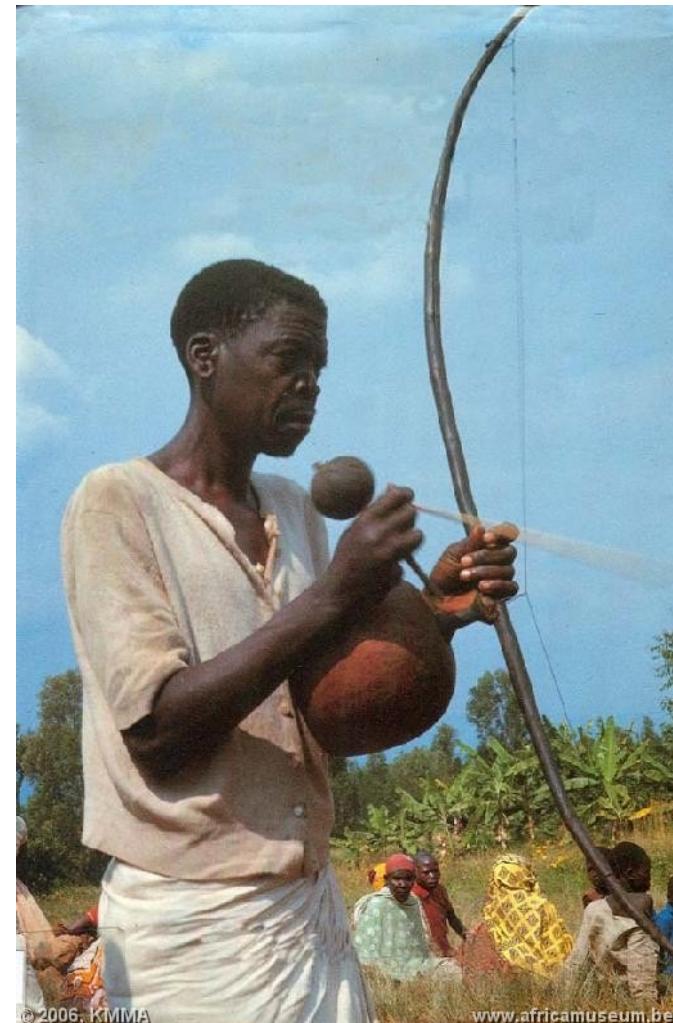
- What is the meaning of **berimbau** based on **context**?

Background music from a **berimbau** offers a beautiful escape.

Many people danced around the **berimbau** player.

I practiced for many years to learn how to play the **berimbau**.

- Idea: **context** makes it easier to understand a word's meaning



© 2006, KIMMA

[www.africamuseum.be](http://www.africamuseum.be)

# Inspiration: Distributional Semantics

“The distributional hypothesis says that the meaning of a word is derived from the context in which it is used, and words with similar meaning are used in similar contexts.”

# Inspiration: Distributional Semantics

- What other words could fit into these contexts?

1. Background music from a \_\_\_\_\_ offers a beautiful escape.
2. Many people danced around the \_\_\_\_\_ player.
3. I practiced for many years to learn how to play the \_\_\_\_\_.

Hypothesis is that words with similar row values have similar meanings

	1.	2.	3.	Contexts
Berimbau	1	1	1	
Soap	0	0	0	
Fire	0	0	0	
Guitar	1	1	1	

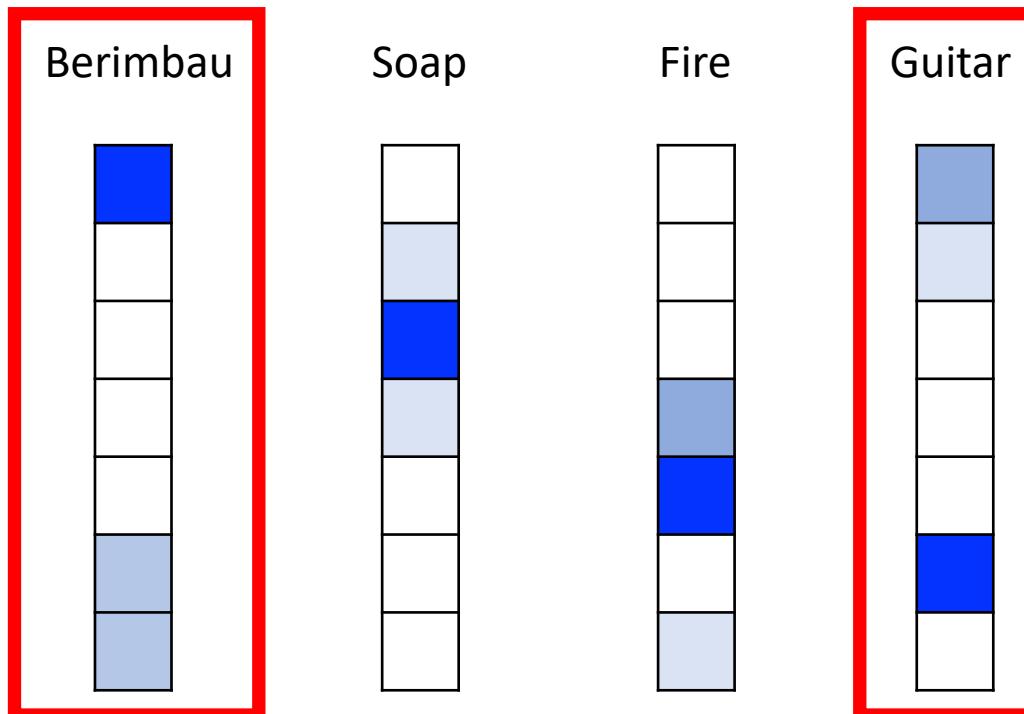
1 if a word can appear in the context  
0 otherwise

# Inspiration: Distributional Semantics

“The distributional hypothesis says that the meaning of a word is derived from the context in which it is used, and words with similar meaning are used in similar contexts.”

# Approach

- Learn a dense (lower-dimensional) vector for each word by characterizing its **context**, which inherently will reflect similarity/differences to other words



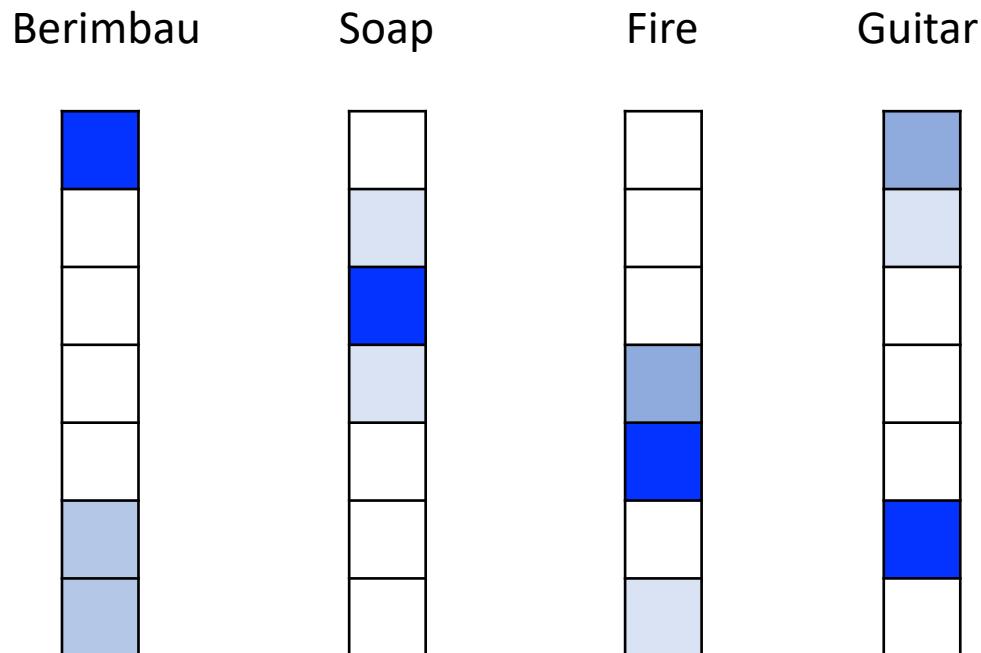
Berimbau and guitar are the closest word pair

The distance between  
each pair of words differs!

Note: many ways to measure  
distance (e.g., cosine distance)

# Approach

- Learn a dense (lower-dimensional) vector for each word by characterizing its **context**, which inherently will reflect similarity/differences to other words

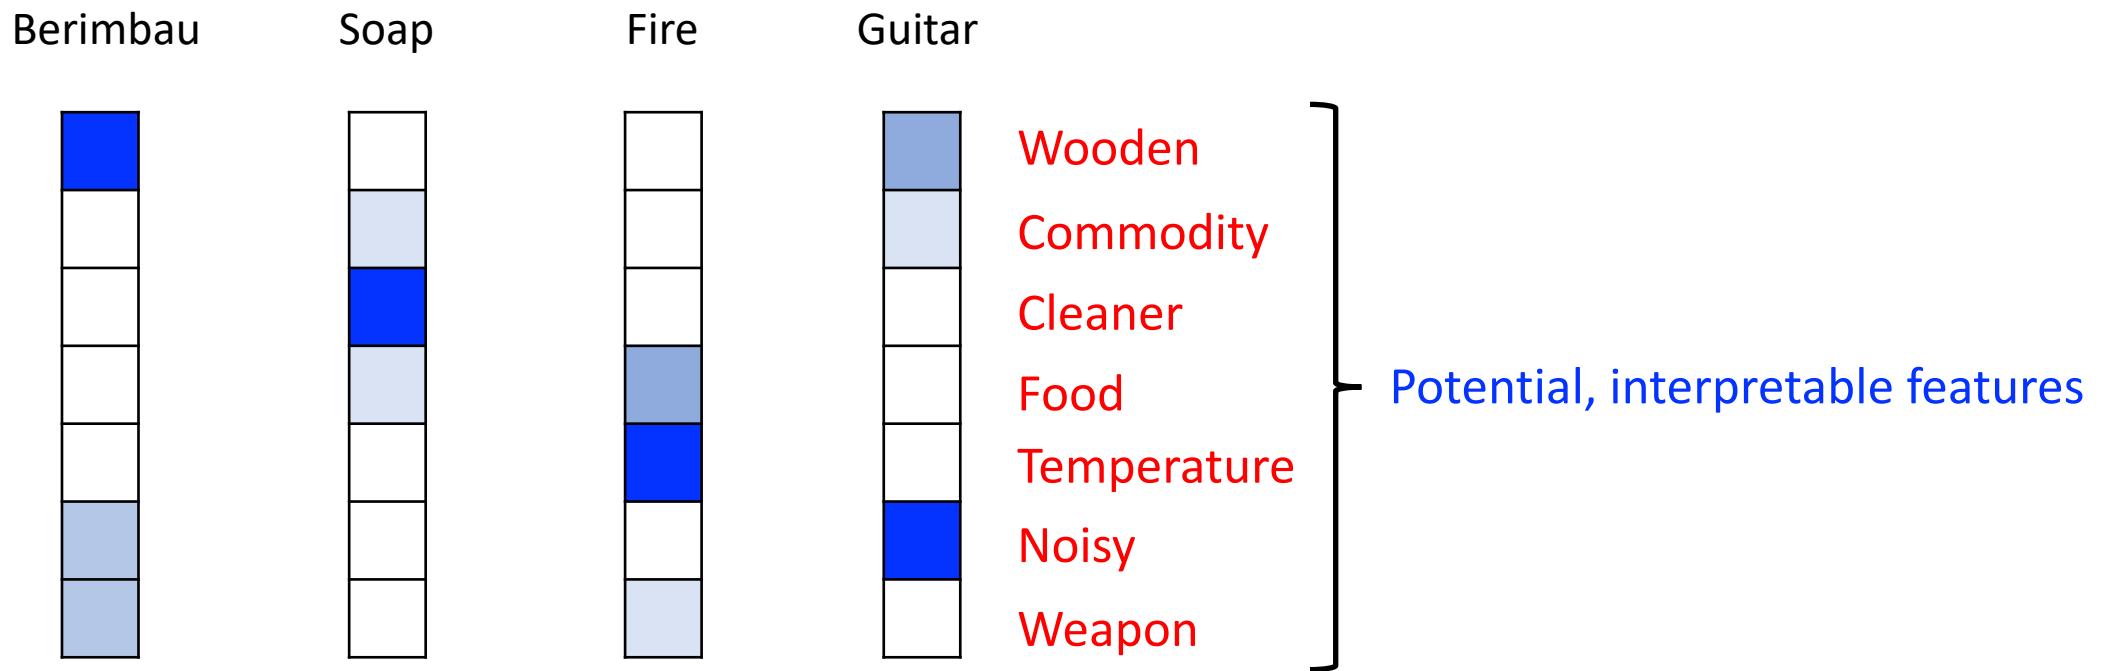


We embed words in a shared space so they can be compared with a few features

What features would discriminate these words?

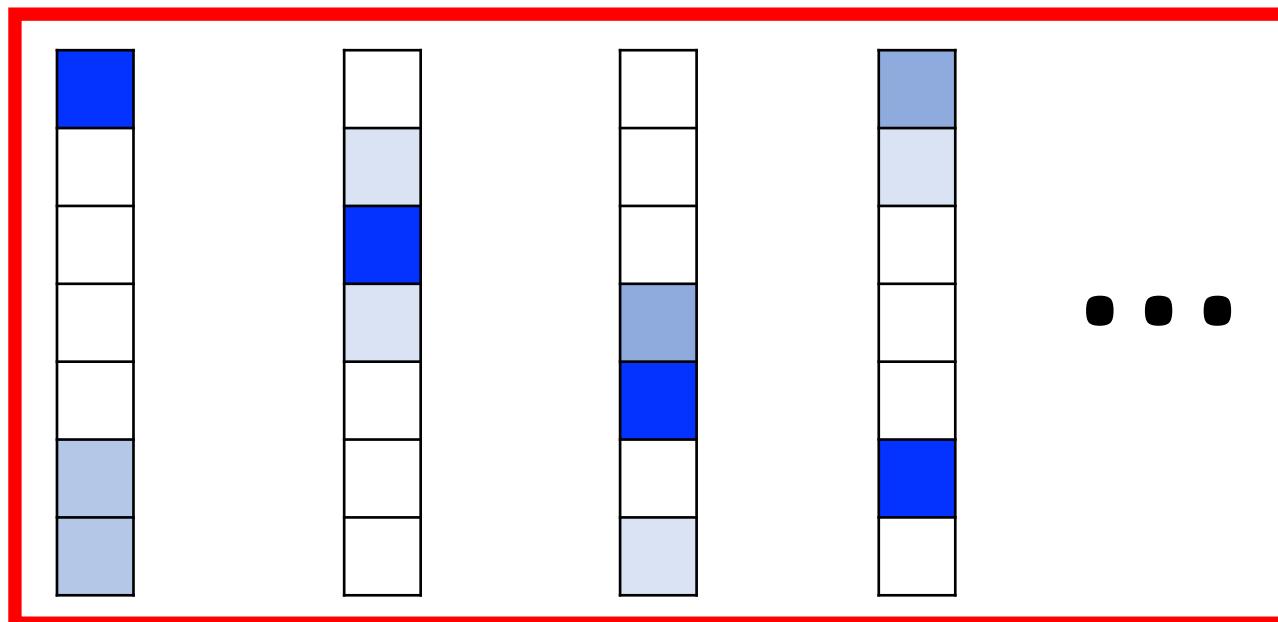
# Approach

- Learn a dense (lower-dimensional) vector for each word by characterizing its **context**, which inherently will reflect similarity/differences to other words



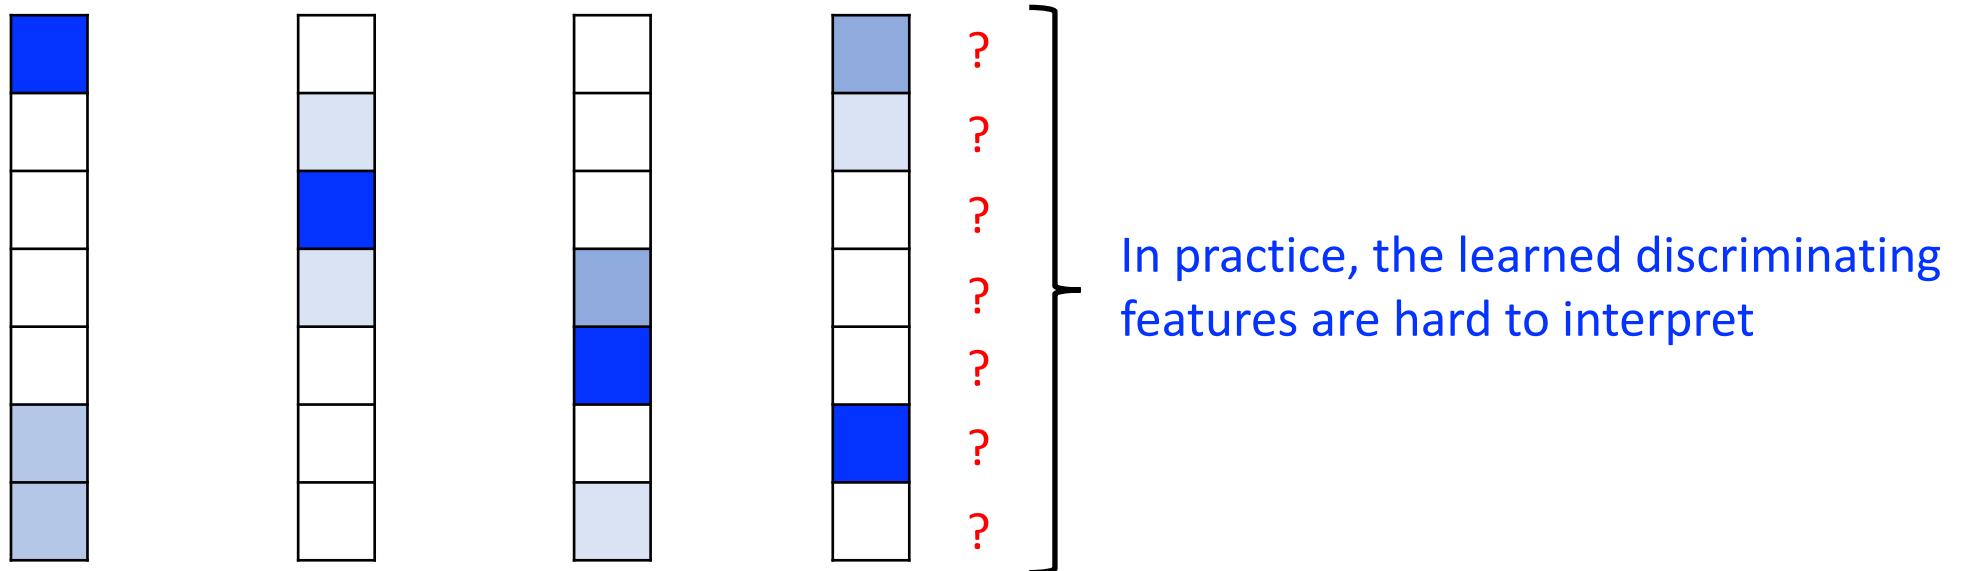
# Approach: Learn Word Embedding Space

- An **embedding space** represents a finite number of words, decided in training
- A **word embedding** is represented as a vector indicating its context
- The dimensionality of all word embeddings in an embedding space match
  - What is the dimensionality for the shown example?



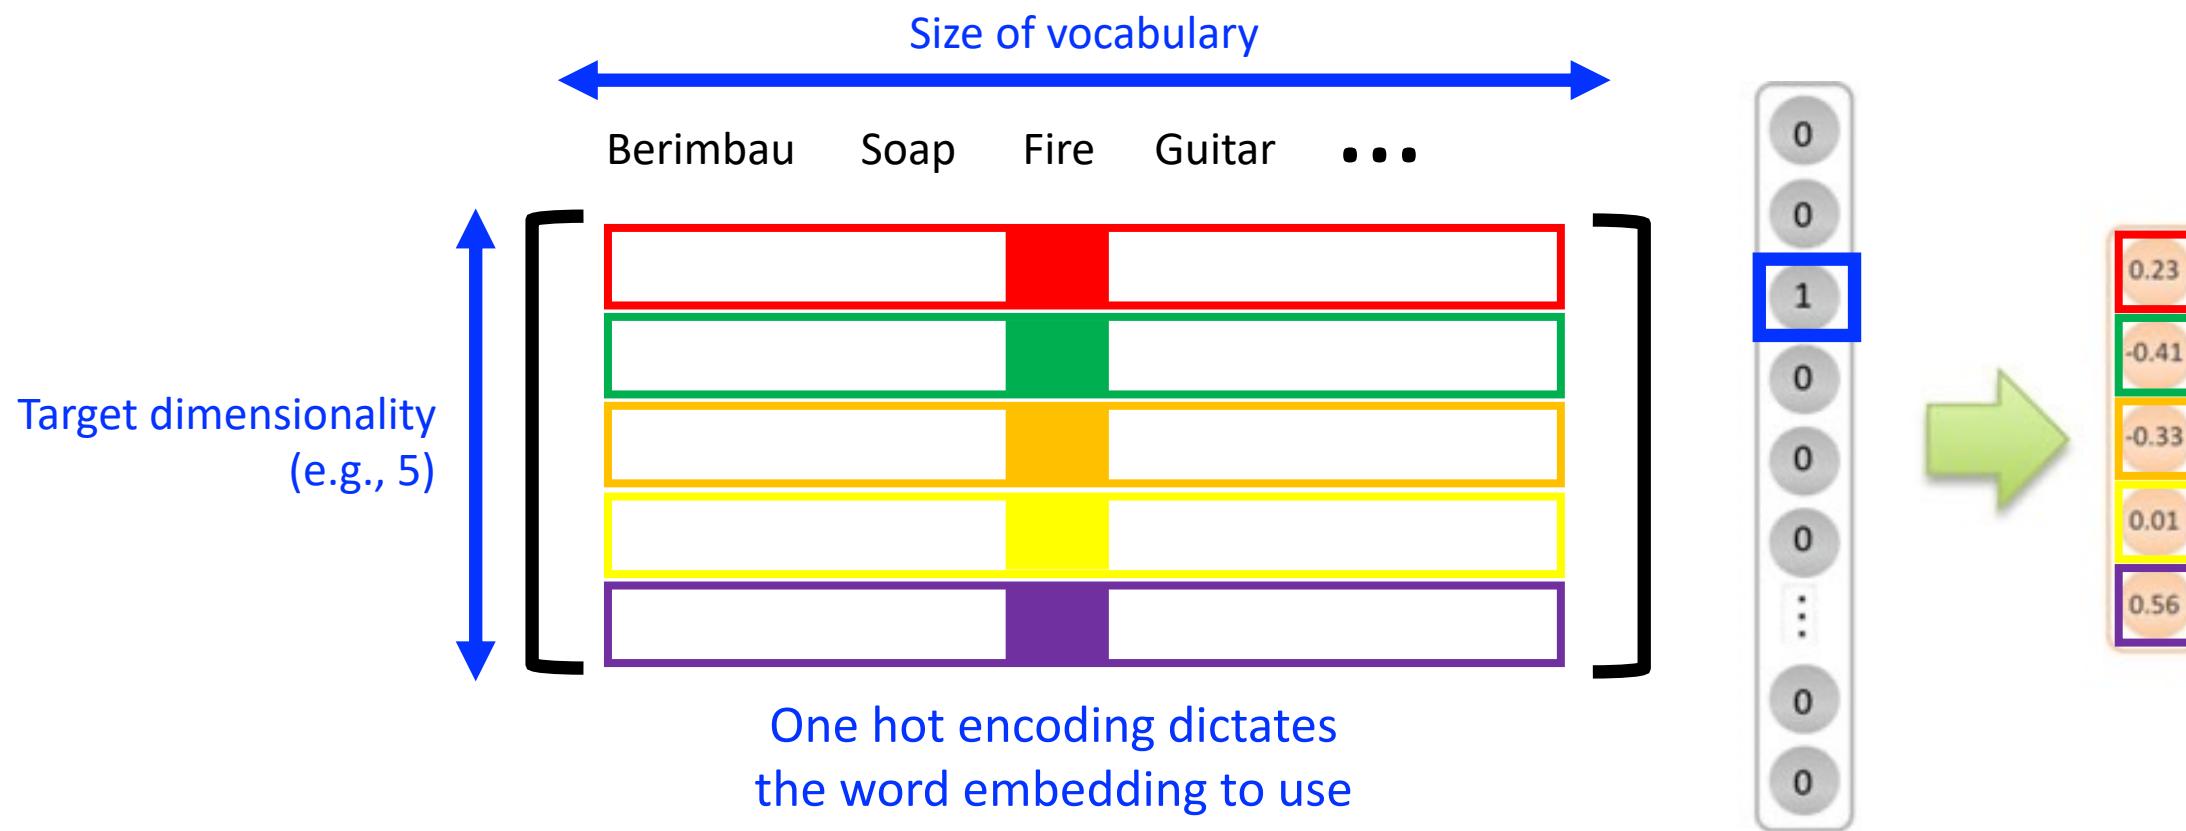
# Approach: Learn Word Embedding Space

- An embedding space represents a finite number of words, defined in training
- A word embedding is represented as a vector indicating its context
- The dimensionality of all word embeddings in an embedding space match



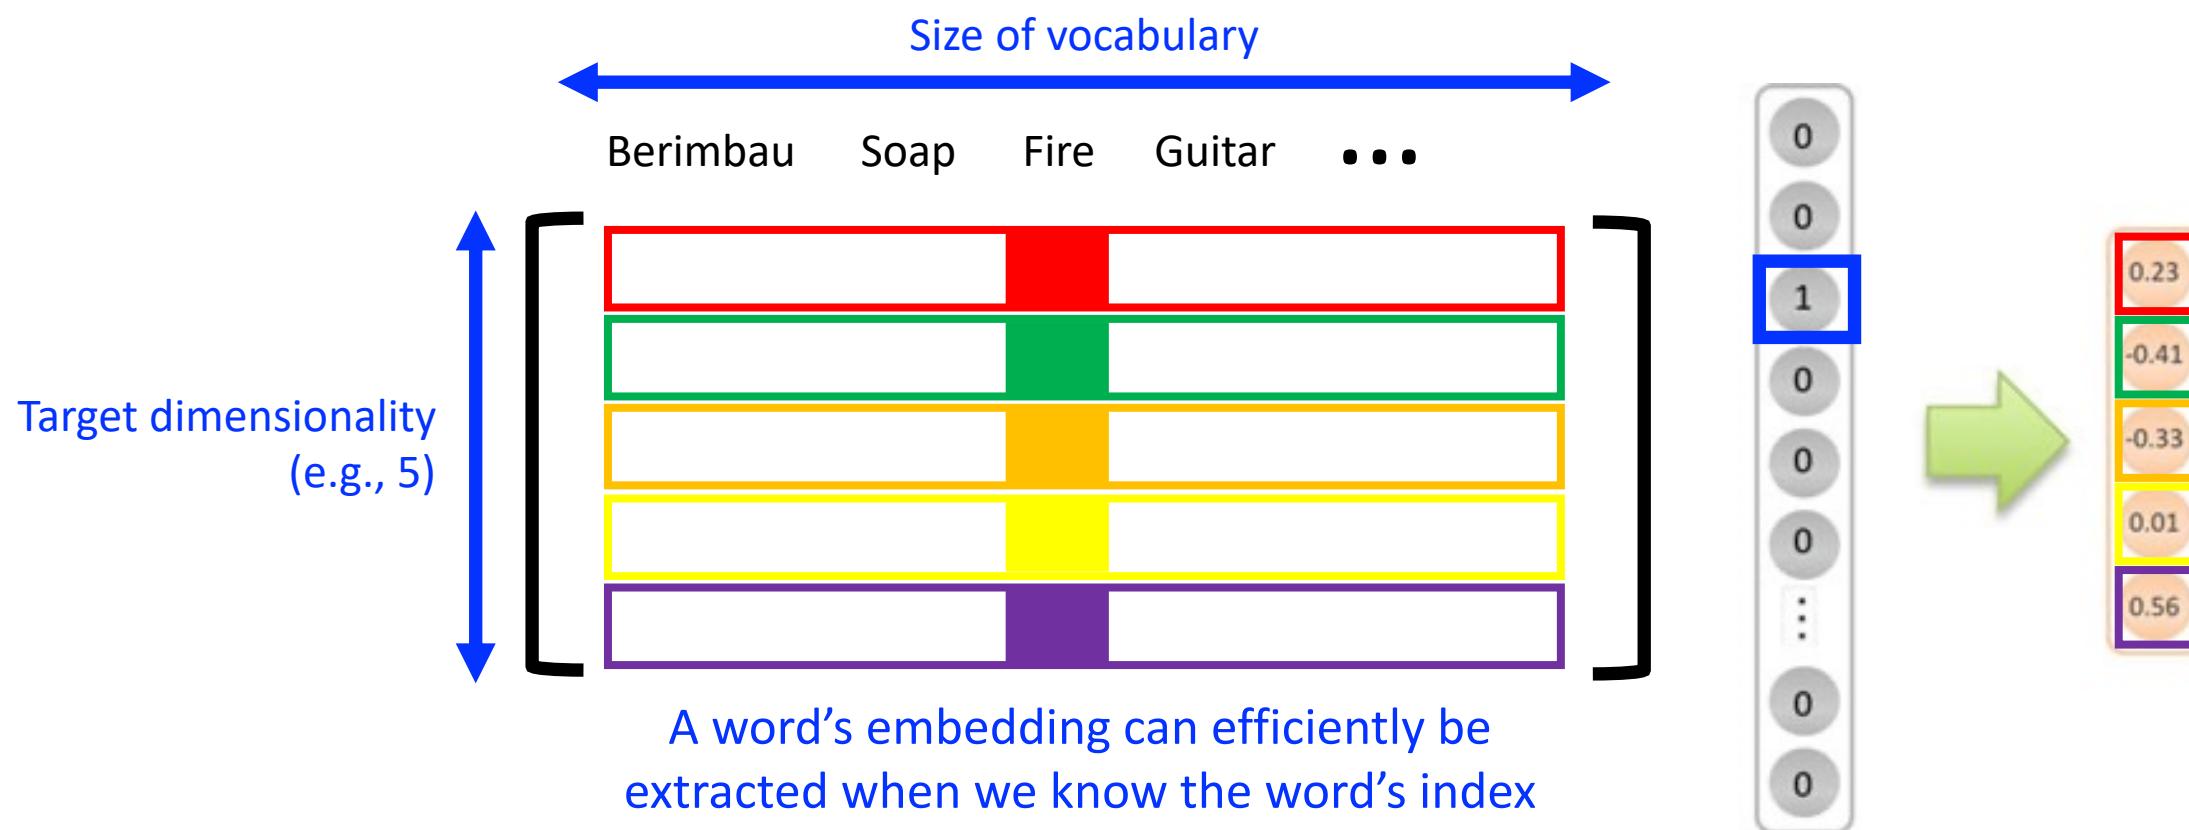
# Embedding Matrix

- The embedding matrix converts an input word into a dense vector



# Embedding Matrix

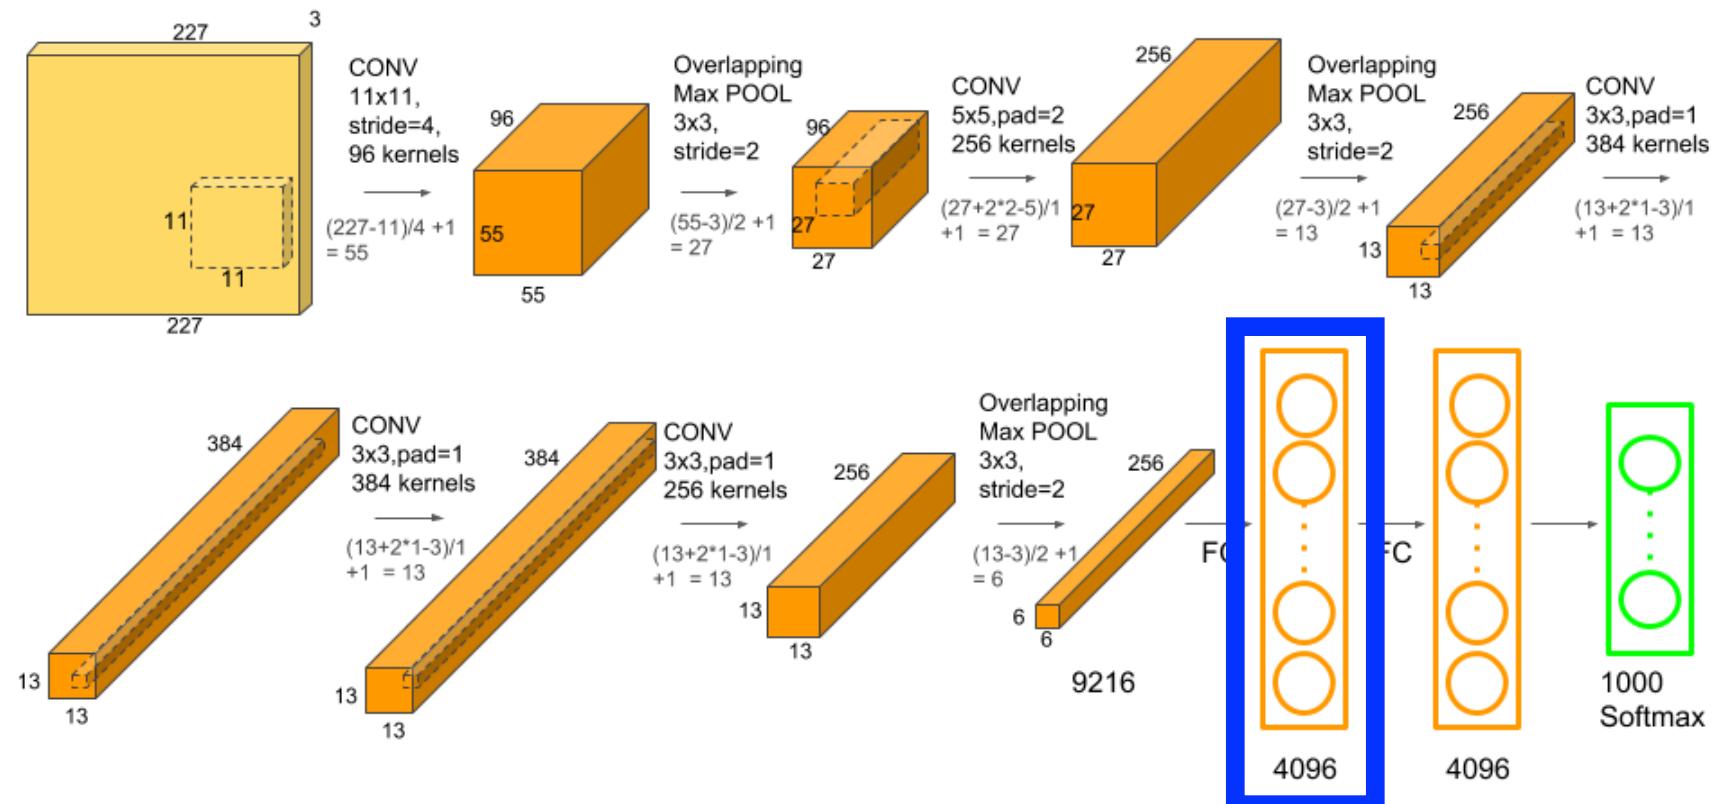
- It converts an input word into a dense vector



# Word Embedding Analogous to a CNN Pretrained Feature

- e.g., FC6 layer of AlexNet

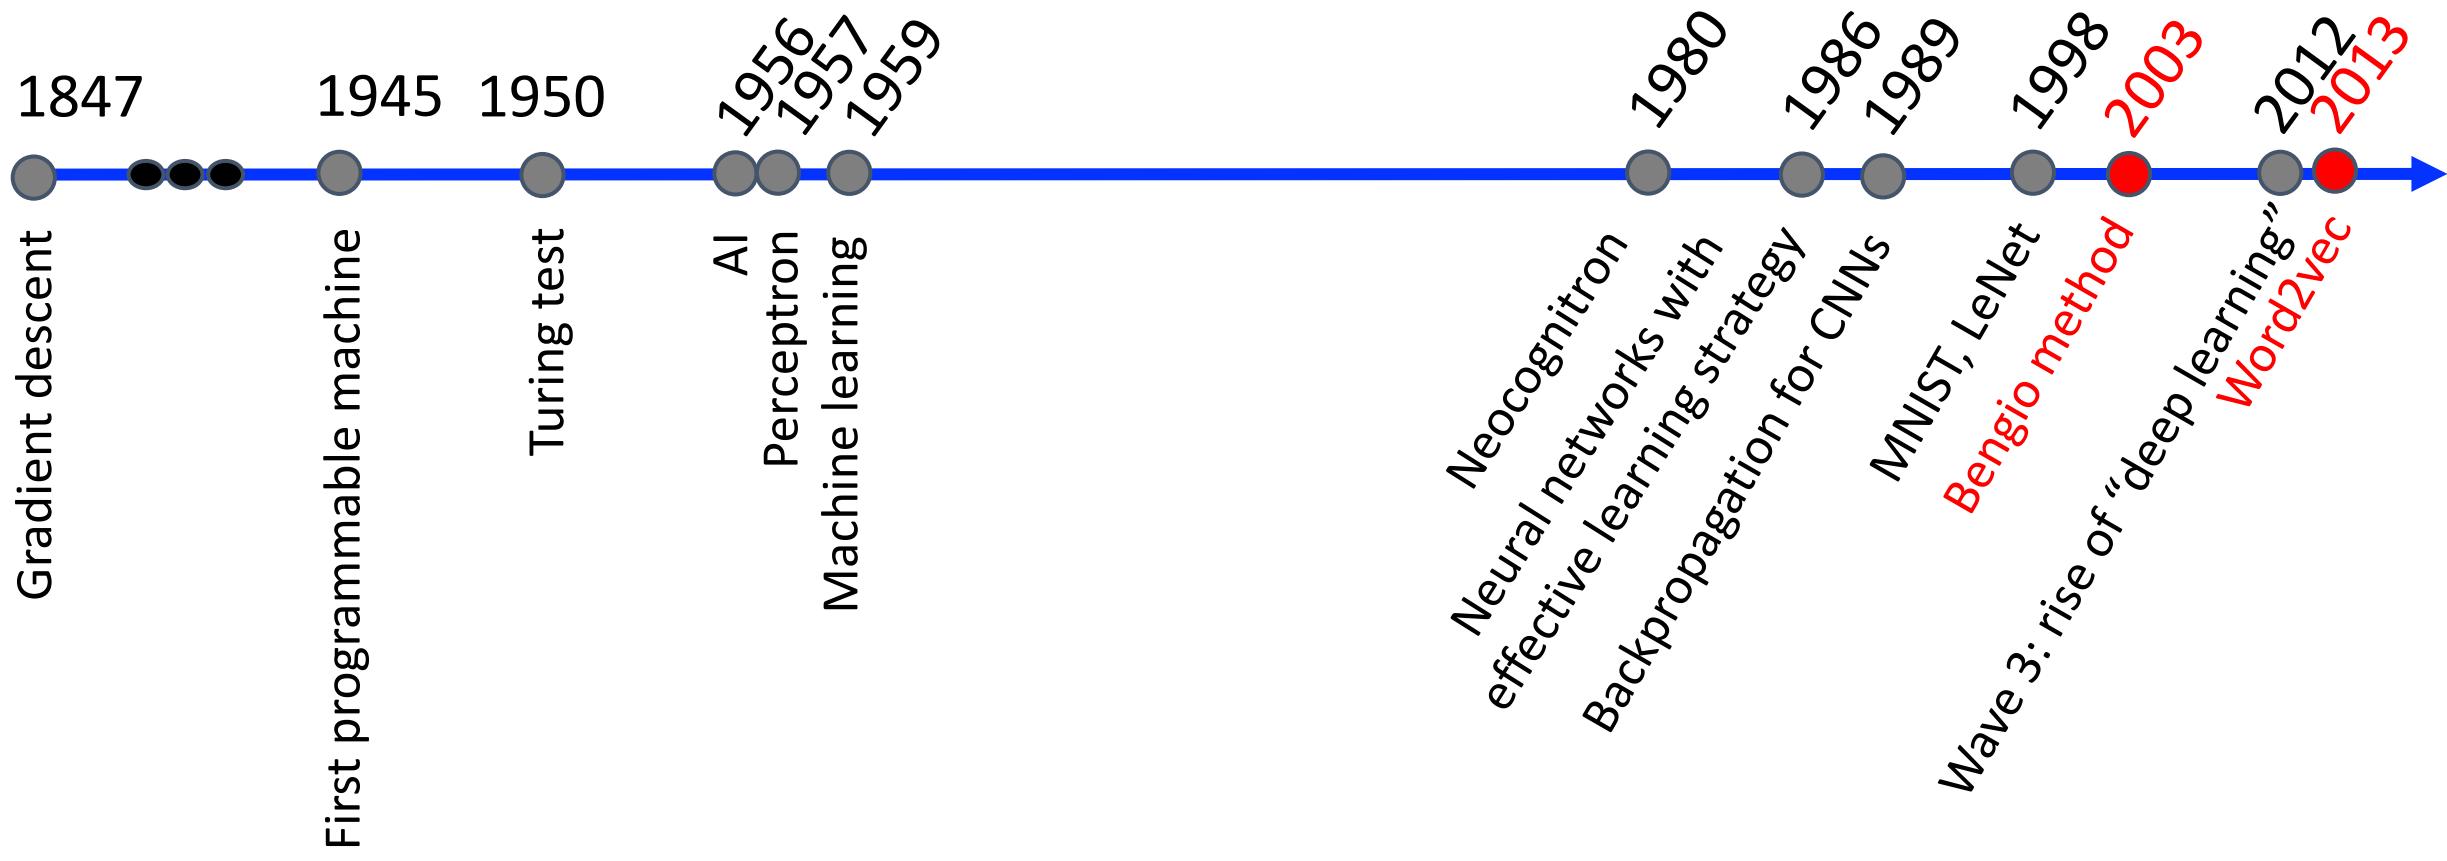
A representation of the data extracted inside a network (rather than the input or predicted output)



# Popular Word Embeddings

- Bengio method
- Word2vec (skip-gram model)
- And more...

# Historical Context



# Popular Word Embeddings

- Bengio method
- Word2vec (skip-gram model)
- And more...

# Idea: Learn Word Embeddings That Help Predict Viable Next Words

e.g.,

1. Background music from a \_\_\_\_\_
2. Many people danced around the \_\_\_\_\_
3. I practiced for many years to learn how to play the \_\_\_\_\_

# Task: Predict Next Word Given Previous Ones

e.g.,

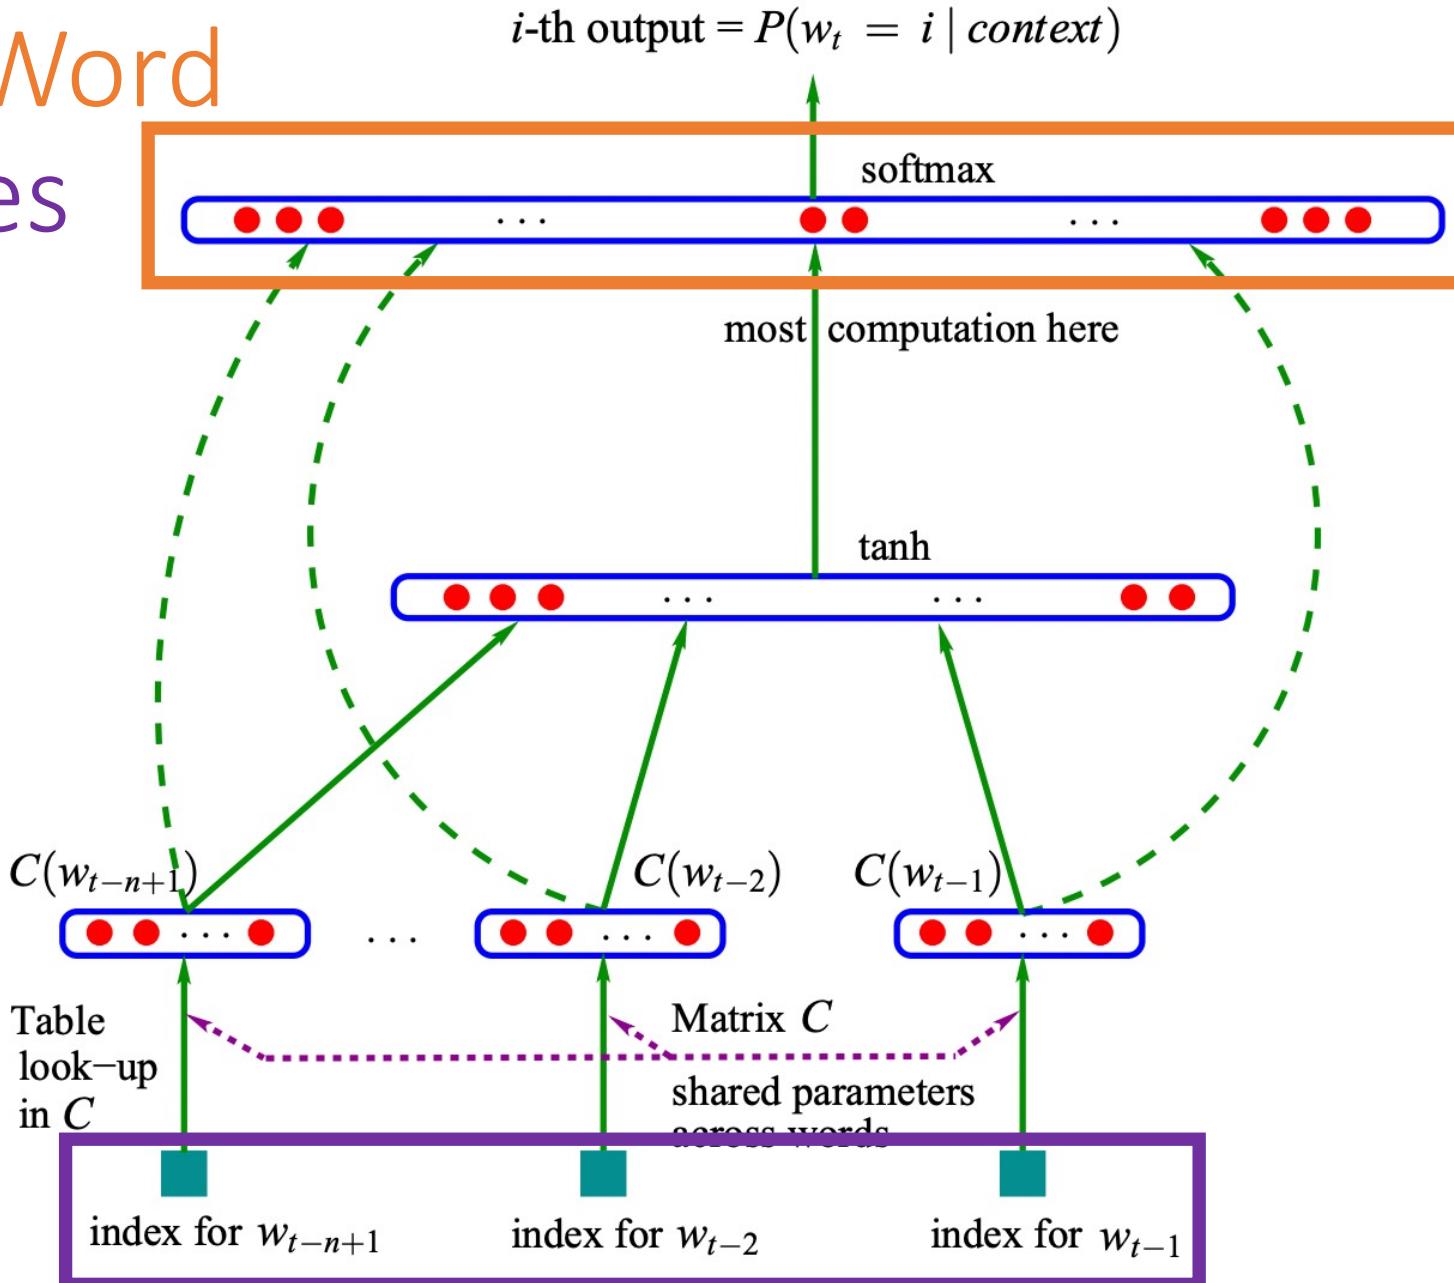
1. Background music from a \_\_\_\_\_
2. Many people danced around the \_\_\_\_\_
3. I practiced for many years to learn how to play the \_\_\_\_\_

# Task: Predict Next Word Given Previous Ones

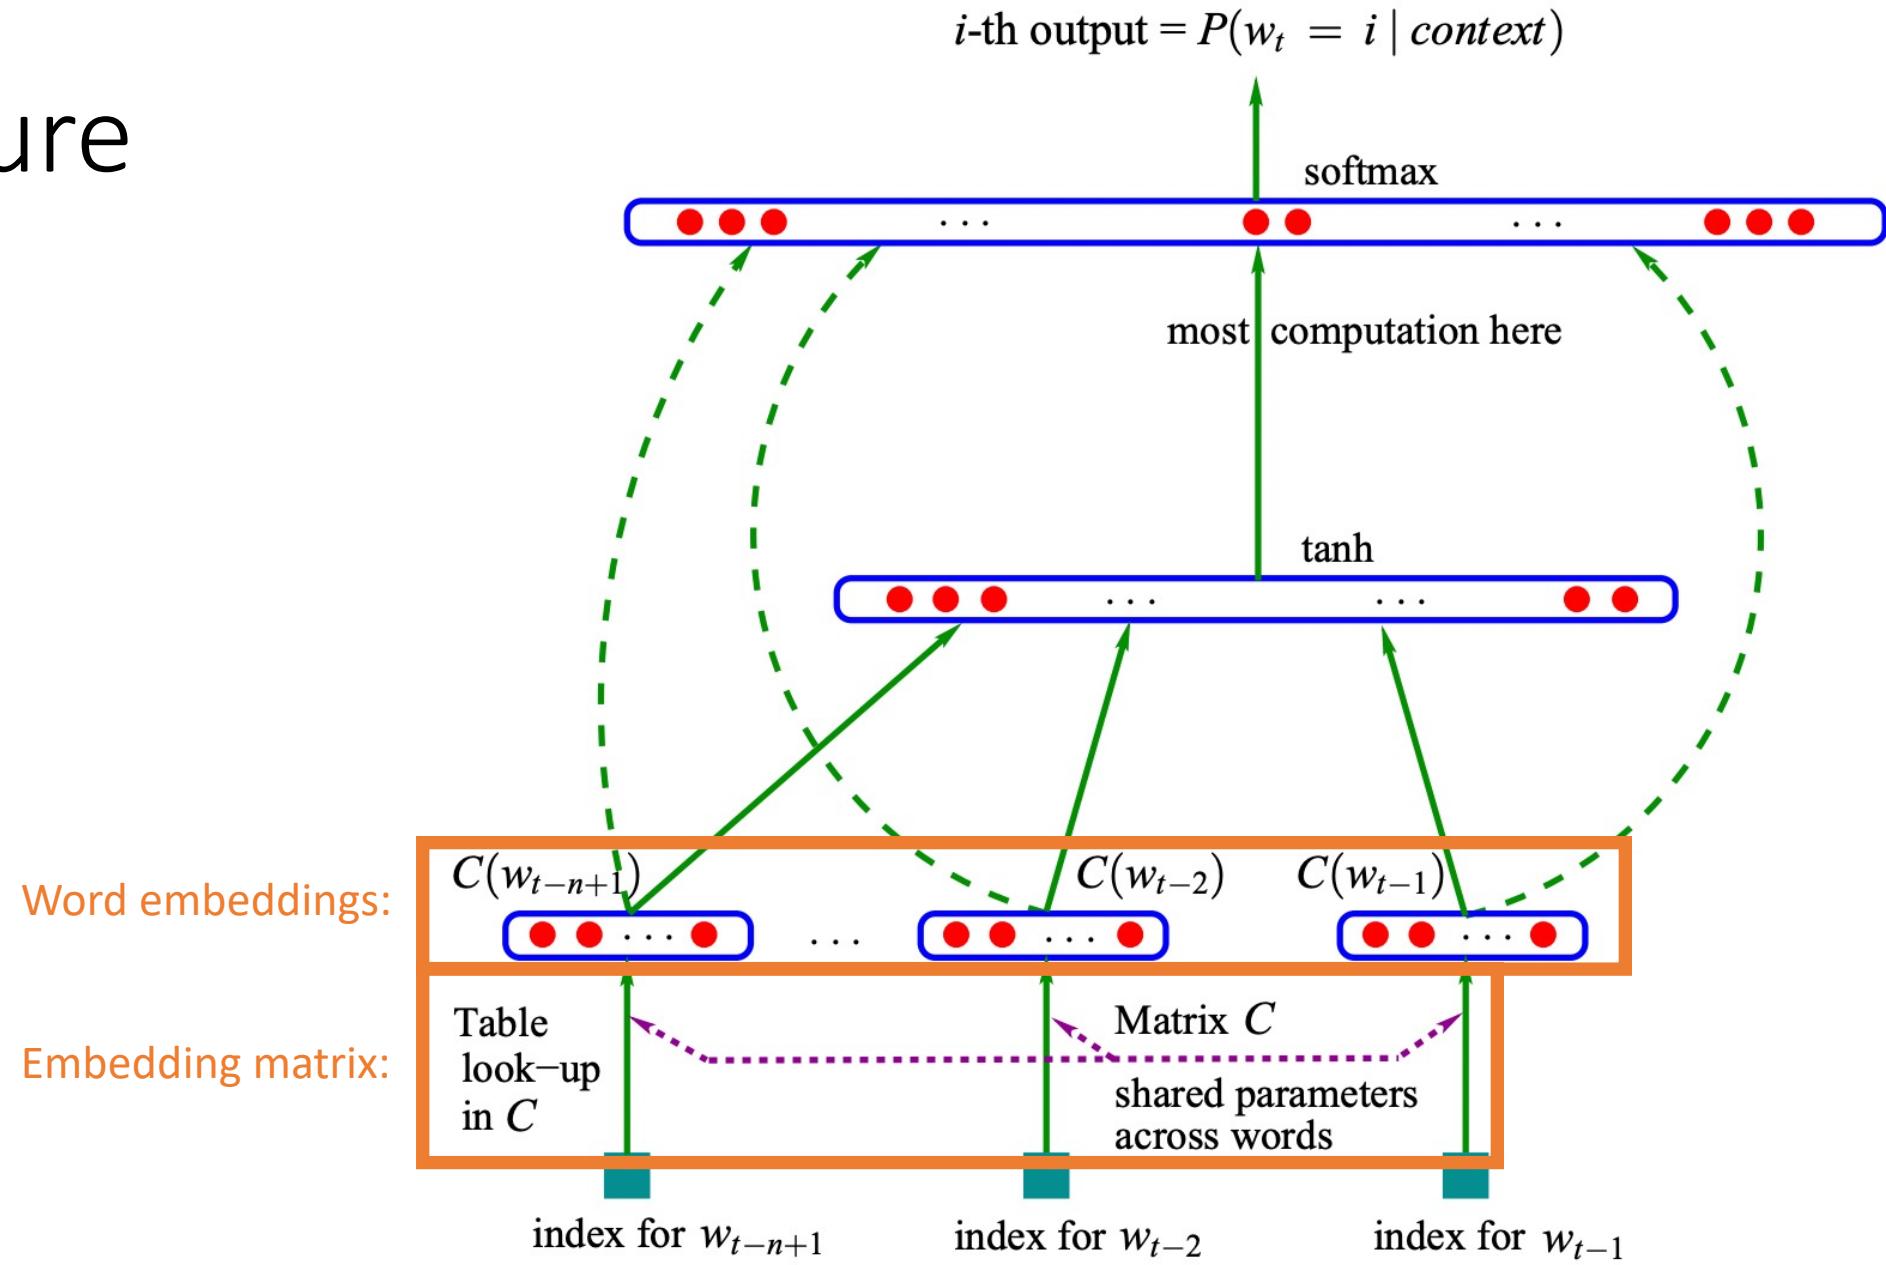
e.g., a vocabulary size of 17,000 was used in experiments

What is the dimensionality of the output layer?

- 17,000 (each indexed position indicates probability of a word)



# Architecture

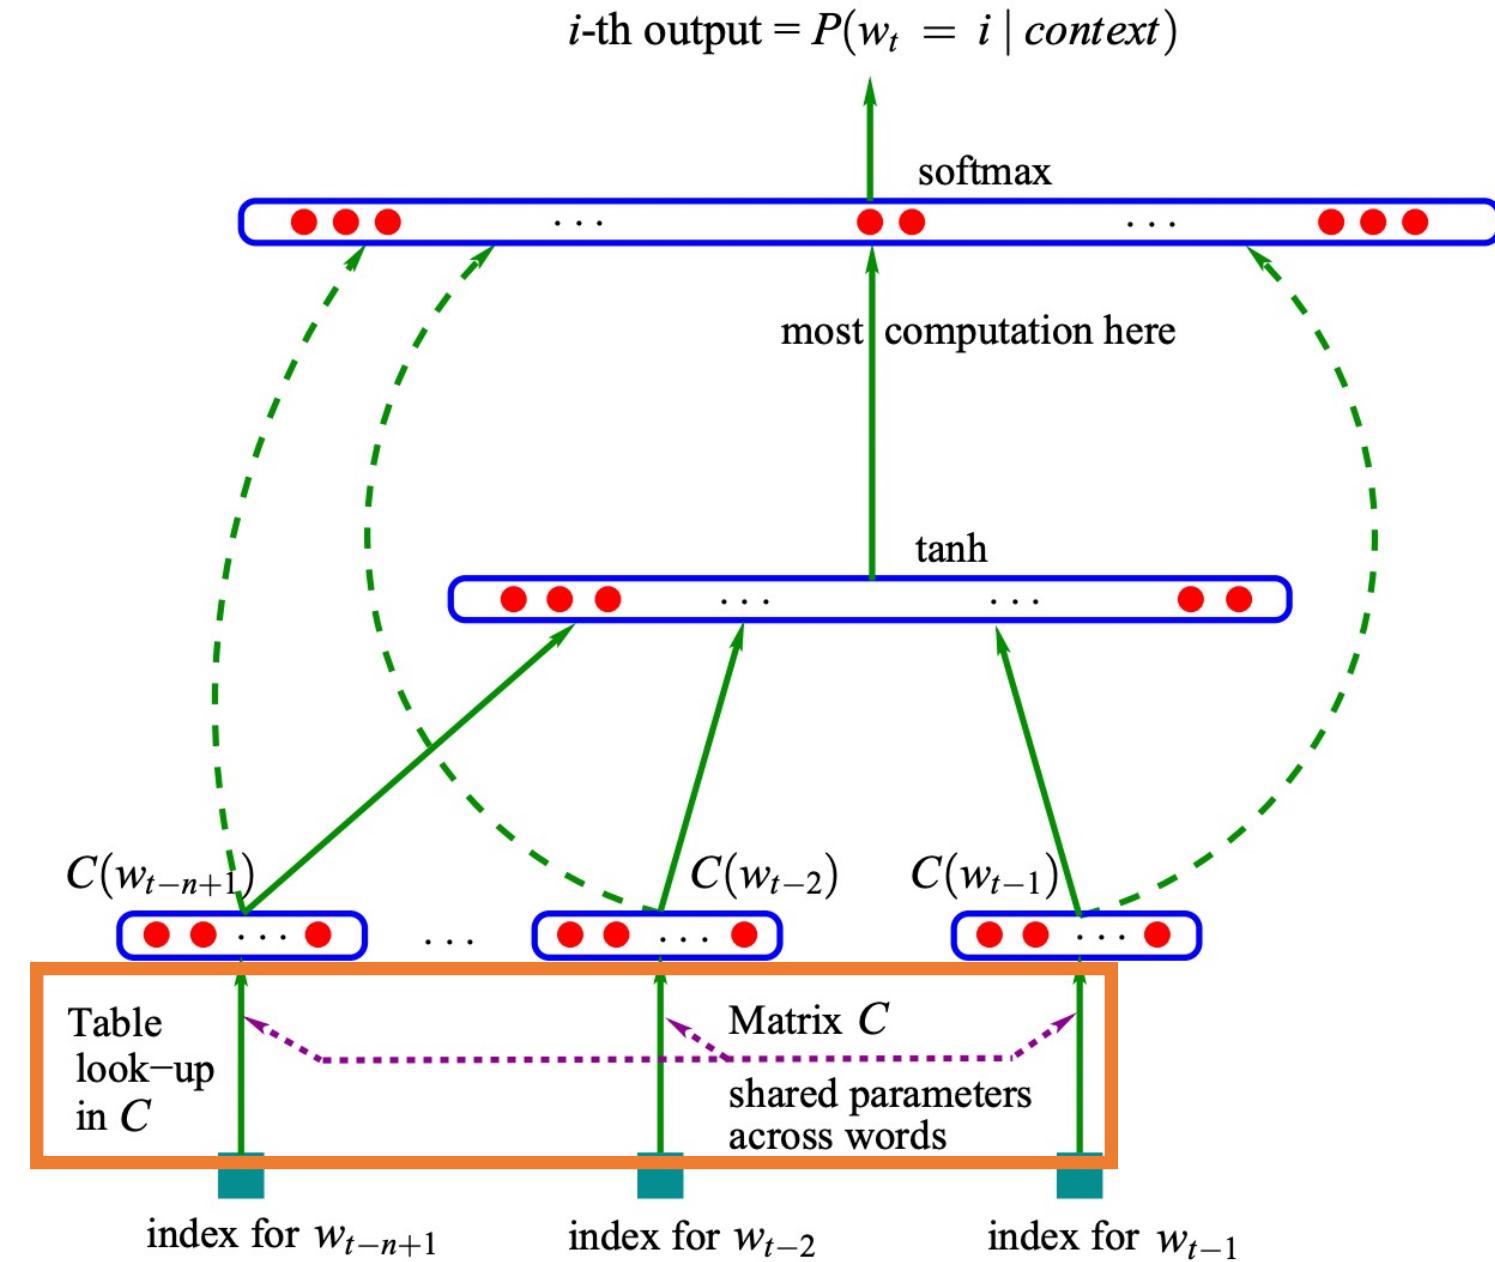


# Architecture

e.g., a vocabulary size of 17,000 was used with embedding sizes of 30, 60, and 100 in experiments

Assume a 30-d word embedding  
- what are the dimensions of the embedding matrix C?

$30 \times 17,000$  (i.e., 510,000 weights)

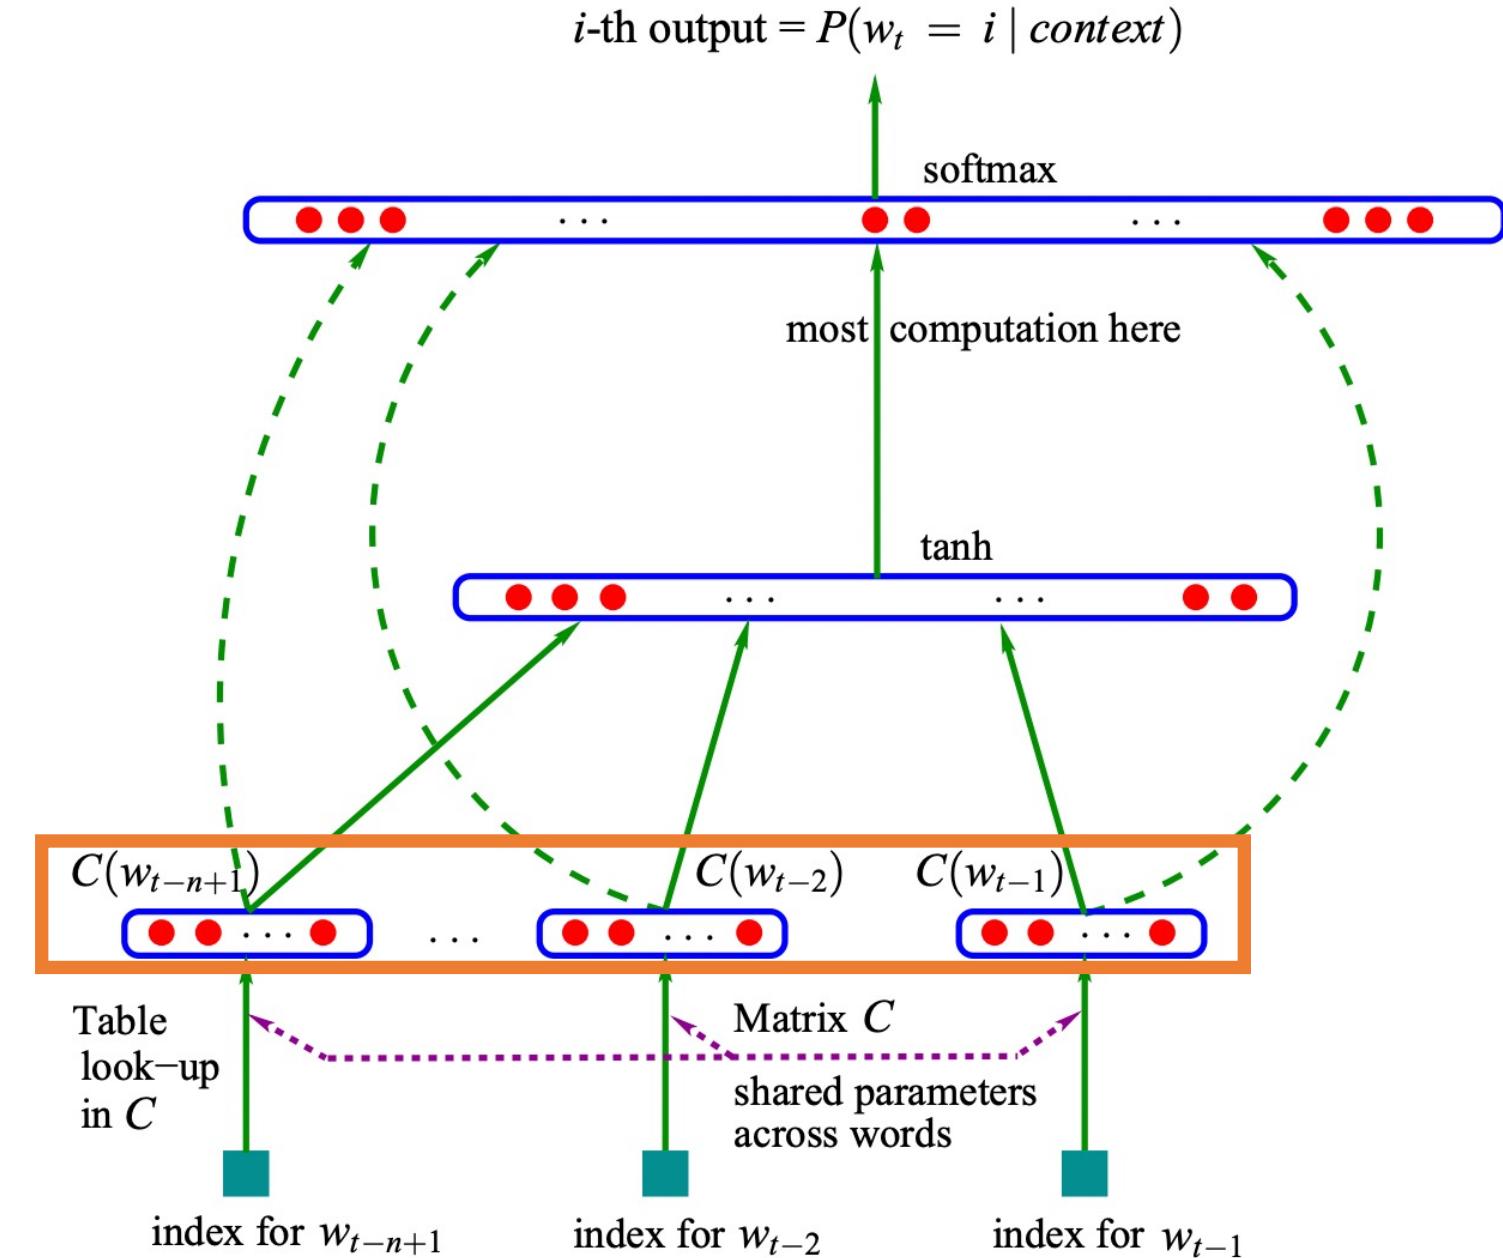


# Architecture

e.g., a vocabulary size of 17,000 was used with embedding sizes of 30, 60, and 100 in experiments

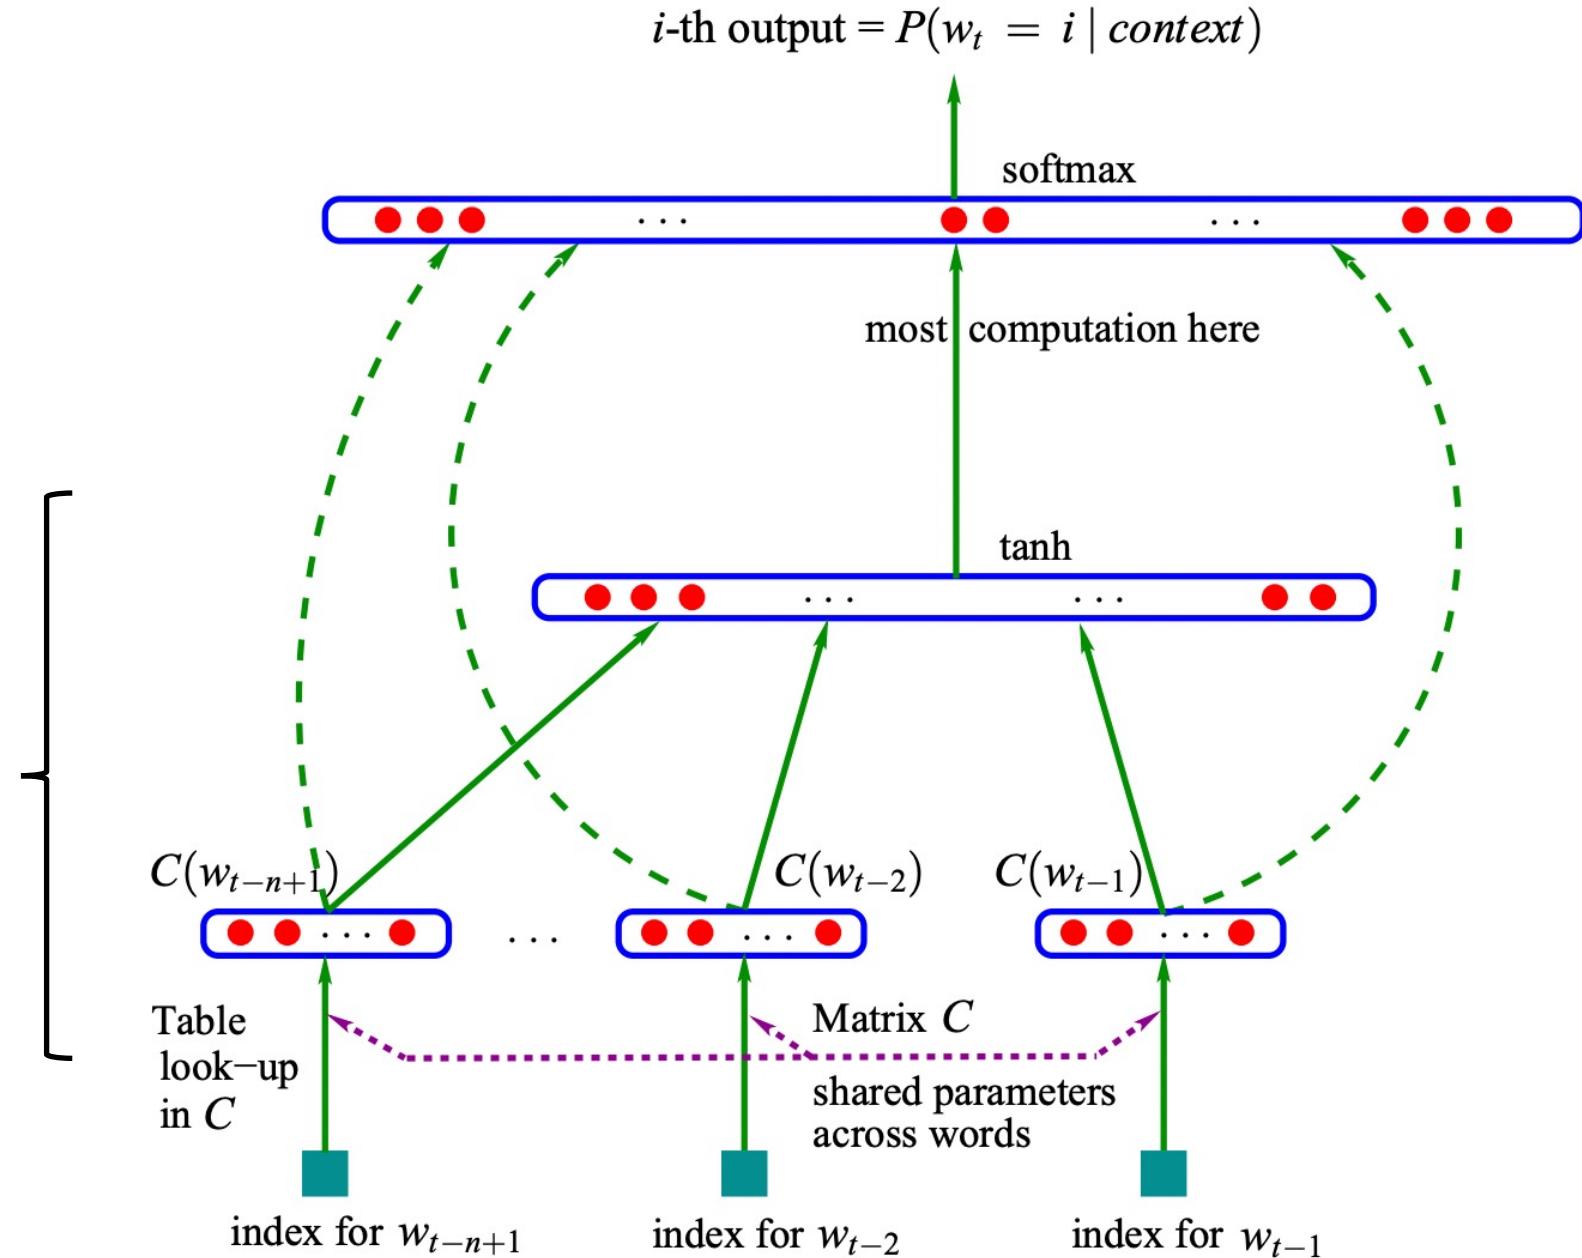
Assume a 30-d word embedding  
- what are the dimensions of each word embedding?

$1 \times 30$



# Architecture

Projection layer followed by a hidden layer with non-linearity

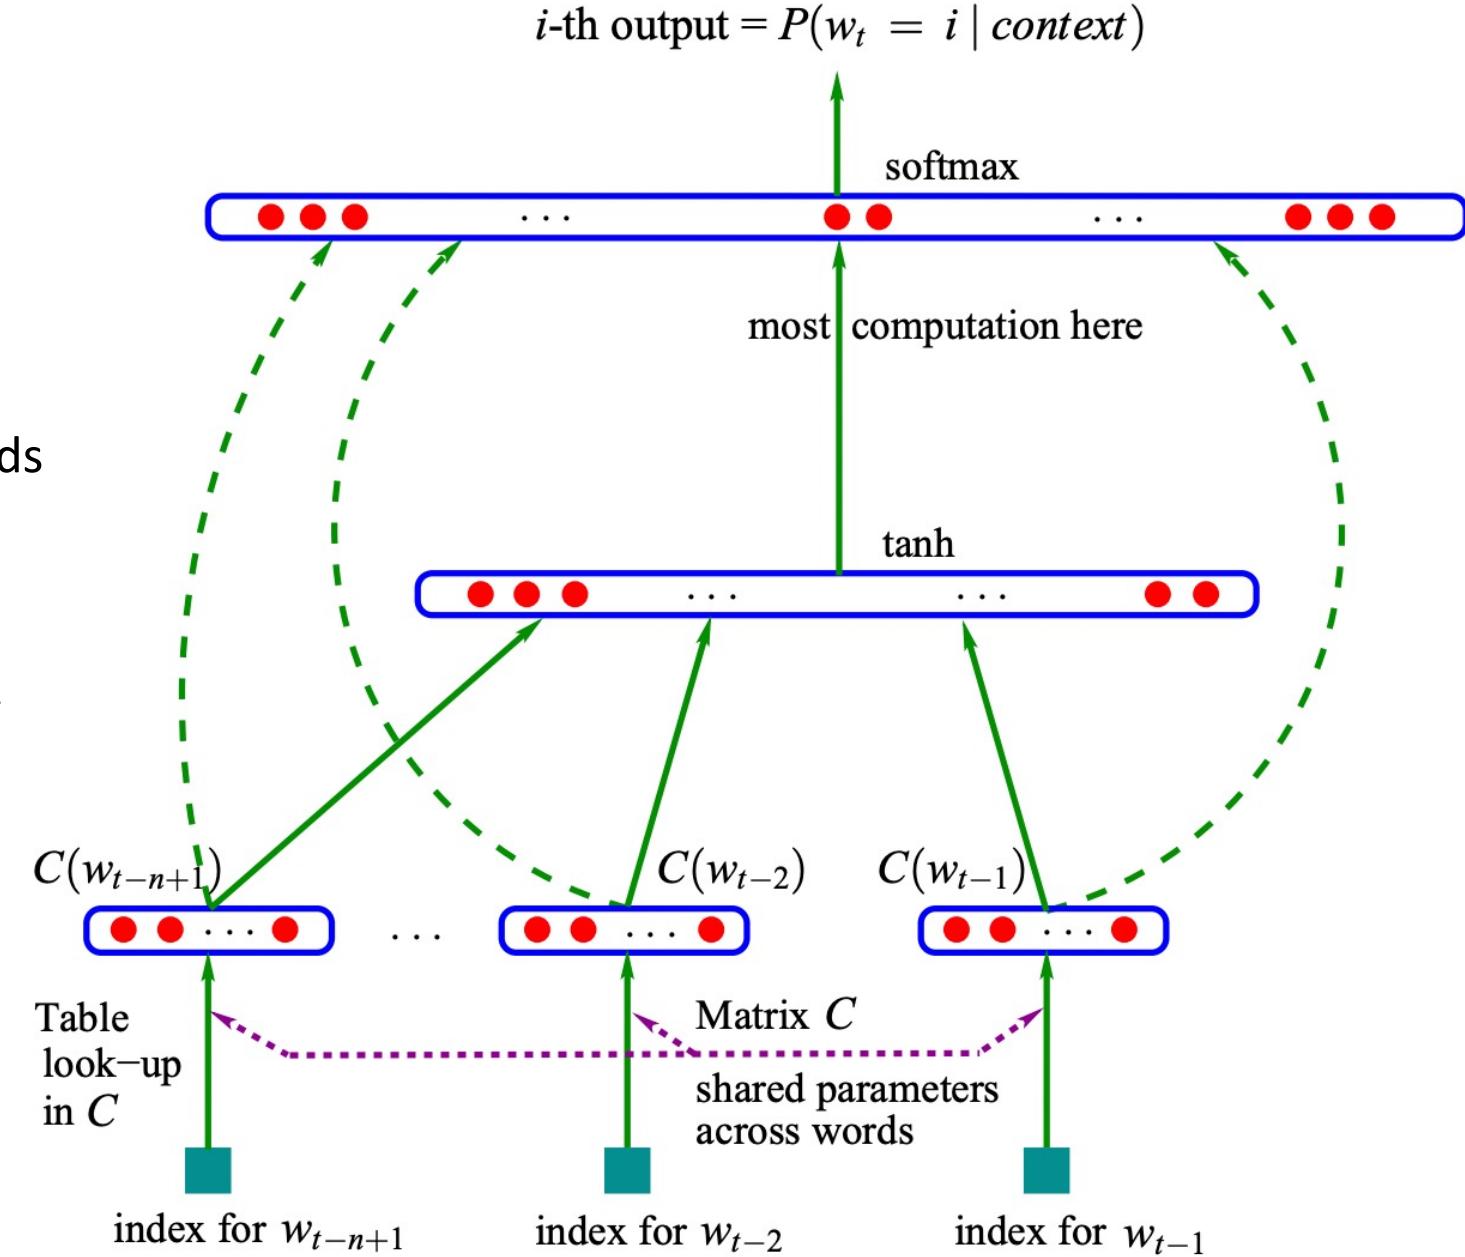


# Training

Use sliding window on input data; e.g., 3 words

Background music from a berimbau offers a beautiful escape...

Input: tried 1, 3, 5, and 8 input words  
and used 2 datasets with ~1 million and  
~34 million words respectively

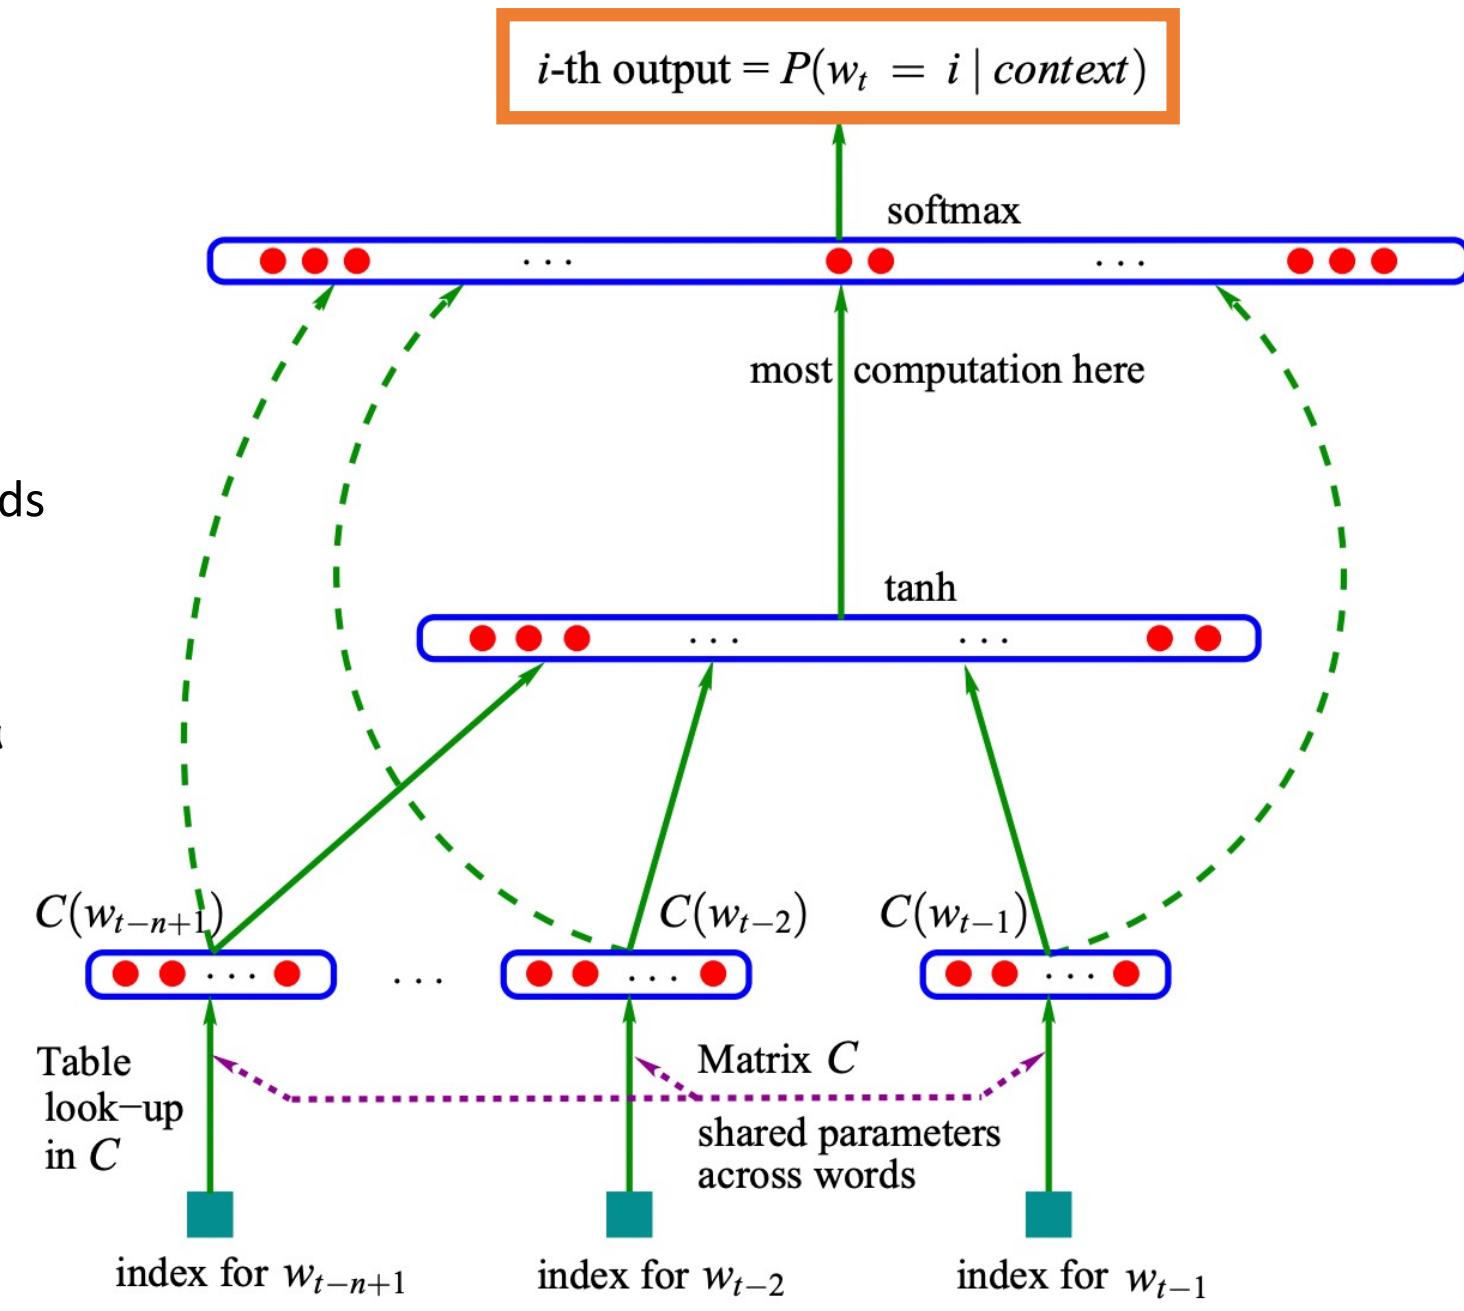


# Training

Use sliding window on input data; e.g., 3 words

Background music from a berimbau offers a beautiful escape...

Input: tried 1, 3, 5, and 8 input words  
and used 2 datasets with ~1 million and  
~34 million words respectively

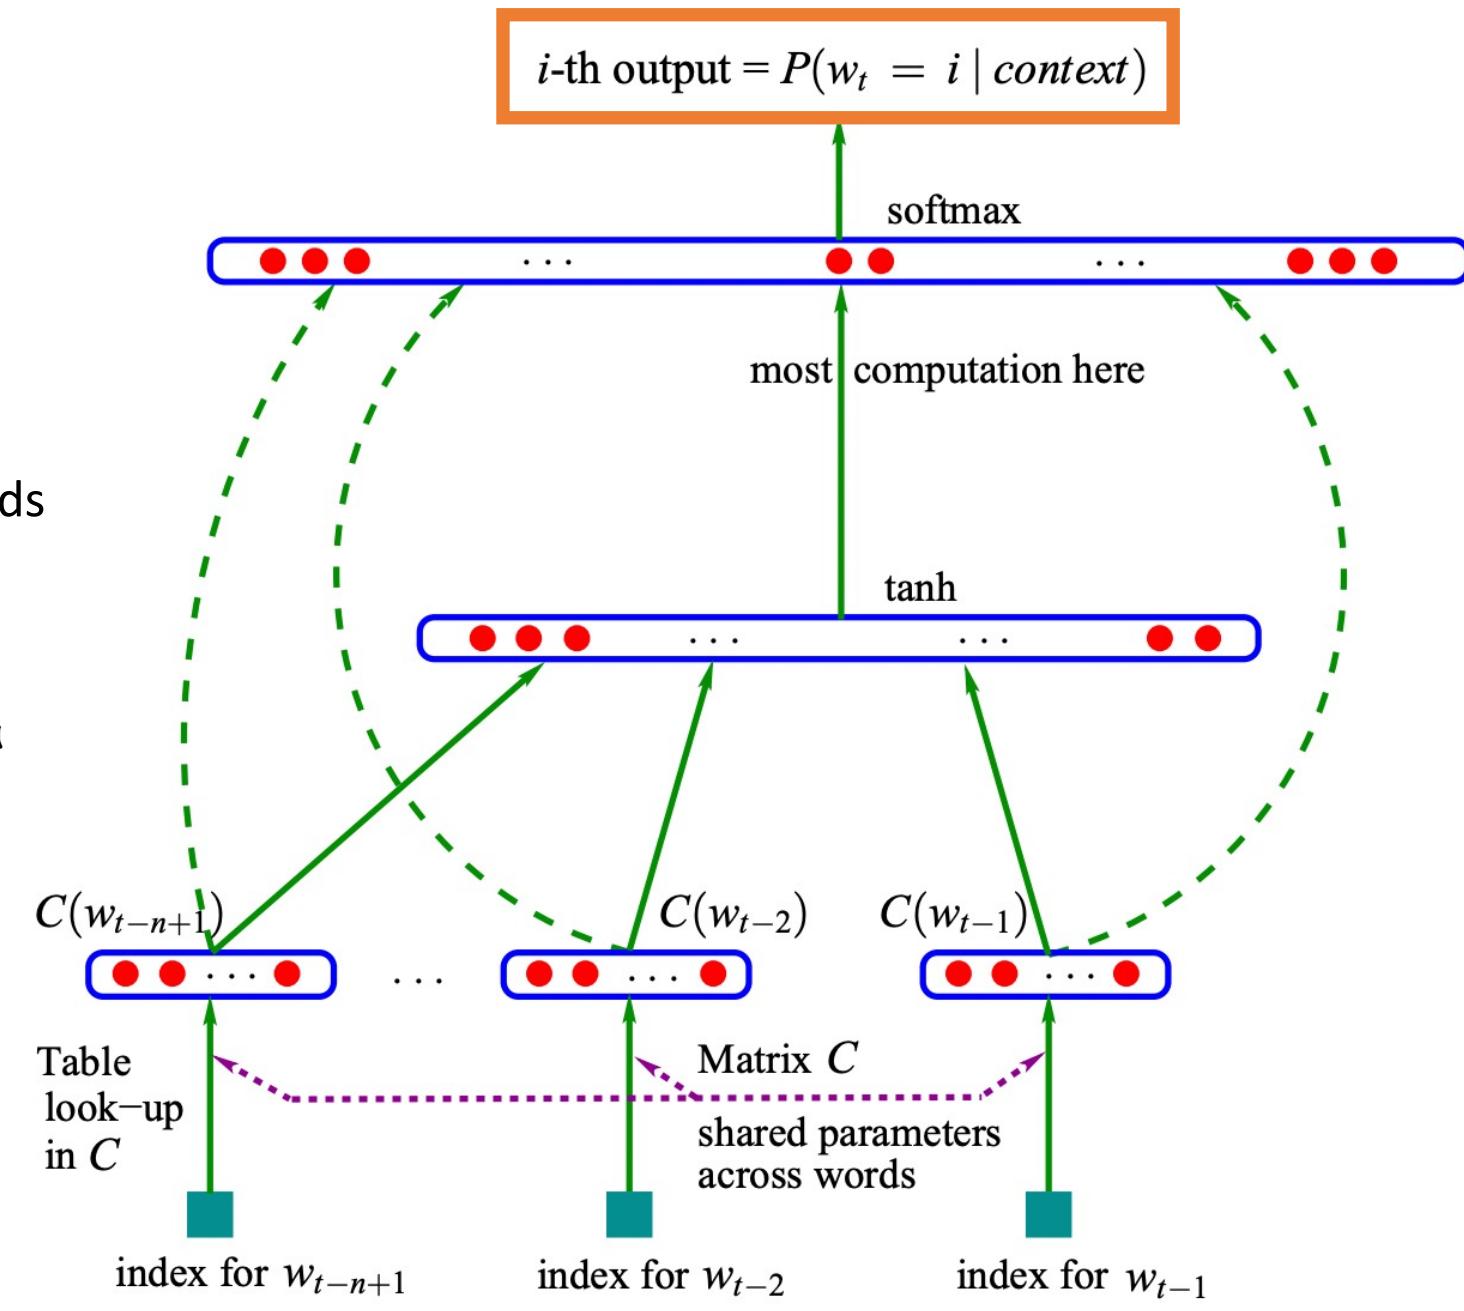


# Training

Use sliding window on input data; e.g., 3 words

Background music from a berimbau offers a beautiful escape...

Input: tried 1, 3, 5, and 8 input words  
and used 2 datasets with ~1 million and  
~34 million words respectively

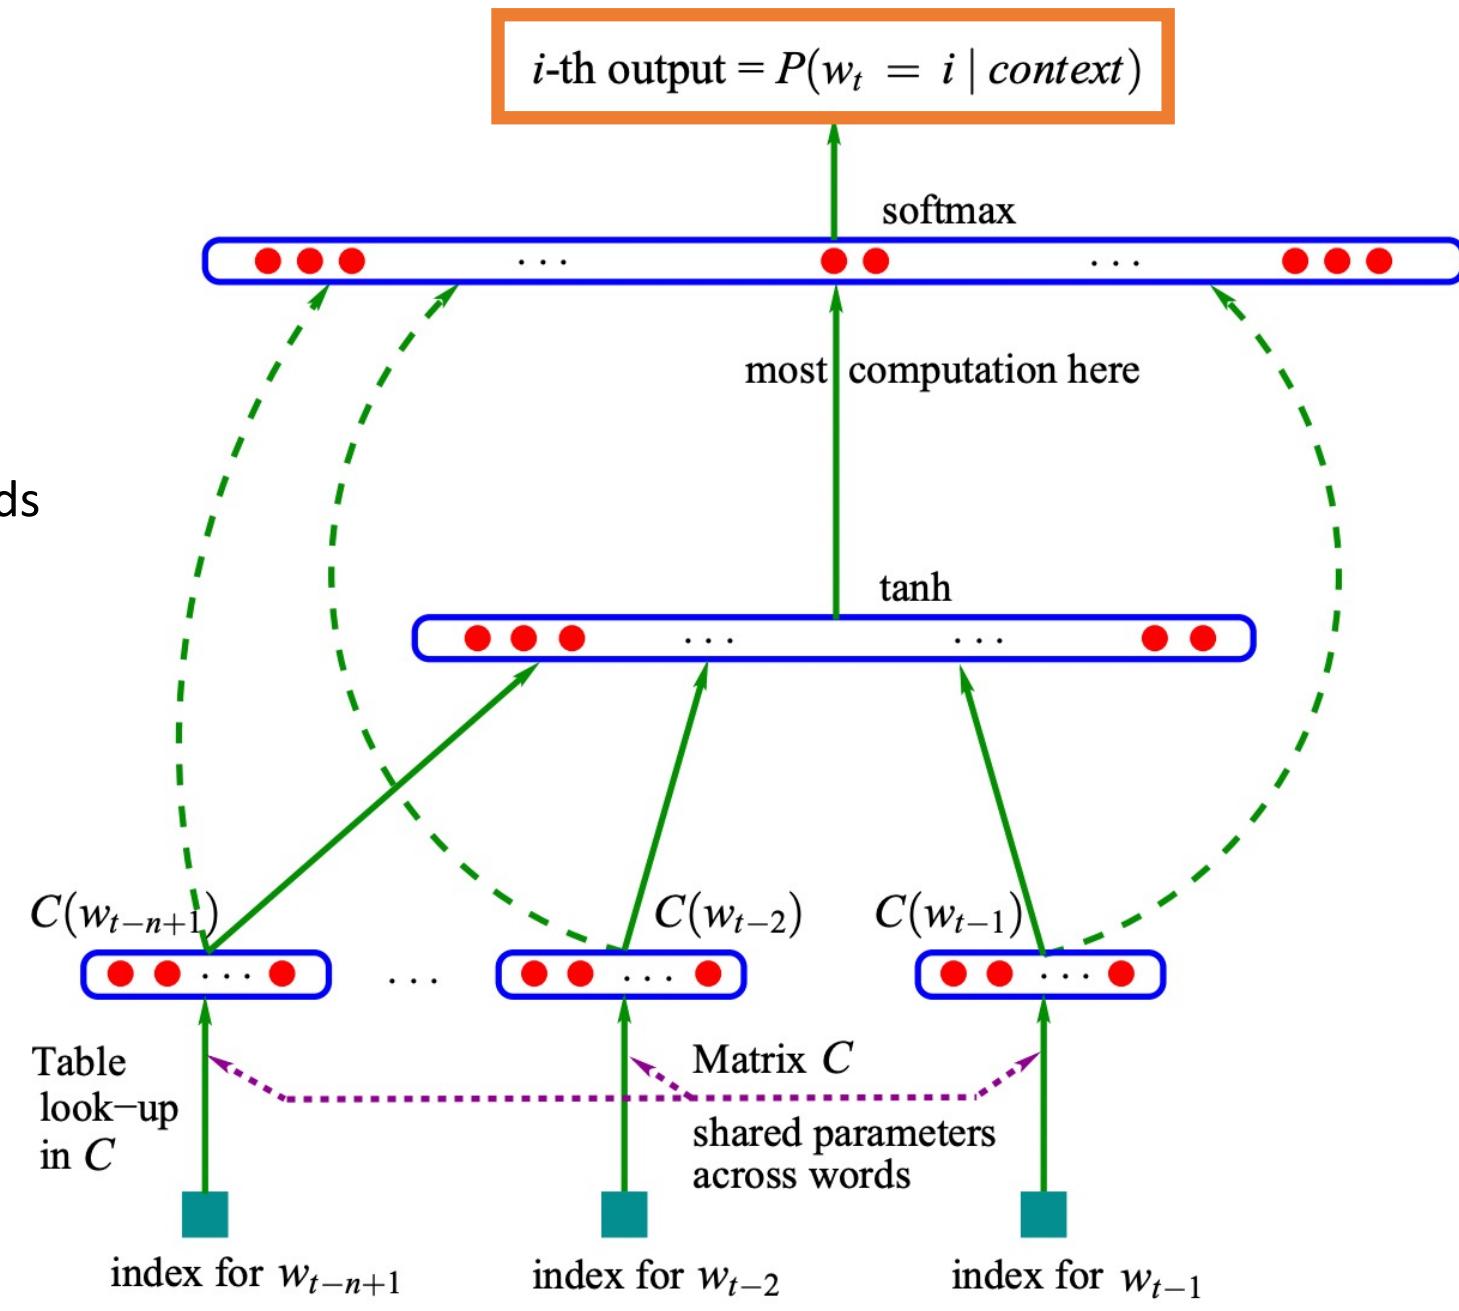


# Training

Use sliding window on input data; e.g., 3 words

Background music from a berimbau offers a beautiful escape...

Input: tried 1, 3, 5, and 8 input words  
and used 2 datasets with ~1 million and  
~34 million words respectively

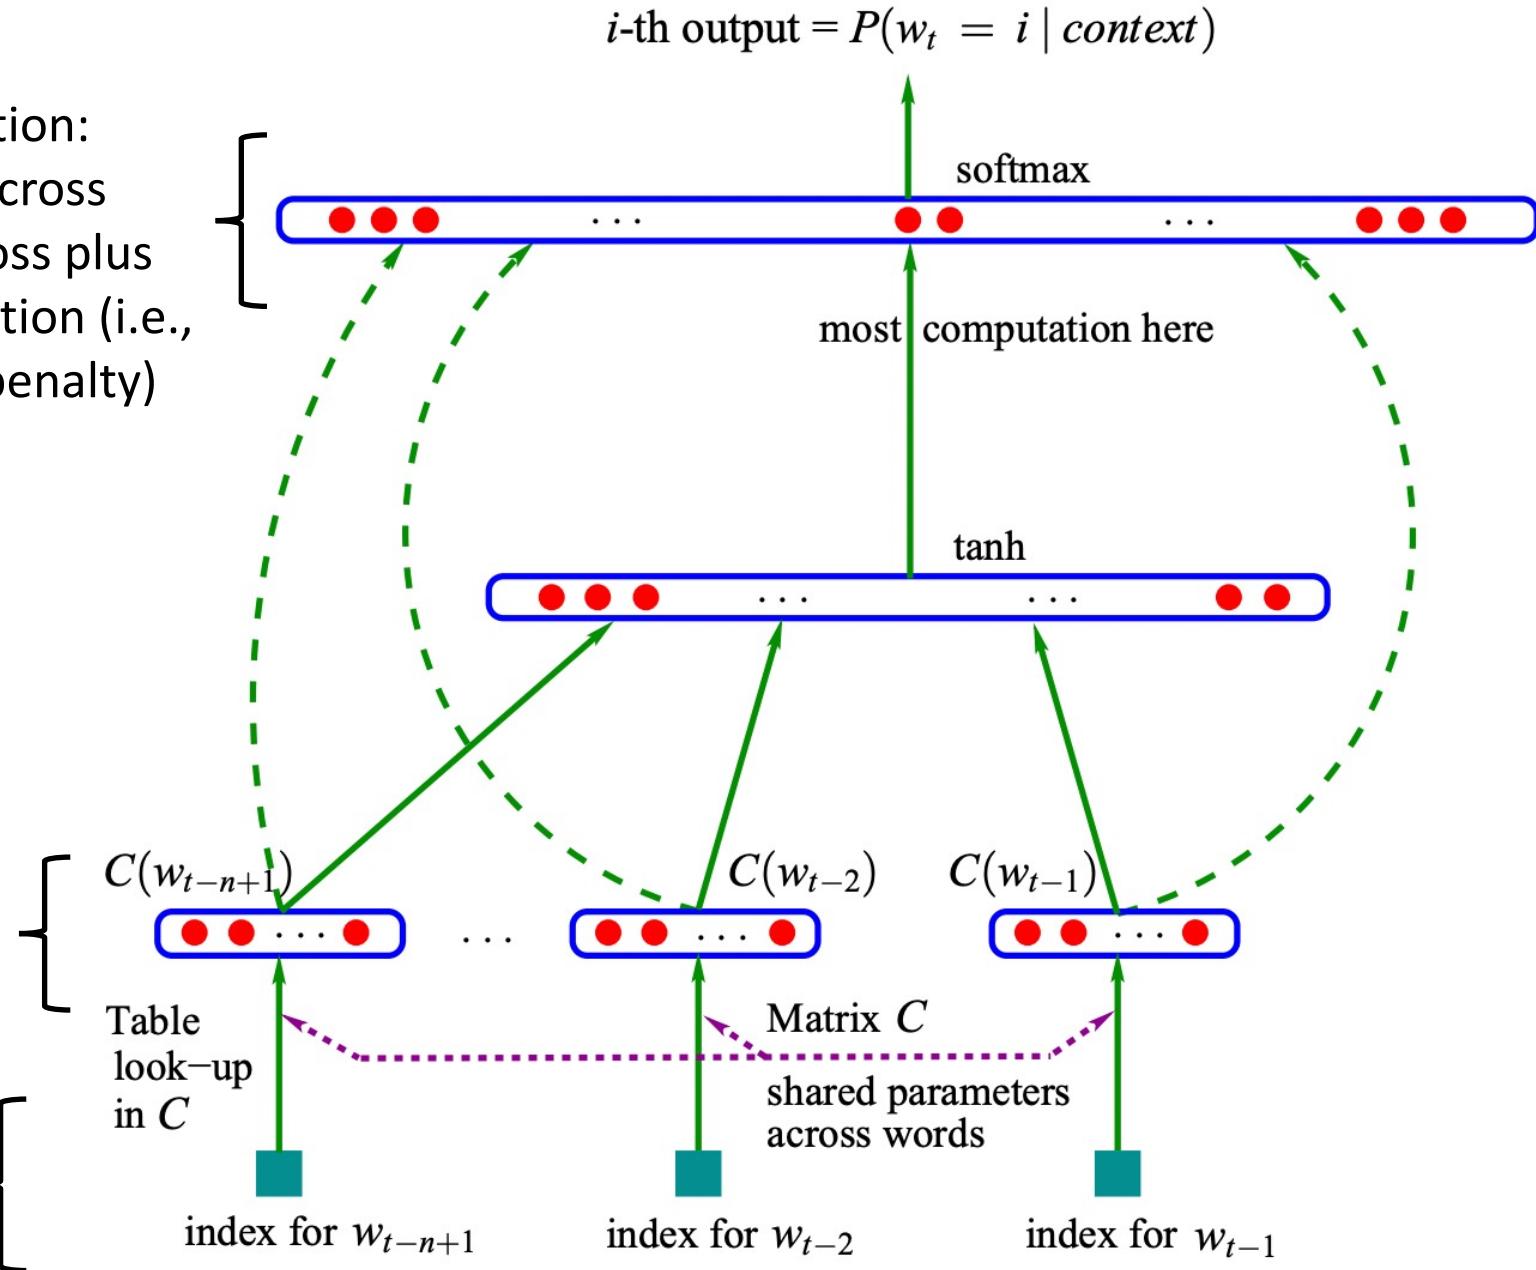


# Training

Input: tried 1, 3, 5, and 8 input words  
and used 2 datasets with ~1 million and  
~34 million words respectively

Word embedding iteratively updated

Cost function:  
minimize cross  
entropy loss plus  
regularization (i.e.,  
L2 norm penalty)



# Summary: Word Embeddings Learn Context of Previous Words Needed to Predict Next Word

e.g.,

1. Background music from a \_\_\_\_\_
2. Many people danced around the \_\_\_\_\_
3. I practiced for many years to learn how to play the \_\_\_\_\_

# Popular Word Embeddings

- Bengio method
- Word2vec (skip-gram model)
- And more...

# Idea: Learn Word Embeddings That Know What Are Viable Surrounding Words

e.g.,

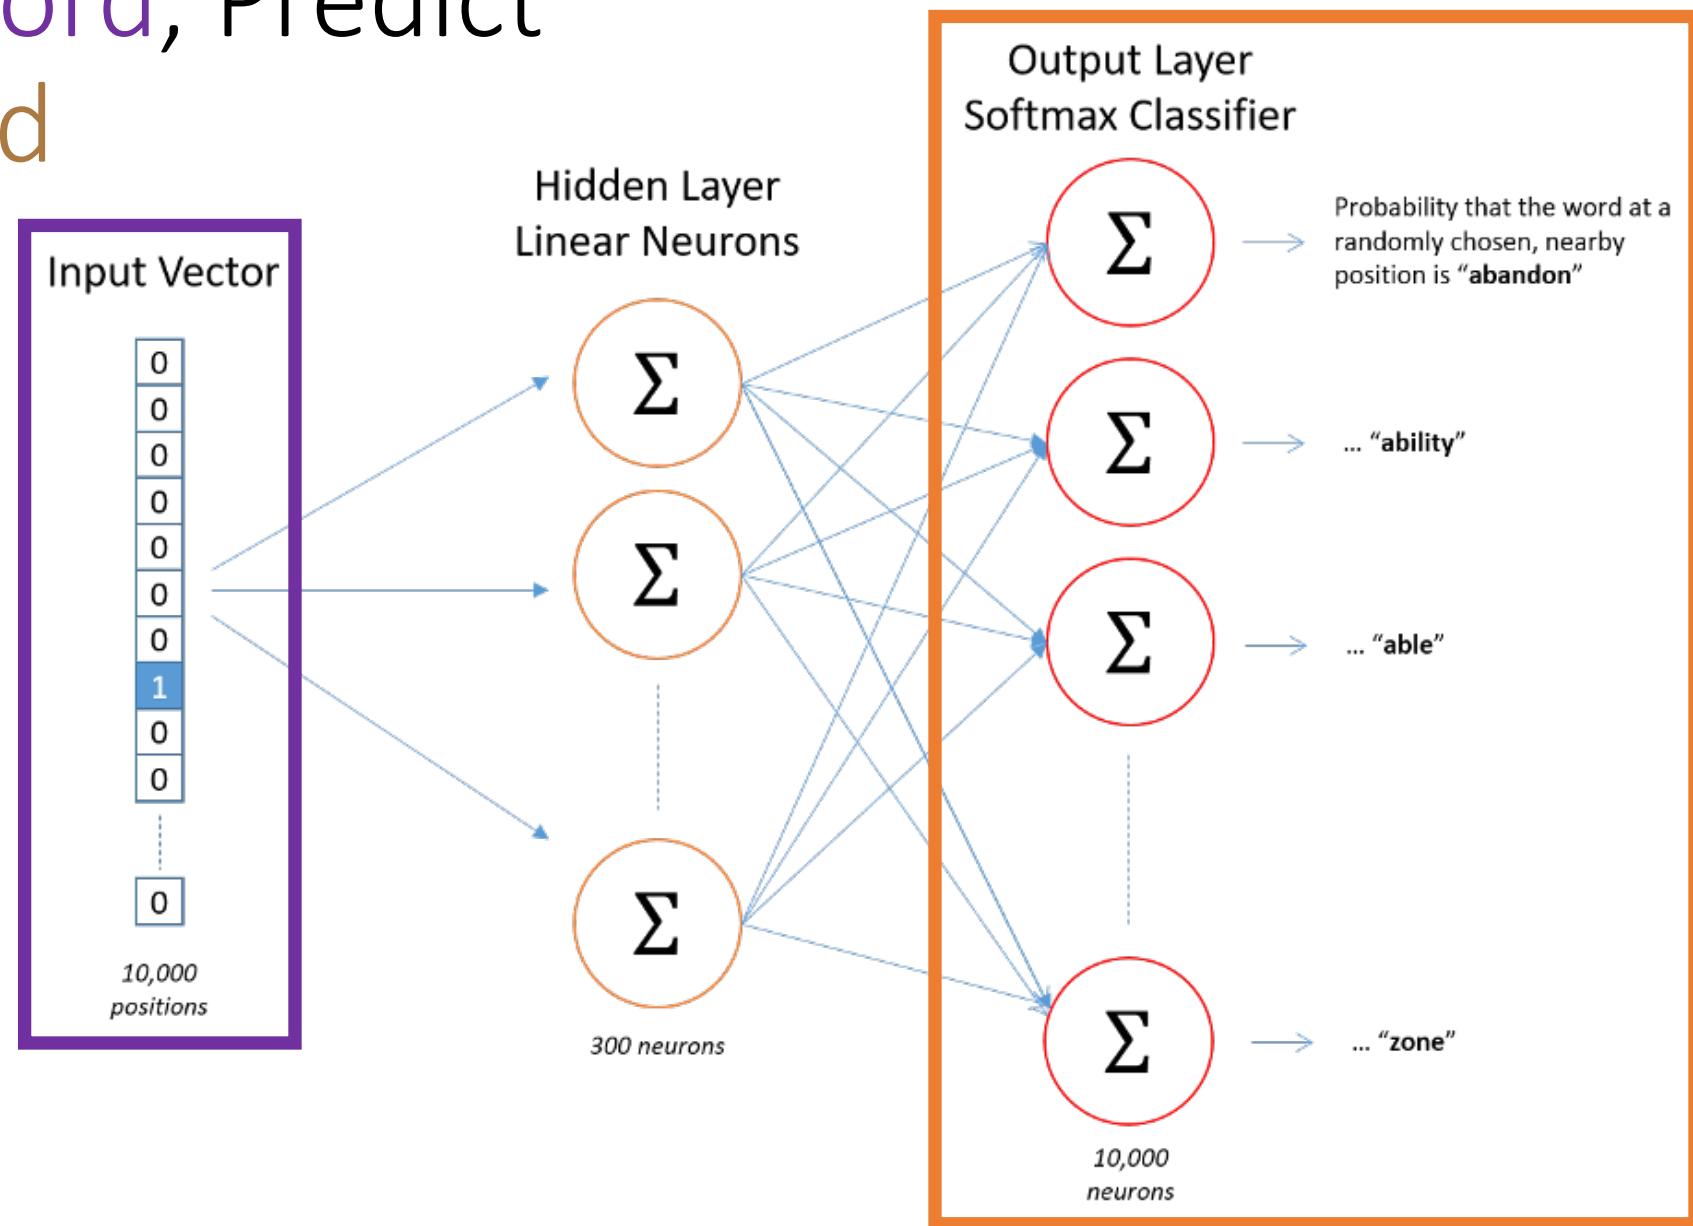
1. \_\_\_\_ berimbau \_\_\_\_
2. \_\_ berimbau \_\_

# Task: Given Word, Predict a Nearby Word

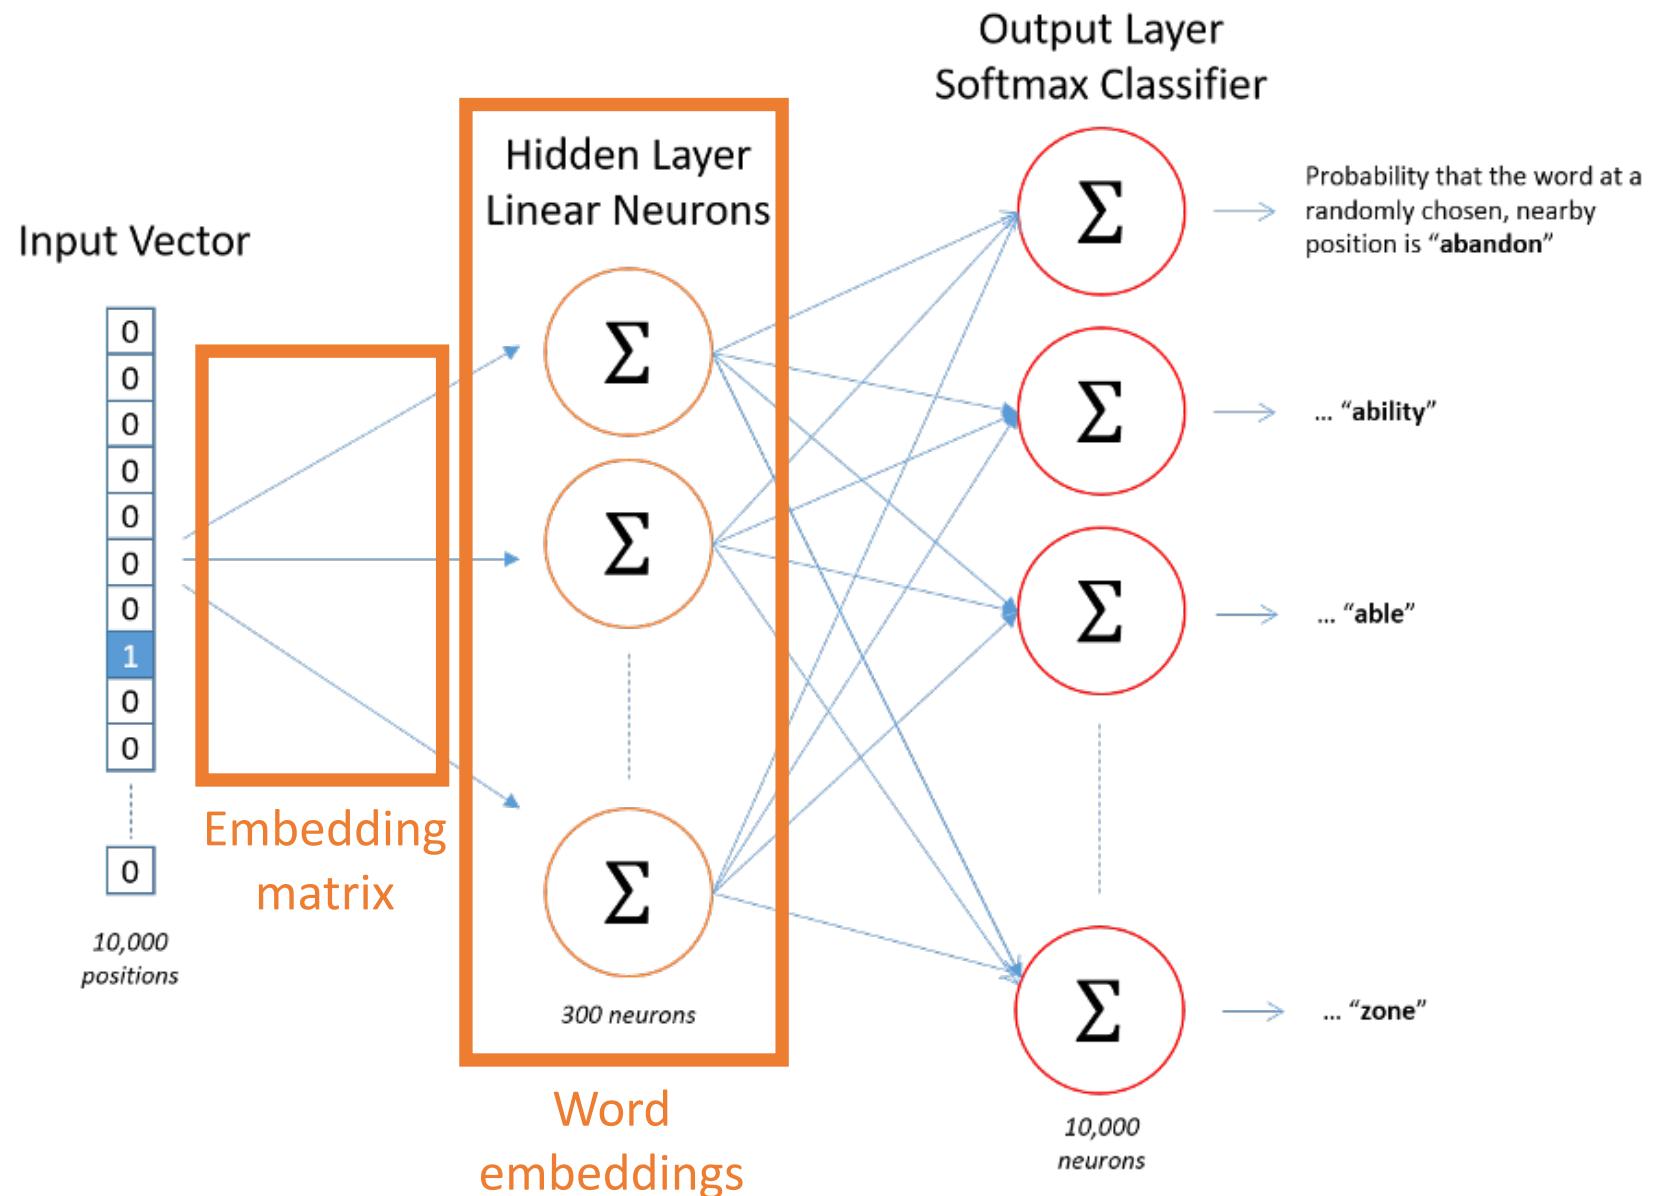
e.g.,

1. \_\_\_ \_\_\_ \_\_\_ \_\_\_ berimbau \_\_\_ \_\_\_ \_\_\_ \_\_\_
2. \_\_\_ berimbau \_\_\_

# Task: Given Word, Predict a Nearby Word



# Architecture

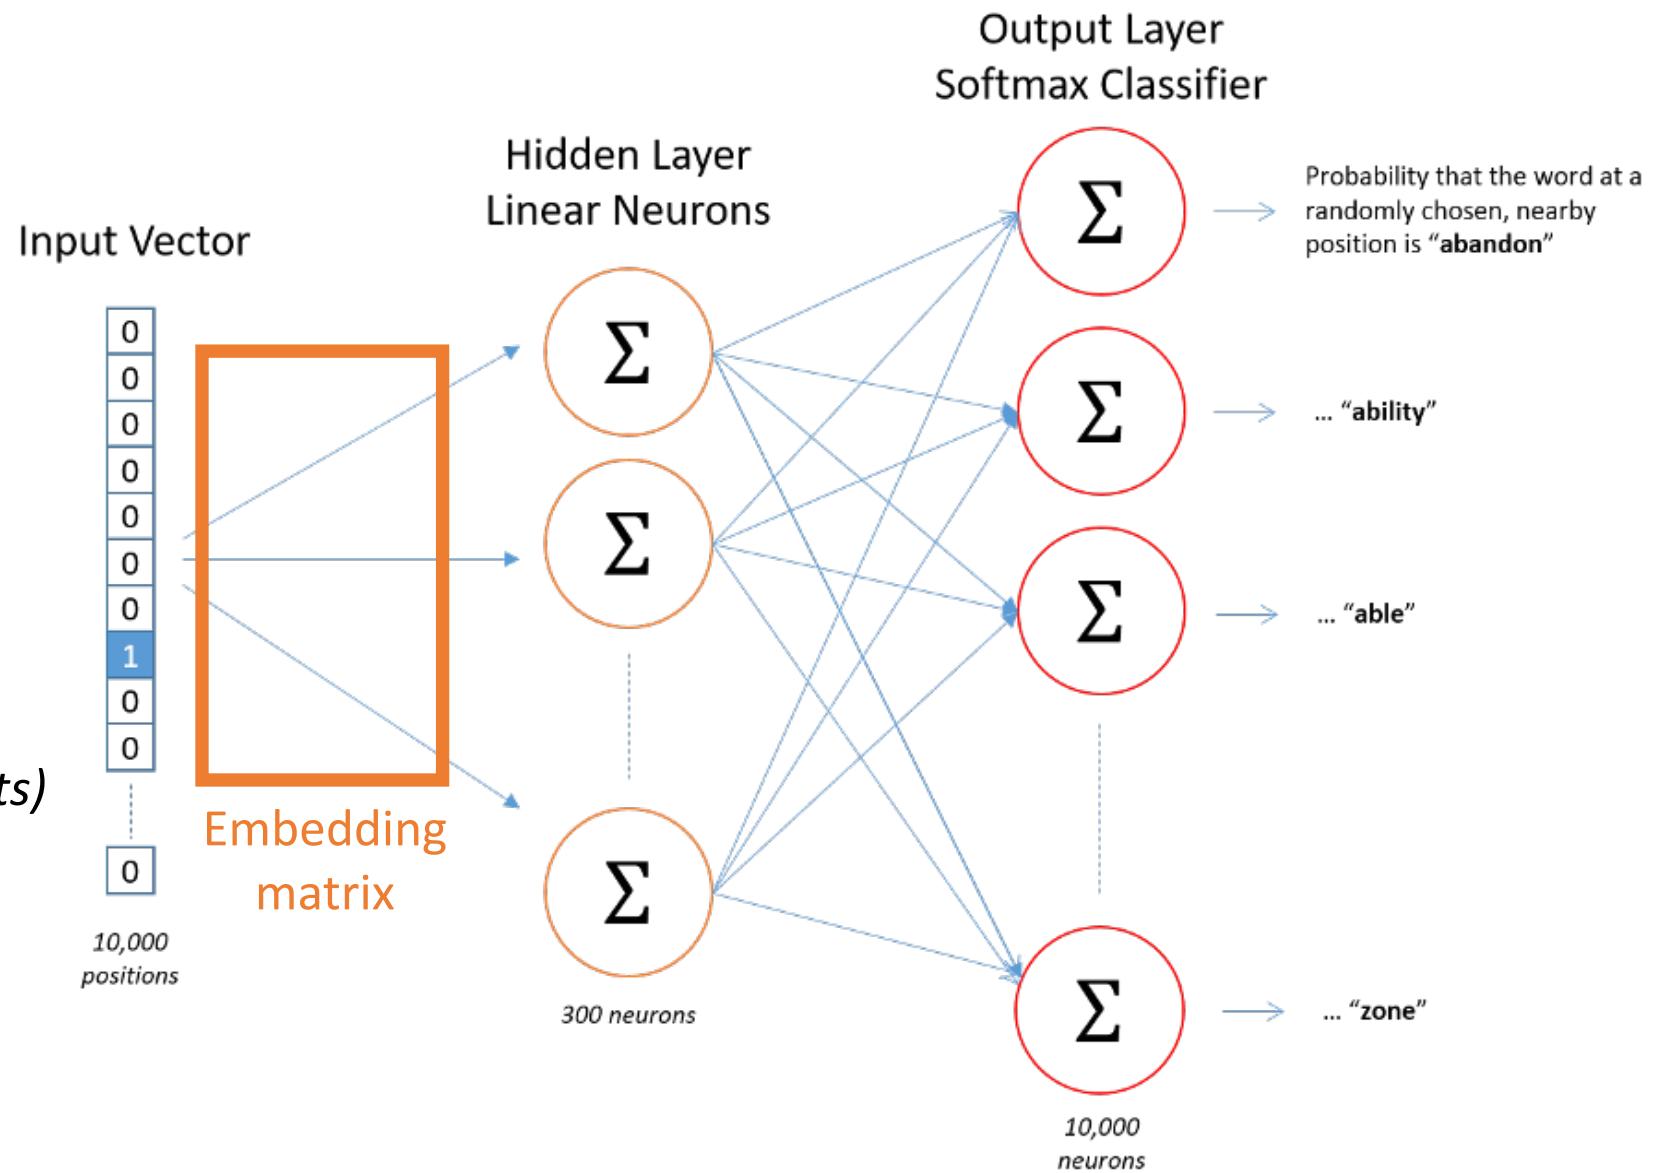


# Architecture

e.g., a vocabulary size of 10,000 is used with embedding sizes of 300

What are the dimensions of the embedding matrix?

$300 \times 10,000$  (i.e., 3,000,000 weights)

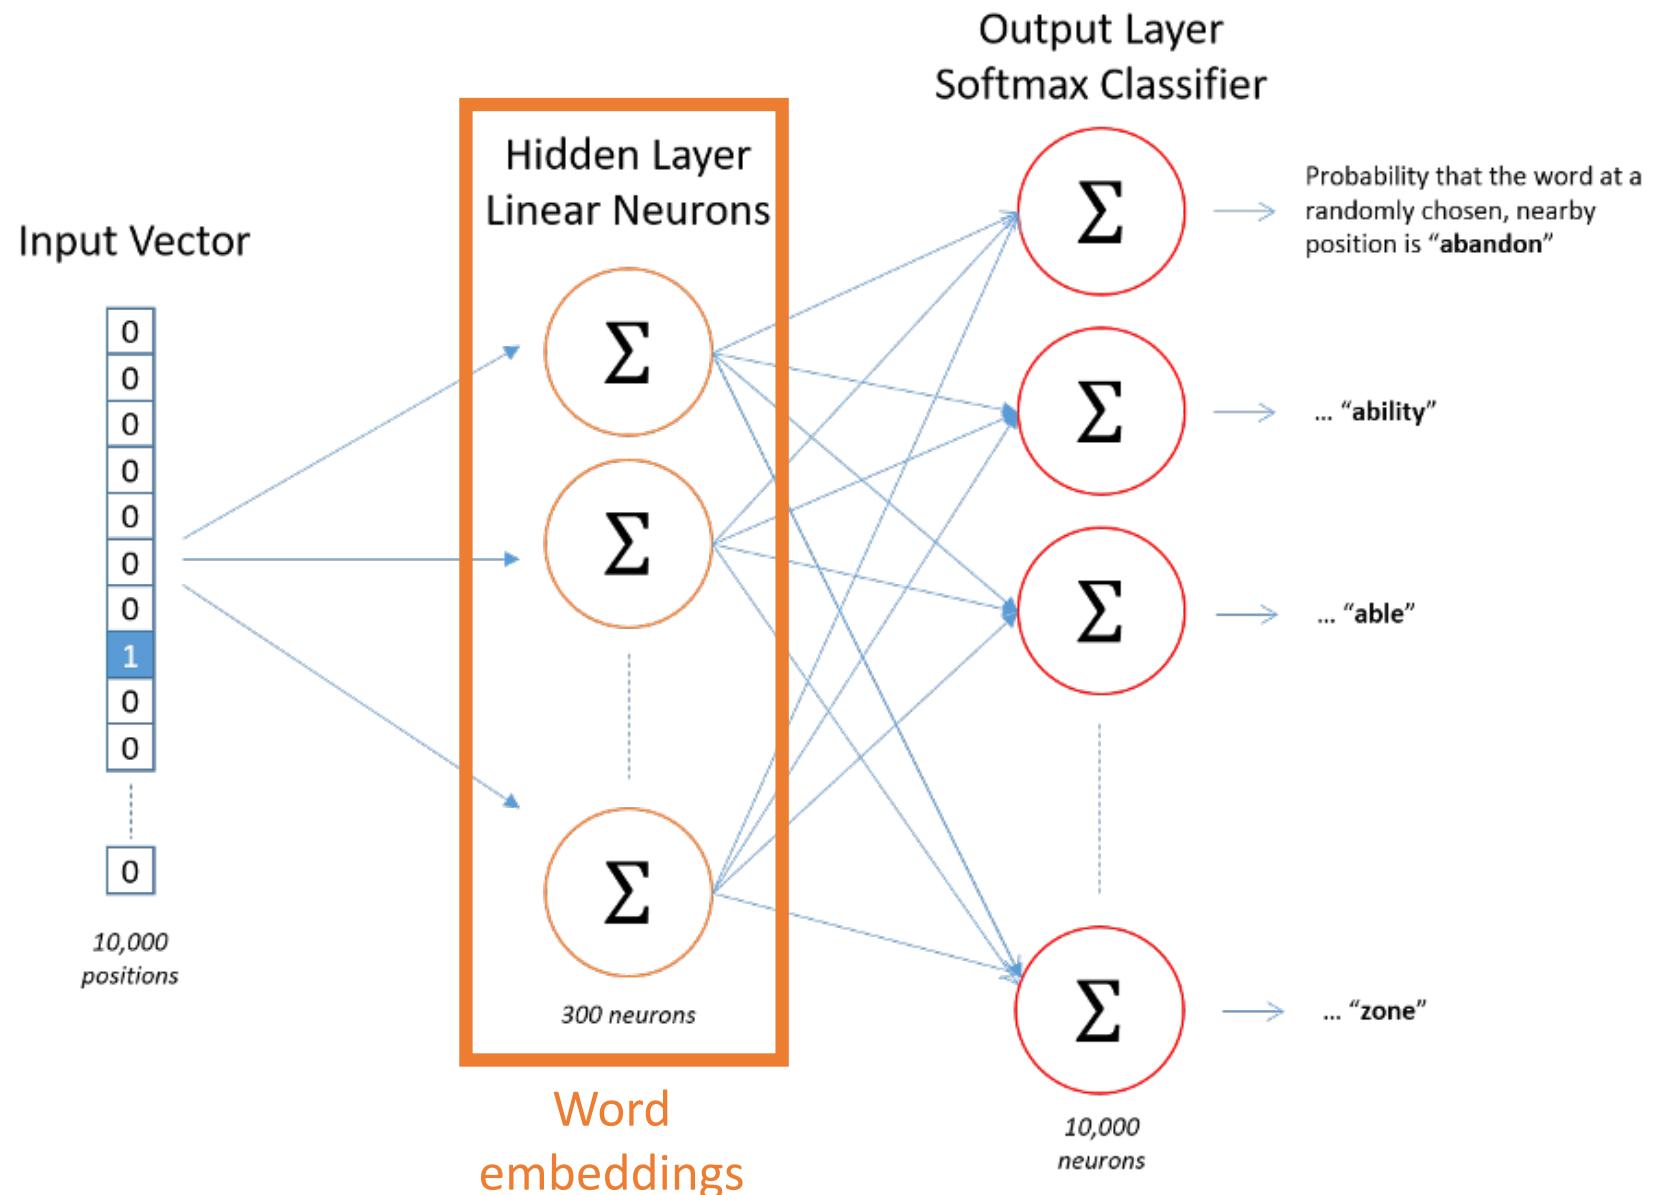


# Architecture

e.g., a vocabulary size of 10,000 is used with embedding sizes of 300

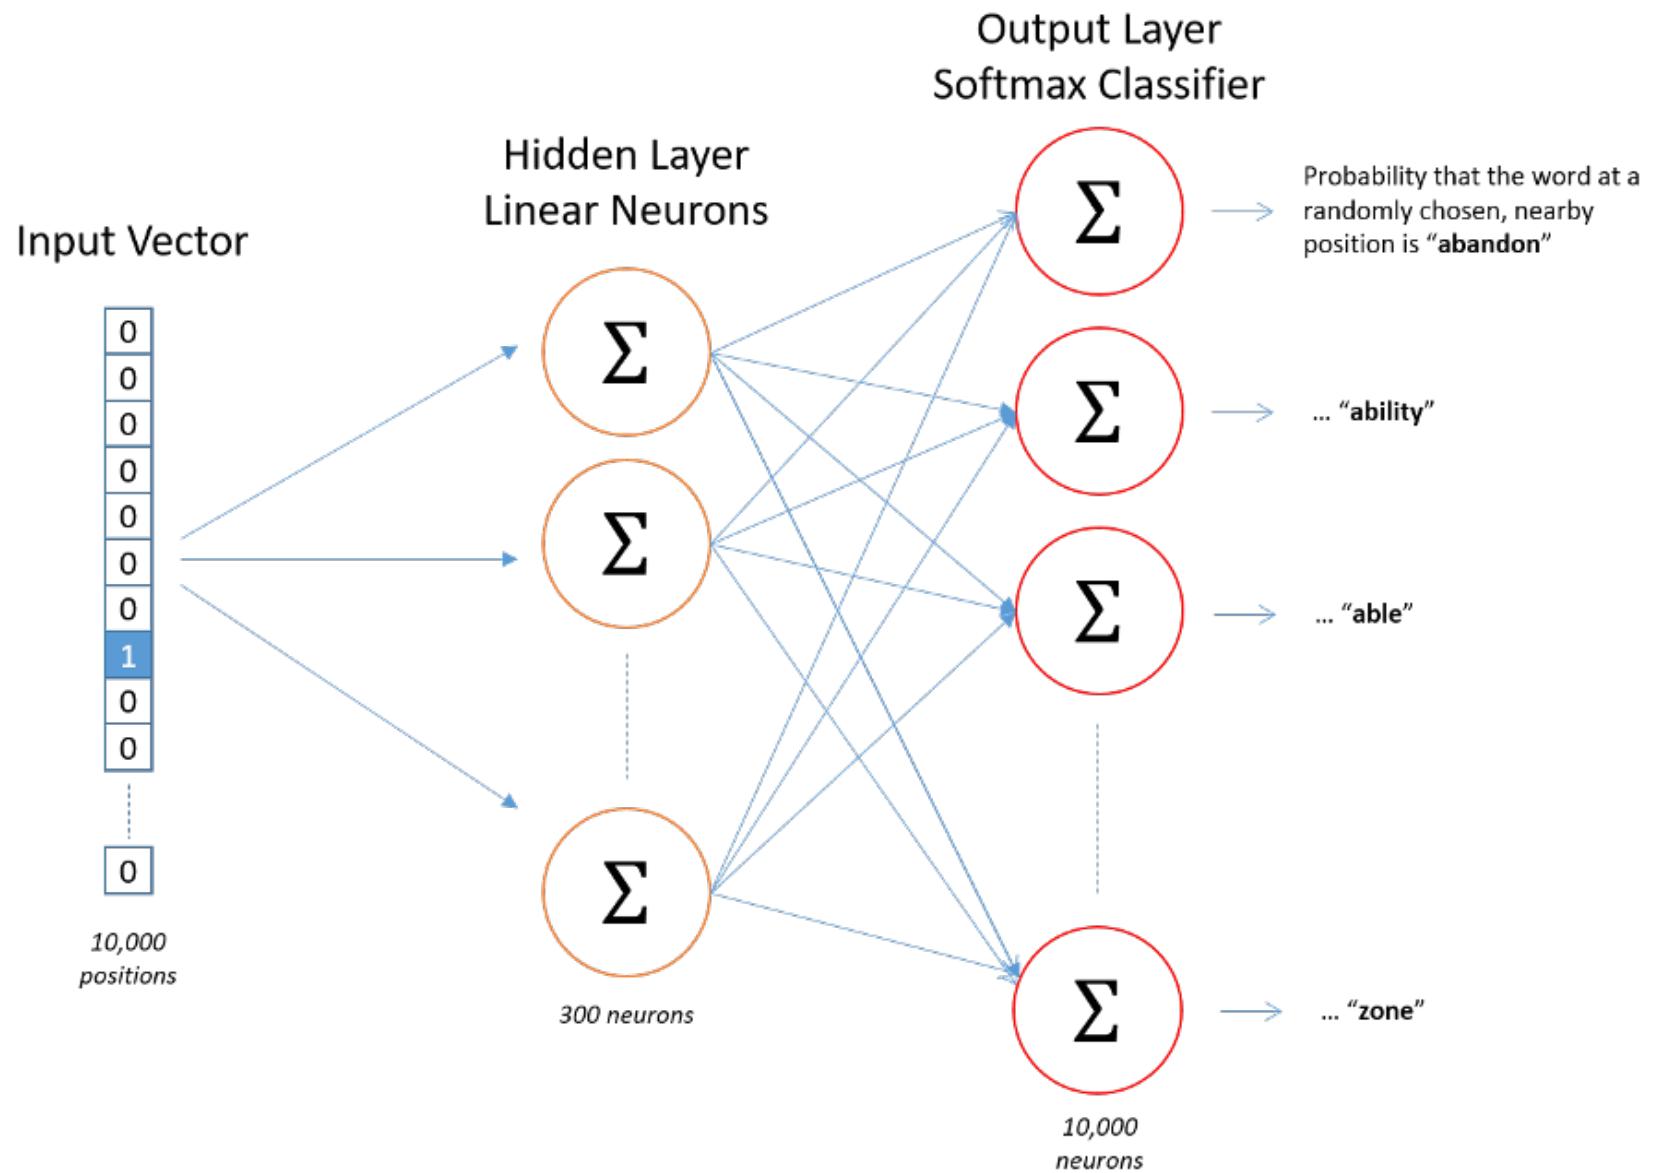
What are the dimensions of each word embedding?

$1 \times 300$



# Architecture

A shallower, simpler architecture than the Bengio approach (i.e., lacks a non-linear hidden layer)!



# Training

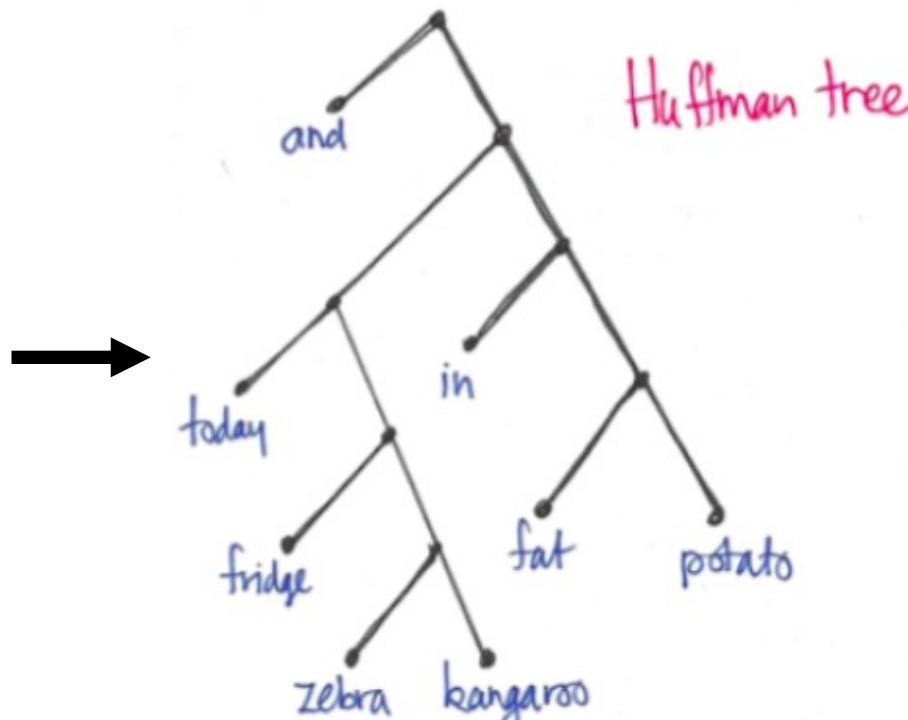
Sliding window run on input data to sample neighbors of each **target word** (e.g., using window size of 2)

Source Text	Training Samples
The <b>quick</b> brown fox jumps over the lazy dog. ➔	(the, quick) (the, brown)
The <b>quick</b> brown <b>fox</b> jumps over the lazy dog. ➔	(quick, the) (quick, brown) (quick, fox)
The <b>quick</b> brown <b>fox</b> jumps over the lazy dog. ➔	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
The quick brown <b>fox</b> jumps <b>over</b> the lazy dog. ➔	(fox, quick) (fox, brown) (fox, jumps) (fox, over)

# Extra Tricks: More Efficient Representations

1. Change output layer to hierarchical softmax

word	count
fat	3
fridge	2
zebra	1
potato	3
and	14
in	7
today	4
kangaroo	2



2. Reformulate problem to perform negative sampling

Binary classification: predict for a given word if another word is nearby

- Positive examples: observed target and neighboring words
- Negative examples: randomly sampled other words

# Hyperparameters: What Works Well?

- Word embedding dimensionality?
  - Dimensionality set between 100 and 1,000
- Context window size?
  - $\sim 10$

# Very Exciting/Surprising Finding

- Vector arithmetic with word embeddings can solves many analogies  
*(Full test list: <http://download.tensorflow.org/data/questions-words.txt>)*
- Semantic relationships (meaning of words in a sentence):
  - Italy + (Paris - France) = Rome
- Syntactic relationships (rules for words in a sentence)
  - smallest + (big – small) = biggest
  - think + (read – reading) = thinking
  - mouse + (dollars – dollar) = mice

# Summary: Word Embeddings Are Learned that Support Predicting Viable Surrounding Words!

e.g.,

1. \_\_\_\_ berimbau \_\_\_\_
2. \_\_\_\_ berimbau \_\_\_\_

# Popular Word Embeddings

- Bengio method
- Word2vec (skip-gram model)
- And more...

# Variants for Learning Word Embeddings

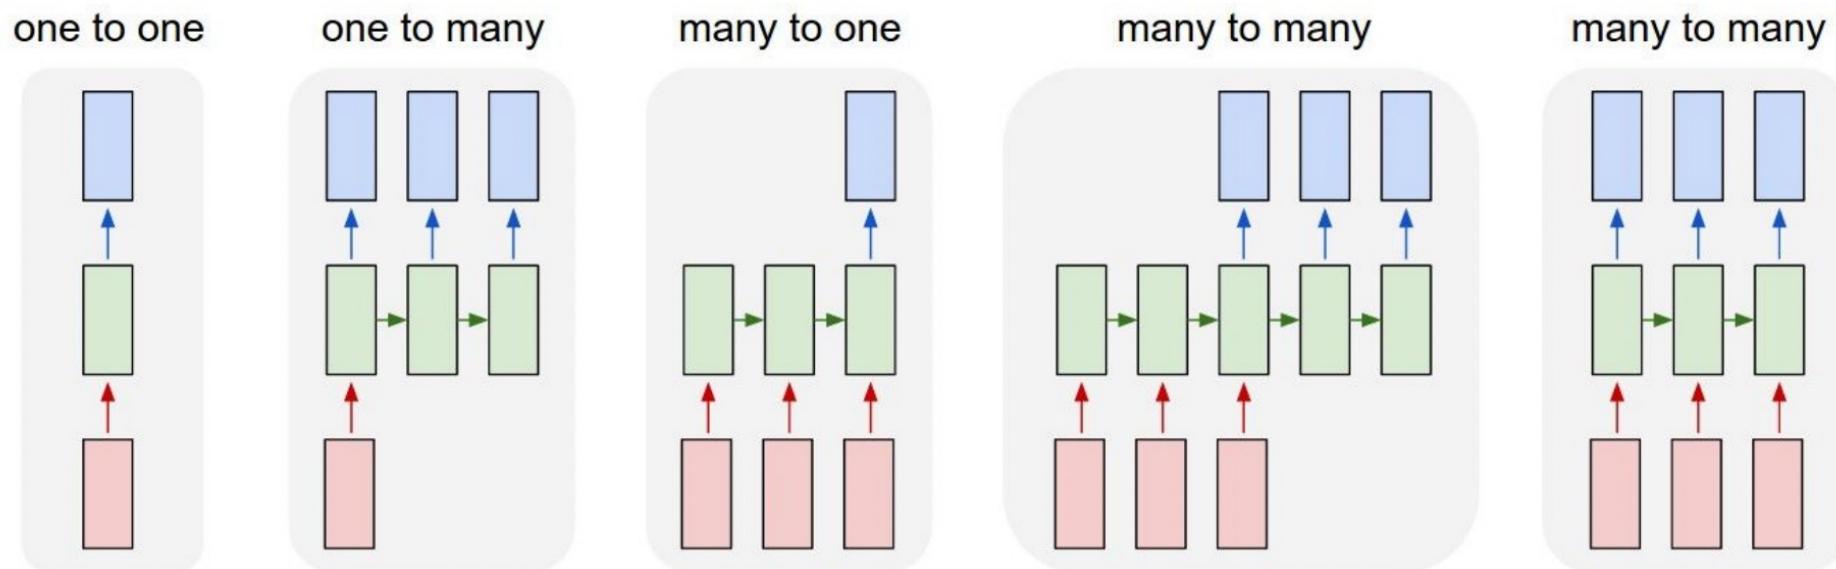
- Capture global context rather than just local context of previous or surrounding words; e.g.,
  - GloVe for Global Vectors (Pennington et al., 2014)
- Capture that the same word can have different word vectors under different contexts; e.g.,
  - Elmo for embeddings from language models (Peters et al., arXiv 2018)
- Support multiple languages; e.g.,
  - Fast-text (Bojanowski et al., 2016)

# Popular Word Embeddings

- Bengio method
- Word2vec (skip-gram model)
- And more...

# Recap of Big Picture

- Convert words into compact vectors as **input** to neural networks; e.g., RNNs



- Implementation detail: may need to learn extra tokens such as “UNK” and “EOS” to represent out of vocabulary words and signify end of the string respectively

- Also, can fine-tune word embedding matrices for different applications

# Word Embedding Limitations/Challenges

- Distinguish antonyms from synonyms
  - Antonyms are learned near each other in the embedding space since they are commonly used in similar contexts: “I **hate** math” vs “I **love** math” or “Take a **right** turn” vs “Take a **left** turn”

- Gender bias:

**Man is to Computer Programmer as Woman is to Homemaker? Debiasing Word Embeddings**

**Tolga Bolukbasi<sup>1</sup>, Kai-Wei Chang<sup>2</sup>, James Zou<sup>2</sup>, Venkatesh Saligrama<sup>1,2</sup>, Adam Kalai<sup>2</sup>**

<sup>1</sup>Boston University, 8 Saint Mary’s Street, Boston, MA

<sup>2</sup>Microsoft Research New England, 1 Memorial Drive, Cambridge, MA

[tolgab@bu.edu](mailto:tolgab@bu.edu), [kw@kwchang.net](mailto:kw@kwchang.net), [jamesyzou@gmail.com](mailto:jamesyzou@gmail.com), [srv@bu.edu](mailto:srv@bu.edu), [adam.kalai@microsoft.com](mailto:adam.kalai@microsoft.com)

# Word Embedding Limitations/Challenges

- Distinguish antonyms from synonyms
  - Antonyms are learned near each other in the embedding space since they are commonly used in similar contexts: “I hate math” vs “I love math” or “Take a right turn” vs “Take a left turn”

- Gender bias:

<b>Extreme <i>she</i></b>	<b>Extreme <i>he</i></b>
1. homemaker	1. maestro
2. nurse	2. skipper
3. receptionist	3. protege
4. librarian	4. philosopher
5. socialite	5. captain
6. hairdresser	6. architect
7. nanny	7. financier
8. bookkeeper	8. warrior
9. stylist	9. broadcaster
10. housekeeper	10. magician

<b>Gender stereotype <i>she-he</i> analogies</b>		
sewing-carpentry	registered nurse-physician	housewife-shopkeeper
nurse-surgeon	interior designer-architect	softball-baseball
blond-burly	feminism-conservatism	cosmetics-pharmaceuticals
giggle-chuckle	vocalist-guitarist	petite-lanky
sassy-snappy	diva-superstar	charming-affable
volleyball-football	cupcakes-pizzas	lovely-brilliant
<b>Gender appropriate <i>she-he</i> analogies</b>		
queen-king	sister-brother	mother-father
waitress-waiter	ovarian cancer-prostate cancer	convent-monastery

# Word Embedding Limitations/Challenges

- Distinguish antonyms from synonyms
  - Antonyms are learned near each other in the embedding space since they are commonly used in similar contexts: “I hate math” vs “I love math” or “Take a right turn” vs “Take a left turn”
- Gender bias
- What other language biases do you think could be learned?

# Today's Topics

- Introduction to natural language processing
- Text representation
- Neural word embeddings
- Programming tutorial

# Today's Topics

- Introduction to natural language processing
- Text representation
- Neural word embeddings
- Programming tutorial

*The End*

# Word Embeddings

Course Instructor: Chandresh  
AI Lab Coordinator @IIT Indore

# Course Content

Module I: History of Deep Learning, Sigmoid Neurons, Perceptrons, and learning algorithms. Multilayer Perceptrons (MLPs), Representation Power of MLPs.,

Module II: Feedforward Neural Networks. Backpropagation. first and second-order training methods. NN Training tricks, Regularization

Module III: Introduction to Autoencoders and their characteristics, relation to PCA, Regularization in autoencoders, and Types of autoencoders, Variational Autoencoder

Module IV: Architecture of Convolutional Neural Networks (CNN), types of CNNs. Image classification, Pre-training vs fine-tuning.- representation learning, Object Detection and Semantic Segmentation

Module V: Architecture of Recurrent Neural Networks (RNN), **Word Embeddings**,  
Backpropagation through time. Encoder-Decoder Models, Attention Mechanism. Advanced Topics:  
Transformers and BERT. Nodule VI: Gen AI- Deep generative models: VAE, GAN,

# Acknowledgement

- Russ Salakhutdinov and Hugo Larochelle's class on Neural Networks
- Neural Networks and Deep Learning course by Danna Gurari University of Colorado Boulder